

Compare Results

Old File:

P4 document Jordi van Liempt.pdf

103 pages (1,17 MB)

22-5-2020 01:01:18

versus

New File:

thesis.pdf

159 pages (5,59 MB)

1-7-2020 01:37:13

Total Changes

1943

Content

504
582
329

Replacements
Insertions
Deletions

Styling and Annotations

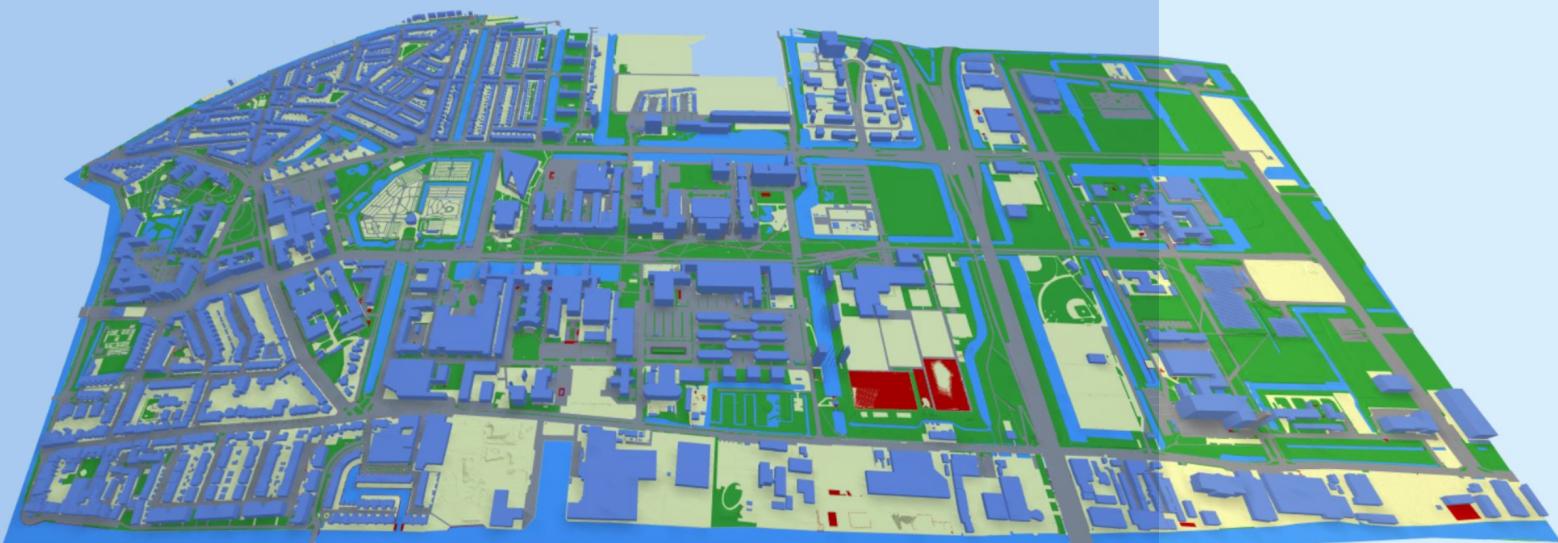
241 Styling
287 Annotations

[Go to First Change \(page 1\)](#)

MSc thesis in Geomatics for the Built Environment

CityJSON: does (file) size matter?

Jordi van Liempt
2020



MSc thesis in Geomatics

CityJSON: does (file) size matter?

Jordi van Liempt

June 2020

A thesis submitted to the Delft University of Technology in partial
fulfillment of the requirements for the degree of Master of Science
in Geomatics

Jordi van Liempt: *CityJSON: does (file) size matter?* (2020)

© This work is licensed under a Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/>.

The work in this thesis was carried out in the:



3D geoinformation group
Department of Urbanism
Faculty of the Built Environment & Architecture
Delft University of Technology

Supervisors: Hugo Ledoux

Balázs Dukai

Co-reader: Ken Arroyo Ohori

Abstract

The possibilities of 3D city models for the analysis of the built environment are increasingly explored, and there is a continuous development in improvements on their inner workings. They are also used more often on the web, for example for visualisation but it is also possible to query, analyse, and edit data in this way. In this case, the network becomes a potential new bottleneck in time performance. Especially when a 3D city model contains a lot of attributes, there is a rapid increase in file size when the area of study is expanded. This presents challenges in efficiency and in this thesis I focus on the improvement of the inner workings of 3D city models to attempt to relieve this problem, in specific for spreading and using them more efficiently on the web.

By investigating and testing different compression techniques on 3D city models stored in the CityJSON format, I have attempted to relieve this problem. CityJSON is already more compact than CityGML and these techniques decrease the file sizes of datasets even further, allowing for faster transmission over a network. But on the other hand, additional steps are needed to process compressed files. The goal of using a compression technique is to result in a net speed gain, meaning that the time that is saved on download time should be larger than the additional time that it costs to process the file before transmission (compressing the data on the server) and after receival (decompressing the data in the client).

There are compression techniques for both general and specific purposes, and I have used combinations of them. Amongst these are Draco compression, zlib, CBOR, and a self-created technique. Specific ones are used for the main characteristics that CityJSON datasets have, which are the JSON structure, feature attributes, and feature geometries. To assess the impact that different combinations of compression techniques have, I take uncompressed CityJSON performance as the baseline and have come up with different performance indicators that include several use cases such as visualisation, querying, analysis, and editing.

I have benchmarked all combinations of compression techniques on the use cases of these performance indicators. For this I have created two types of server implementations: one with which datasets are compressed beforehand and are processed in the client based on the request made by the user, and one where the data is processed first and only then compressed and transmitted to the server. In the results, you can see the best-performing compression type per use case.

The benchmarking is performed on a variety of datasets that are split into four categories: larger datasets, larger datasets without attributes, smaller datasets, and smaller datasets without attributes. This ultimately makes for use cases that are very specific and choosing suitable compression types requires finding out which ones perform relatively well in most cases, and are not difficult to implement in order to keep CityJSON a simple file format. It turns out that Draco compression can give good results in specific situations, but in general is not good to use. Not only regarding performance, but also from a developer point of view. CBOR, zlib, and a combination of these two are easy to use and generally affect the performance of CityJSON on the web in a good way.

Contents

1 Introduction	1
1.1 Motivation and problem statement	1
1.2 Research question	3
1.3 Scope	4
1.4 Thesis overview and preview of results	4
2 Theoretical background on compression	7
2.1 Data compression in general	7
2.1.1 Lossless versus lossy compression	7
2.2 Generic compression techniques	7
2.2.1 Binary file structure	8
2.2.2 Run-length encoding	8
2.2.3 LZ77	8
2.2.4 Huffman coding	9
2.2.5 DEFLATE and zlib	10
2.2.6 Smaz	10
2.3 Compression of geometries, meshes, and graphs	10
2.3.1 Quantisation	10
2.3.2 Delta encoding	10
2.3.3 Edgebreaker	11
2.3.4 Sequential connectivity compression	12
2.3.5 Parallelogram prediction	12
3 Related work	15
3.1 CityJSON	15
3.2 Draco	17
3.3 glTF	18
3.4 Cesium 3D Tiles	18
3.4.1 Batched 3D Model	19
3.5 Esri I3S	19
3.6 Binary JSON	20
3.6.1 CBOR	20
4 Methodology	23
4.1 Implementation of compression techniques	24
4.2 Performance of (compressed) datasets	26
4.3 Testing platform and benchmarking	27
4.4 Dataset preparation	29
4.5 Overview	29
5 Implementation and experiments	31
5.1 Flask server	31
5.1.1 Three.js	32
5.1.2 cgio	33
5.1.3 Draco	33
5.1.4 Testing platform overview	33

Contents

5.2 Operations	34
5.2.1 Visualisation	34
5.2.2 Querying	34
5.2.3 Spatial analysis	35
5.2.4 Editing (attributes)	35
5.2.5 Editing (geometry)	36
5.3 Testing and benchmarking with Selenium	36
5.4 Compression and decompression	37
5.4.1 original-zlib, original-cbor, and original-cbor-zlib	37
5.4.2 original-cbor-smaz	38
5.4.3 original-replace, original-replace-zlib, original-cbor-replace, and original-cbor-replace-zlib	38
5.4.4 original-cbor-replace-huff	39
5.4.5 Draco-compressed equivalents	39
5.5 Dataset description	40
5.6 Experiments	40
5.6.1 Data loss	41
5.6.2 Querying compressed CityJSON	41
5.6.3 Draco compression level	42
5.7 Computer specifications	44
6 Benchmarking results and analysis	45
6.1 Encoding performance	46
6.2 Decoding performance	48
6.3 File size	50
6.4 Visualisation	52
6.4.1 Compression in advance	52
6.4.2 Compression on the fly	54
6.5 Querying	56
6.5.1 Querying one feature	56
6.5.2 Querying all features	59
6.6 Spatial analysis	64
6.6.1 Buffering one feature	64
6.6.2 Buffering all features	68
6.7 Editing (attributes)	72
6.7.1 Editing attributes one feature	72
6.7.2 Editing attributes all features	75
6.8 Editing (geometry)	79
6.8.1 Editing geometry one feature	79
6.8.2 Editing geometry all features	84
6.9 Conclusion and recommendations	89
7 Conclusion	93
7.1 Performed work and summary of achieved results	93
7.2 CityJSON on the web and compression	94
7.3 Limitations	95
7.4 Future work	95
8 Streaming	97
8.1 OGC API	97
8.2 Streaming CityJSON	98
8.3 Streaming CityJSON and compression	98

A Datasets and characteristics	101
A.1 Delft	102
A.2 Den Haag	103
A.3 Montréal	105
A.4 New York	106
A.5 Rotterdam	108
A.6 Singapore	109
A.7 TU Delft campus	110
A.8 Zürich	112
B List of used tools	113
B.1 Shapely	113
B.2 cpio	113
B.3 Flask	113
B.4 three.js	113
B.5 harp.gl	113
B.6 Turf.js	113
B.7 Proj4js	113
B.8 Selenium	114
B.9 Draco	114
C Reproducibility self-assessment	115
D Full benchmark results	117

List of Figures

2.1 Visualisation of an example Huffman tree, created with the string "this is an example of a huffman tree" as input. Source: Dcoetzee [2007]	9
2.2 Visualisation of the creation of a triangle spanning tree, not showing the dual graph representation. Figure taken from Taubin et al. [1998]	10
2.3 The five cases defined for the Edgebreaker method. X is the current triangle in the iteration, v the vertex that is shared between X and the triangles to its left and right, and blue triangles are unvisited. Figure taken from Rossignac et al. [2003]	11
2.4 Illustrated how a new vertex can be predicted using parallelogram prediction. r^p is the vertex as predicted with the parallelogram rule, r_{cp}^p is as predicted with the crease angle taken into account. Figure adapted from Touma and Gotsman [1998]	13
3.1 The general structure of the Draco file format (.drc). Figure taken from Galligan [2019]	18
3.2 The conceptual structure of the Graphics Library Transmission Format (glTF) format (cropped from Khronos Group [2019])	18
4.1 The general research workflow	23
4.2 The workflow of the benchmarking. This is done for every compression method and all datasets.	28
5.1 Diagram showing an overview of the structure of the testing platform.	32
5.2 Benchmark of mean encoding times per different Draco compression level	42
5.3 Mean file size multipliers per Draco compression level	43
5.4 Benchmark of mean decode times per different Draco compression level	43
6.1 Boxplots of the encoding time performance of different compression types.	46
6.2 Boxplots of the decoding time performance of different compression types.	48
6.3 Boxplots of file size multipliers indicating to which extent compression types compress the data.	50
6.4 Boxplots of time performance of (in advance) compressed datasets on visualisation compared to uncompressed CityJSON.	52
6.5 Boxplots of time performance of (on the fly) compressed datasets on visualisation compared to uncompressed CityJSON.	54
6.6 Boxplots of time performance of (in advance) compressed datasets on querying one feature compared to uncompressed CityJSON.	56
6.7 Boxplots of time performance of (on the fly) compressed datasets on querying one feature compared to uncompressed CityJSON.	58
6.8 Boxplots of time performance of (in advance) compressed datasets on querying all features compared to uncompressed CityJSON.	60
6.9 Boxplots of time performance of (on the fly) compressed datasets on querying all features compared to uncompressed CityJSON.	62
6.10 Boxplots of time performance of (in advance) compressed datasets on buffering one feature compared to uncompressed CityJSON.	64
6.11 Boxplots of time performance of (on the fly) compressed datasets on buffering one feature compared to uncompressed CityJSON.	66
6.12 Boxplots of time performance of (in advance) compressed datasets on buffering all features compared to uncompressed CityJSON.	68

List of Figures

6.13 Boxplots of time performance of (on the fly) compressed datasets on buffering all features compared to uncompressed CityJSON.	70
6.14 Boxplots of time performance of (in advance) compressed datasets on editing attributes of one feature compared to uncompressed CityJSON.	72
6.15 Boxplots of time performance of (on the fly) compressed datasets on editing attributes of one feature compared to uncompressed CityJSON.	73
6.16 Boxplots of time performance of (in advance) compressed datasets on editing attributes of all features compared to uncompressed CityJSON.	75
6.17 Boxplots of time performance of (on the fly) compressed datasets on editing attributes of all features compared to uncompressed CityJSON.	77
6.18 Boxplots of time performance of (in advance) compressed datasets on editing geometry of one feature compared to uncompressed CityJSON.	79
6.19 Boxplots of time performance of (on the fly) compressed datasets on editing geometry of one feature compared to uncompressed CityJSON.	82
6.20 Boxplots of time performance of (in advance) compressed datasets on editing geometry of all features compared to uncompressed CityJSON.	84
6.21 Boxplots of time performance of (on the fly) compressed datasets on editing geometry of all features compared to uncompressed CityJSON.	87
6.22 Overview of the performance of all compression types by dataset type. The results are averaged over every operation.	89
6.23 Overview of the performance of all compression types by dataset type. The results are averaged over every operation.	90

List of Tables

1.1 The six performance indicators on which variants of compressed CityJSON are assessed, and the sections where they are addressed	4
2.1 Table that accompanies Figure 2.1. It contains all characters of the original string with their frequencies and the binary code with which they are mapped. Source: Dcoetzee [2007]	9
3.1 Types of city objects (features) that CityJSON natively supports (from Ledoux et al. [2019])	16
4.1 The six performance indicators on which variants of compressed CityJSON are assessed	24
4.2 Combinations of compression methods, separated by type of geometry.	25
4.3 Names for all operations in the results.	27
5.1 Chosen datasets and their characteristics, sorted by file size	40
5.2 Specifications of computer used for benchmarking	44
6.1 Median performance with visualise on large datasets, compression in advance	53
6.2 Median performance with visualise on large datasets without attributes, compression in advance	53
6.3 Median performance with visualise on small datasets, compression in advance	53
6.4 Median performance with visualise on small datasets without attributes, compression in advance	53
6.5 Median performance with visualise on large datasets, compression on the fly	55
6.6 Median performance with visualise on large datasets without attributes, compression on the fly	55
6.7 Median performance with visualise on small datasets, compression on the fly	55
6.8 Median performance with visualise on small datasets without attributes, compression on the fly	55
6.9 Median performance with queryone on large datasets, compression in advance	57
6.10 Median performance with queryone on large datasets without attributes, compression in advance	57
6.11 Median performance with queryone on small datasets, compression in advance	57
6.12 Median performance with queryone on small datasets without attributes, compression in advance	57
6.13 Median performance with queryone on large datasets, compression on the fly	59
6.14 Median performance with queryone on large datasets without attributes, compression on the fly	59
6.15 Median performance with queryone on small datasets, compression on the fly	59
6.16 Median performance with queryone on small datasets without attributes, compression on the fly	59
6.17 Median performance with queryall on large datasets, compression in advance	61
6.18 Median performance with queryall on large datasets without attributes, compression in advance	61
6.19 Median performance with queryall on small datasets, compression in advance	61
6.20 Median performance with queryall on small datasets without attributes, compression in advance	61

List of Tables

6.21 Median performance with queryall on large datasets, compression on the fly	63
6.22 Median performance with queryall on large datasets without attributes, compression on the fly	63
6.23 Median performance with queryall on small datasets, compression on the fly	63
6.24 Median performance with queryall on small datasets without attributes, compression on the fly	63
6.25 Median performance with bufferone on large datasets, compression in advance	65
6.26 Median performance with bufferone on large datasets without attributes, compression in advance	65
6.27 Median performance with bufferone on small datasets, compression in advance	65
6.28 Median performance with bufferone on small datasets without attributes, compression in advance	65
6.29 Median performance with bufferone on large datasets, compression on the fly	67
6.30 Median performance with bufferone on large datasets without attributes, compression on the fly	67
6.31 Median performance with bufferone on small datasets, compression on the fly	67
6.32 Median performance with bufferone on small datasets without attributes, compression on the fly	67
6.33 Median performance with bufferall on large datasets, compression in advance	69
6.34 Median performance with bufferall on large datasets without attributes, compression in advance	69
6.35 Median performance with bufferall on small datasets, compression in advance	69
6.36 Median performance with bufferall on small datasets without attributes, compression in advance	69
6.37 Median performance with bufferall on large datasets, compression on the fly	71
6.38 Median performance with bufferall on large datasets without attributes, compression on the fly	71
6.39 Median performance with bufferall on small datasets, compression on the fly	71
6.40 Median performance with bufferall on small datasets without attributes, compression on the fly	71
6.41 Median performance with editattrone on large datasets, compression in advance	73
6.42 Median performance with editattrone on small datasets, compression in advance	73
6.43 Median performance with editattrone on large datasets, compression on the fly	74
6.44 Median performance with editattrone on small datasets, compression on the fly	74
6.45 Median performance with editattrall on large datasets, compression in advance	76
6.46 Median performance with editattrall on small datasets, compression in advance	76
6.47 Median performance with editattrall on large datasets, compression on the fly	78
6.48 Median performance with editattrall on small datasets, compression on the fly	78
6.49 Median performance with editgeomone on large datasets, compression in advance	80
6.50 Median performance with editgeomone on large datasets without attributes, compression in advance	80
6.51 Median performance with editgeomone on small datasets, compression in advance	81
6.52 Median performance with editgeomone on small datasets without attributes, compression in advance	81
6.53 Median performance with editgeomone on large datasets, compression on the fly	83
6.54 Median performance with editgeomone on large datasets without attributes, compression on the fly	83
6.55 Median performance with editgeomone on small datasets, compression on the fly	83
6.56 Median performance with editgeomone on small datasets without attributes, compression on the fly	83
6.57 Median performance with editgeomall on large datasets, compression in advance	85
6.58 Median performance with editgeomall on large datasets without attributes, compression in advance	85

6.59 Median performance with editgeomall on small datasets, compression in advance	86
6.60 Median performance with editgeomall on small datasets without attributes, compression in advance	86
6.61 Median performance with editgeomall on large datasets, compression on the fly	88
6.62 Median performance with editgeomall on large datasets without attributes, compression on the fly	88
6.63 Median performance with editgeomall on small datasets, compression on the fly	88
6.64 Median performance with editgeomall on small datasets without attributes, compression on the fly	88
8.1 File size in MB for CityJSONCollections of datasets and compressed versions	99
D.1 Full benchmark results of compression in advance, datasets with attributes	121
D.2 Full benchmark results of compression in advance, datasets without attributes	125
D.3 Full benchmark results of compression on the fly, datasets with attributes	129
D.4 Full benchmark results of compression on the fly, datasets without attributes	133

Acronyms

GIS geographical information system	94
JSON JavaScript Object Notation	1
WebGL Web Graphics Library	113
blob Binary Large OBject	25
CBOR Concise Binary Object Representation	2
BSON Binary JSON	8
glTF Graphics Library Transmission Format	xi
OGC Open Geospatial Consortium	1
I3S Indexed 3D Scene layer	2
BSON Binary JSON	8
CBOR Concise Binary Object Representation	2
UBJSON Universal Binary JSON Specification	8
LZ77 Lempel-Ziv 1977	8
GML Geography Markup Language	1
b3dm Batched 3D Model	2
LoD Level of Detail	2
DBMS Database Management System	33
REST Representational state transfer	31
CRS coordinate references system	18
CQL Contextual Query Language	97
WFS Web Feature Service	1
WMS Web Map Service	1
CLI command line interface	113
HTML HyperText Markup Language	34

1 Introduction

1.1 Motivation and problem statement

The possibilities of 3D city models for the analysis of the built environment are increasingly explored, and there is a continuous development in improvements on their inner workings. A list of applications domains and use cases for which they are exclusively suitable, or more so than traditional 2(.5)D data, includes estimation of solar irradiation, visibility analysis, energy demand estimation Biljecki et al. [2015]. While many of these applications foremost need information on the geometries of the objects that are included in the study area, some require (detailed) semantics on their parts as well. Especially in the latter case this makes for a rapid increase in file size when the area of study is expanded. At the same time, there is an increase in popularity of the dissemination and direct use of geoinformation on the web [Dempsey, 2018; Mango, 2017]. This presents challenges in efficiency, as the network becomes a potential new bottleneck in time performance.

In this thesis I focus on the improvement of the inner workings of 3D city models to attempt to relieve this problem, in specific for spreading and using them more efficiently on the web. I do this by investigating and testing different compression techniques on 3D city models stored in the CityJSON [CityJSON, 2019] format. These techniques decrease the file sizes of datasets allowing for faster transmission over a network, but on the other hand add additional steps to process them. The goal of using a compression technique is to result in a net speed gain, meaning that the time that is saved on download time should be larger than the additional time that it costs to decompress the file after receival.

The creation of CityJSON in 2017 exemplifies the continuous developments on 3D city models, it being an encoding that is used to store the geometries and semantics of 3D cities and landscapes [Ledoux et al., 2019]. It follows the Open Geospatial Consortium (OGC) standardised CityGML data model but aims to be a better alternative to it by using the JavaScript Object Notation (JSON) encoding, which for example has the advantage of making it more suitable as an international 3D geoinformation standard for web use. The main reason for its creation is that Geography Markup Language (GML) in general is verbose and more complex to work with. CityJSON is more compact since JSON files are already smaller than GML, but also because it has better compression possibilities than CityGML (by storing geometries in the Wavefront OBJ [Reddy] style), resulting in on average 6 times smaller file sizes. In addition to that, JSON is smoother to use than GML files as it is a simpler file format, making it easier to read for humans and to be used by developers.

The compactness is an especially useful characteristic for web use since it decreases the time of data transfer over a network, and the use of compression techniques can decrease the sizes of datasets even further. Besides the potential improvement in data transmission efficiency that I examine in this thesis, CityJSON has not yet been fully tested in a web environment. Currently, 2D geoinformation is often disseminated through Web Feature Service (WFS) and Web Map Service (WMS) services, and the OGC recently has initiated the OGC API [Open Geospatial Consortium, 2019c] that specifies how online access should be provided to collections of geospatial data. Because it is comprised of multiple blocks and allows for extensions, its capabilities can flexibly be increased, giving way to the support for 3D geoinformation and performing (spatial) operations on it through the web [Open Geospatial Consortium, 2019c]. Efforts are already made to make CityJSON compliant with the OGC API specifications [Ledoux, 2020]. With this potential in mind, in this thesis I already consider

1 Introduction

the influence of the use of different compression techniques on the performance of a variety of data operations with CityJSON.

In order to further improve the efficiency of CityJSON, this thesis investigates compression techniques that can be implemented to decrease its average file size and potentially improve its web use speed. Compression techniques can be all-purpose, but can also be created for a specific type of data. An example of the former that I test is zlib, with which a complete CityJSON dataset can be compressed at once. The downside of using such a technique is that the file has to be decompressed completely before it is useable. This can be inefficient in cases where this is not needed, for example when querying only one feature. In addition to that, the decompression time can be relatively long as shown in Section 6.2. The decompression time is more important than compression time because the combination of a small file size and good decompression performance decreases the time it takes for a file to be useable after receipt, while compression is (in most cases) part of dataset preparation.

CityJSON has several characteristics for which specific compression techniques can be suitable. The main ones that are considered are the geometries of features, their attributes and semantics, and the JSON encoding itself. Using specific techniques on them can allow for partial decoding of compressed datasets and yield different performance results. With this it is important that no information is lost, as techniques can be lossy or lossless. The decompression speed of an algorithm is presumed to be more important than the compression speed, because decompressing is always necessary before a file is useable, while the compression is mostly part of data preparation.

A way that could be incorporated for the compression of the geometries is compression with the Draco library [Google, 2019b], which is an open-source library and can be used to compress and decompress 3D geometric meshes. The library makes use of several lossless compression techniques to encode the vertices and connectivity information of meshes, such as delta and entropy coding for the former and the Edgebreaker algorithm and parallelogram prediction for the latter (see Sections 2.3.3 and 2.3.4). Draco is specifically designed to improve the efficiency of their storage and transmission [Google, 2019b]. However, it assumes simple triangular geometries as input, while CityJSON can contain complex collections of volumes that potentially have holes and cavities. Therefore, it requires testing to find out if it is suitable to use.

In addition, the attributes and semantics can be compressed using techniques such as zlib, Huffman coding, and smaz (see Subsections 2.2.5, 2.2.4, 2.2.6). These reduce the redundancy in data in different ways, which I explain in Chapter 2. Lastly, the JSON encoding is human-readable which comes at the cost of a slower performance and higher file size than a binary format. There are binary formats such as Concise Binary Object Representation (CBOR) into which JSON can be encoded that approach the latter's key-value structure, which would yield performance benefits at the cost of being able to easily inspect a file without decoding it.

Different file formats already exist for the web use of 3D city models—Batched 3D Model (b3dm) and Indexed 3D Scene layer (i3S) [Analytical Graphics, Inc., 2019; Esri, 2019b], see Sections 3.4.1 and 3.5—and they focus on high performance, which this thesis aims at as well. But for performance they do not only regard the time that it simply takes to load a file but also incorporate other features that benefit user experience such as tiling, which enables only specific parts of the dataset to be loaded, or rendered in a different Level of Detail (LoD). In addition, they combine binary structures and JSON and allow for the use of Draco geometries. However, the focus on high performance makes them relatively complex to use for developers, as opposed to CityJSON that aims to be developer-friendly. While CityJSON or compressed variants might not be able to outperform these other file types, it is still suitable to use on the web when the best performance is not necessary, because of its ease of use. This does not mean that its performance should be neglected—compression can be of added value, but only if CityJSON's simplicity is mostly retained.

Besides that, the purpose of tiling is mostly to improve the streaming and visualisation of geospatial data [Analytical Graphics, Inc., 2018]. There are however more use cases for which a geoinformation

web service could be utilised, some of which could be easier to implement with CityJSON. Users might want to download a subset of the data per specified attributes or bounding box to avoid having to download unneeded information, editing functionalities can come in handy for crowd-sourced initiatives such as OpenStreetMap, and the ability to do spatial analysis can help it function as a cloud-based GIS application. Therefore the influence of using compression techniques is tested for several types of data operations.

Concretely, I test a variety of combinations of compression techniques on multiple CityJSON datasets, and assess them on their performance when executing different data operations with them. The operations are divided into five categories: visualisation, querying, spatial analysis, editing of attributes, and editing of geometries. Besides that the difference in file size is considered for lower storage space and the lossiness induced by the combination of compression techniques as all important information should be retained. The tested datasets have different characteristics since they can influence the workings of compression techniques, and I introduce a web application that is created for the testing of implementations. This testing platform is a Flask application that acts as a server to process requests, which in turn renders a web page that works with Three.js and Harp.gl to render datasets.

1.2 Research question

The main research question of the thesis is as follows: *To what extent can the implementation of different compression techniques improve CityJSON's performance on the web, considering file size, visualisation, querying, spatial analysis, and editing performance?*

To assess the impact that different compression techniques have, I take uncompressed CityJSON performance as the baseline and performance indicators are defined as derived from Section 1.1, shown in Table 1.1. I name it uncompressed for simplicity but all tested CityJSON datasets include the "transform" member, meaning that the origin of the set of vertices is moved to (0, 0, 0) and that their coordinates are represented as integers (see Section 3.1 for a more detailed explanation)—which is in fact a form of compression already. Some datasets originally had "transform" already and Draco requires vertices to be stored in this way to compress them, which is why I have chosen to make the datasets more uniform as to have one factor less to consider when interpreting the results.

For the visualisation, Three.js is used and the performance depends on the speed of the parsing of CityJSON geometries to a Three.js mesh and rendering it. As for querying and spatial analysis, the performance can be measured by the difference between the moment that a request has been made to perform the operation and the return of results. The main difference here is that the former needs object attributes to be decoded, while the latter requires their geometries. Since the decoding time and the ability to decode objects separately are characteristics that set compression methods apart, it does not matter which type of query or spatial analysis is tested. However, because of the latter characteristic it is fairer to test the performance of this on both a single and on all objects of the dataset. The editing is done in a similar way, with the difference being that the edited information needs to be compressed again, for which the execution time of the algorithm becomes relevant. For these four performance indicators the execution time of corresponding operations—which are defined in Chapter 5—is benchmarked.

Besides these, there are two characteristics that do not involve a time benchmark. The file size can be measured by the amount of bytes that a file takes up, with smaller files having the advantage that less storage space is needed and that they are transmitted over the web more quickly. The last point is the assessment of the lossiness of a compression type. No vital information from the original dataset should be lost—for instance loss of object ordering is acceptable while reduced coordinate accuracy and precision is not—as this would render the compression unusable for certain applications.

Compression type performance indicators	Section
Visualisation time	6.4
Querying time	6.5
Spatial analysis time	6.6
Editing time	6.7 & 6.8
File size compression	6.3
Lossiness	5.6.1

Table 1.1: The six performance indicators on which variants of compressed CityJSON are assessed, and the sections where they are addressed

1.3 Scope

The average execution time of the (de)compression algorithms are indirectly assessed with the indicators. Decompression is always done before a file can be used, and compression is done after editing a file. For this reason these are not seen as separate performance indicators. They are still benchmarked however because it can give more insight on what causes the differences in performance between compression types. Additionally a significantly longer execution time can make an implementation less desirable.

What I do not cover in the thesis is the compression of textures and materials. They are structured differently from the other parts of CityJSON and would therefore have to be investigated thoroughly to find suitable compression techniques, and most existing datasets (as retrieved from CityJSON [2020]) do not contain them. For this reason I consider it as future work.

The OGC API is mostly introduced (in Section 8.1) to give context about the current developments on the dissemination of geospatial information. It is not implemented in this thesis, but Section 5.6 shows an experiment on the compression of datasets that is useful when streaming datasets. Features can be compressed individually and streamed following the specification, so I did initial compression implementations in this way to see if it is worth it to investigate it further.

Tiling is explained as well (see Section 3.4). While being an important topic related to efficient web use of geoinformation, it is too complex to assess different ways to implement it into CityJSON within the time frame of the research and see if it is worthwhile to do it. However, because of the relevance, it is discussed whether or not it is viable to incorporate a tiling specification into CityJSON as an idea for future work, and how compression could be related to it.

Lastly, the "web use experience" part of the main research question does not encompass the assessment by humans. I am assuming that the six aforementioned objective main criteria of Table 1.1 can sufficiently represent the user experience, with the presumption that they prefer faster performance.

1.4 Thesis overview and preview of results

Chapter 2 discusses theory on compression and techniques that are either used in the thesis or concepts that are needed to explain other techniques. Subsequently Chapter 3 gives an explanation on CityJSON and existing file formats and concepts that are related to efficient web use of 3D models. Chapter 4 gives an overview of the methodology used to prepare CityJSON datasets, compress them, and assess their performance. Then, Chapter 5 explains with more details how the research is carried out by explaining variants of compressed CityJSON and the testing platform, and experiments that require further exploration, are done to make decisions in the implementation, or help

explaining the results. After that Chapter 6 shows the results of the six performance indicators for different compression implementations. As the last main part of the thesis, I discuss in Chapter 7 what has been achieved, what I think would be the best compression options to consider to actually implement, and how CityJSON could evolve with tiling and the OGC API. Lastly there is Chapter 8 on the streaming of CityJSON, which explains a different way of serving CityJSON data than is used in this thesis, and how compression can be useful for it.

The thesis shows results as good as 9% of the time to perform an operation with a compressed dataset as compared to doing the same thing with original CityJSON. That is the case when buffering all features of large datasets. However, there are also use cases for which it is not (very) beneficial to use compression, such as querying or buffering one feature. In general you see that operations on larger datasets benefit the most from compression. The influence of specific dataset characteristics however proves hard to analyse. In the end it is best to use compression techniques that are easy to implement and that work well all-round, rather than very good in specific use cases.

2 Theoretical background on compression

2.1 Data compression in general

Data compression has the goal to arrive at a compact representation of data by reducing its redundancy, while being suitable to use for specific needs [Sayood, 2017]. This latter part indicates that not always the most compact representation is the most optimal. There is always a trade-off between compactness and processing speed (compression and decompression), between which choices need to be made. In the case of this thesis, it is the most beneficial to have a balance between these such that the net speed gain is optimal, as explained in Section 1.1. The technique being lossless or lossy (is explained in Section 2.1.1) is of importance as well. With geoinformation in mind, a high decompression speed can give a better user experience as the rendering speed is increased, but loss of precision and accuracy that alters topological relationships is not acceptable.

All data compression techniques exploit either the structure or redundancy of the data or both. This means that techniques have to be designed in such a way to optimally make use of this, and prior knowledge on the data allows for the use of a more suitable technique (in speed, lossiness, or storage space) rather than using an all-purpose method [Sayood, 2017].

2.1.1 Lossless versus lossy compression

Lossless compression methods will always result in encodings that are able to be restored to their original form. This means that both the full information that it contained and its exact structure need to be retained for a method that is truly lossless. Text for example should virtually always be compressed losslessly, as small deviations can give outcomes that have a considerably different meaning [Sayood, 2017]. Examples of losslessly compressed file formats are RAR and ZIP, and run-length encoding, Huffman coding, and LZ77 are examples of lossless compression techniques which are mentioned in Section 2.2.

With lossy techniques however, data can potentially be compressed further to an even smaller representation. A method is labeled as such when information that is deemed as unimportant will be lost in the process, resulting in data that is different from the original as it is encoded and subsequently decoded. Examples are the removal of audio and visual frequencies that (most) humans are not able to perceive [Sayood, 2017]. Popular file formats such as MP3 and JPG are lossy, while examples of their lossless counterparts are FLAC and PNG. In the context of geoinformation, it may be acceptable for coordinates to be compressed in a lossy manner when they are of higher precision than necessary.

2.2 Generic compression techniques

In this Section, techniques that can be used or taken inspiration from for the compression of CityJSON are discussed, as well as basic methods that need to be covered since they are used to explain other techniques.

2.2.1 Binary file structure

A binary file can be defined as "any file that contains at least some data that consists of sequences of bits that do not represent plain text" [The Linux Information Project, 2006]. It usually is more compact than human-readable files, which is why it is suitable to use for compressed data. Additionally, by the addition of the type and length of a part of the data it can be read faster by machines [BSON, 2019].

There are several existing options to almost directly encode JSON into binary. The following ones are included in the JSON for Modern C++ library [Lohmann, 2020] and can therefore be considered as the most important ones: Binary JSON (bson) [BSON, 2019], CBOR [CBOR, 2019], MessagePack [MessagePack, 2019], and Universal Binary JSON Specification (ubjson) [UBJSON, 2020].

It is possible to write an original binary format specification to encode CityJSON, but to keep it close to the original format it is chosen to only investigate the aforementioned methods that can directly encode the format. In Section 3.6.1 CBOR is further explained, which is chosen as the one to test on CityJSON since it turned out to be the most efficient one regarding both file size and processing speed according to a benchmark of Zderadicka [2017].

2.2.2 Run-length encoding

Run-length encoding takes advantage of the occurrence of the same symbol in sequence. Following its basic principles, such a repeated character is encoded by the character itself, preceded by an integer indicating the number of times that it is repeated. In order to make the encoding more sophisticated, a negative integer can be used within the code to denote the amount of next characters that only occur once. This avoids having to always place a "1" in front of characters that are not repeated [Bourke, 1995].

As an example, the first string below would be encoded as the second string (adapted from Bourke [1995]).

Original string:

```
abcdfffffcbbbbabcdef
```

Encoded string:

```
-3 a b c 6 d -1 c 4 b -6 a b c d e f
```

2.2.3 LZ77

Similarly to run-length encoding, Lempel-Ziv 1977 (LZ77) exploits the sequential repetition of the data input. However, rather than working with single characters, it finds repeated blocks. While reading the data, the algorithm keeps record of a set amount of previous characters named sliding window. As the window slides, a new block identical to a block elsewhere in the window is replaced by two numbers that represent a distance (D) and a length (L). The former states the distance into the sliding window where the identical block starts, and the latter describes the length for which the block is the same [Feldspar, 1997].

An example taken from Feldspar [1997], below it is shown how a string is encoded in LZ77 with D being the amount of characters back where the repetition starts, and L being the number of characters from that part that is repeated.

Original string:



2.2.4 Huffman coding

Huffman coding is a general purpose and lossless method that attempts to find a minimum-redundancy code for a text that contains a minimised amount of symbols. As redundancy is eliminated, it results in a code that can be decoded back into the original text. This is done by first using the algorithm to produce a full binary tree (Huffman tree) with the to be coded text as input, which then outputs a tree containing all symbols that are present in the text based on the frequency in which they occur. Subsequently, the text is encoded based on this tree with more frequently occurring symbols being represented by less bits using binary code [Huffman, 1952].

On Figure 2.1 an example of a Huffman tree is shown, with the frequencies enclosed in boxes and the characters underneath it. It is accompanied by Table 2.1 in which all characters are shown with their frequencies and their binary encoding. Since the encoding of characters differs per Huffman tree, the tree needs to be available when data needs to be decoded. The tree uses up additional space and thus to be efficient, the encoded text should either be large or a general tree has to be created that is deemed suitable for to-be encoded strings.

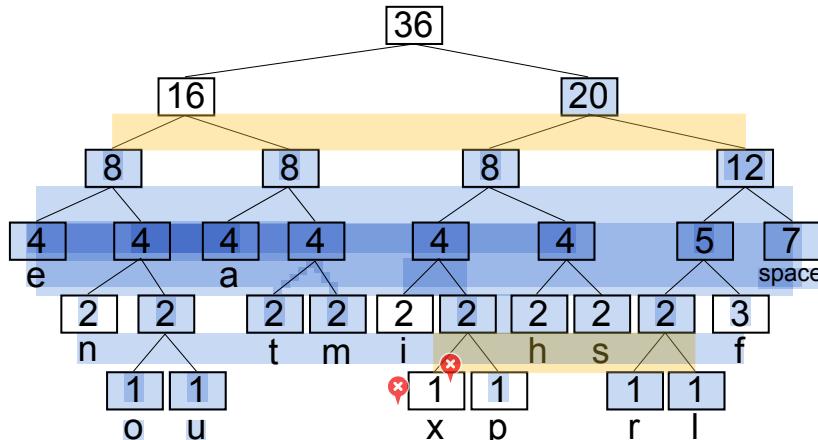


Figure 2.1: Visualisation of an example Huffman tree, created with the string "this is an example of a huffman tree" as input. Source: Dcoetzee [2007]

Char	Freq	Code	Char	Freq	Code	Char	Freq	Code
space	7	111	m	2	0111	p	1	10011
a	4	010	n	2	0010	r	1	11000
e	4	000	s	2	1011	u	1	00111
f	3	1101	t	2	0110	x	1	10010
h	2	1010	l	1	11001			
i	2	1000	o	1	00110			

Table 2.1: Table that accompanies Figure 2.1. It contains all characters of the original string with their frequencies and the binary code with which they are mapped. Source: Dcoetzee [2007].

2.2.5 DEFLATE and zlib

Zlib is a free general purpose compression library that uses the Deflate method, which in turn is based on Huffman coding and LZ77 compression. The input data is split into a series of blocks. Each block is compressed separately by using LZ77. Subsequently, Huffman coding is used on the encoded blocks. For the latter part, zlib has two different implementations that are performed on individual blocks: either with creation Huffman trees with the algorithm which is then stored alongside the data, or by using default trees that are defined in DEFLATE which removes the need of storing the extra information [Feldspar, 1997].

2.2.6 Smaz

Even though it can be used on any kind of natural language text, Smaz is a compression library for the specific purpose of working well on small strings. General purpose compression techniques tend to have a larger overhead in order to be able to work well with dynamic input, which is why Smaz is potentially more suitable for this specific case [Sanfilippo, 2019a]. It works with a static codebook containing frequent-occurring English (fragments of) words and bigrams. This makes it likely to work best on English strings, however, this is subject to testing. The parts are encoded in binary, with the most frequent parts being in the smallest representation [Sanfilippo, 2019b]. Its general idea is thus similar to Huffman coding (see 2.2.4).

2.3 Compression of geometries, meshes, and graphs

2.3.1 Quantisation

Quantisation is a technique that can be regarded as lossy compression. It is used for signal processing, converting continuous data from an analogue source to a digital representation. But it can also be applied to data that is already digitalised, for instance a list of coordinates as is exemplified by Draco in this thesis. It works by mapping a dictionary containing all values that the data encompasses to a dictionary reduced in size, to which original values are joined many-to-one. This means for example that a value that occurs once will be mapped to a value that is close to it but occurs more frequently, losing some of the original information [Sayood, 2017]. Vector quantisation can be used to compress coordinates [Rossignac, 2001].

2.3.2 Delta encoding

Another used name for this concept is delta compression. It is for instance used for version control, the distribution of software patches, and other transmission of data over the web. The idea behind it is to only store or transmit the delta (difference) of the data to data that is already stored or received. The complete data can subsequently be reconstructed by adding the delta to the existing data that it is connected to [Suel, 2019]. The technique can be applied within a complete set of data as well as is for example done in Deering [1995], where vertices are stored in an array as a vector of the delta difference with the previous vertex. Both transmitting deltas and storing deltas in a complete dataset are relevant for this thesis in the context of streaming and the compression of geometry as well as feature attributes.

2.3.3 Edgebreaker

Edgebreaker is one of the techniques used by Draco to compress 3D triangle meshes. A mesh that consists of triangles can be represented by the vertex data, which is made up of the coordinates of all vertices that the mesh encompasses, and the connectivity, which defines the relation between the triangles (or faces) and the vertices that belong to it. This is how Wavefront OBJ files are structured as well, which is followed by CityJSON. The Edgebreaker algorithm compresses the connectivity data of meshes [Rossignac et al., 2003].

Compression

To encode this information, first a spiraling triangle spanning tree is created for the mesh [Rossignac et al., 2003]. Analogised with peeling an orange, it is done by cutting around the peel in a spiraling way, after which it can be laid out flat. This is visualised in Figure 2.2. With a mesh this would result into a curled string of triangles. Ultimately, the dual graph of this string represents the spanning tree [Taubin et al., 1998].

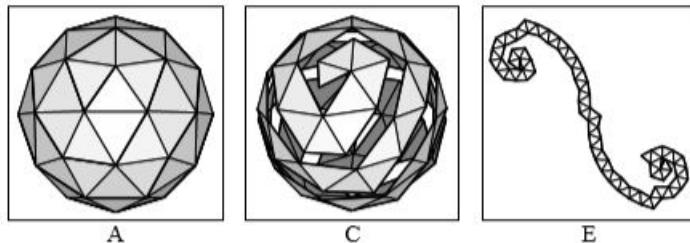


Figure 2.2: Visualisation of the creation of a triangle spanning tree, not showing the dual graph representation. Figure taken from Taubin et al. [1998]

Consequently, the algorithm traverses the triangles of the spanning tree, keeping track of all visited triangles and the vertices they are bounded by. Each triangle is assigned one character of the set {C, L, E, R, S}, which denotes one of the five cases (see Figure 2.3) that are defined in the specification of the algorithm. C means that the shared vertex v, between the current triangle and the ones to its left and right, has not been visited at all. As for the other four cases, vertex v has been visited and the triangle left or right (L or R), both of them (E) or neither of them (S) has been traversed [Rossignac et al., 2003; Rossignac, 2001].

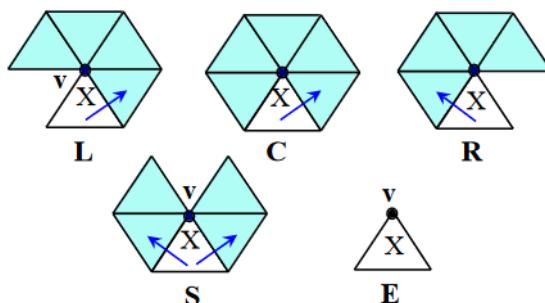


Figure 2.3: The five cases defined for the Edgebreaker method. X is the current triangle in the iteration, v the vertex that is shared between X and the triangles to its left and right, and blue triangles are unvisited. Figure taken from Rossignac et al. [2003]

2 Theoretical background on compression

During the traversal of the triangles two tables (M and U) are created to respectively store the visited vertices and triangles, and two others (V and O) are used to keep references to vertices and their opposite corners.

Optionally, a prediction method for CLERS symbols of succeeding unencoded triangles can be used, which is based on the vertex valences [Google, 2020c]. The valency (or degree) of a vertex is the amount of edges that is emanating from it. The addition of this method can decrease the file size in similar fashion to parallelogram prediction (see Section 2.3.5), as symbols would not have to be stored explicitly.

Decompression

Decompressing the mesh the V (vertice references) and O (opposite corners references) tables will be reconstituted, and after that the table containing the vertex positions. The CLERS string is iteratively decoded to create a triangulated polygon that approaches C in Figure 2.2, adding one triangle at a time. An S triangle will invoke the start of the decompression of one string of triangles (as there are multiple ones because of recursive compression). Encountering an E means that the end of the string is reached in the current recursion iteration. Depending on the other symbols that are encountered, the bounding edges shown in C in Figure 2.2 are zipped up by matching pairs of edges. In this way, ultimately shape A in Figure 2.2 is reconstituted. After the decompression of the connectivity information, the positions of the vertices are decompressed [Rossignac et al., 2003].

2.3.4 Sequential connectivity compression

An alternative method to edgebreaker that Draco can use is their own sequential connectivity method. It either uses delta encoding (storing the difference of vertex coordinates to the ones of the previous vertex [Deering, 1995] to store the IDs of vertices after which they are compressed by entropy encoding (of which Huffman coding is an example— 2.2.4), or the IDs are simply stored using everytime the amount of bits that the largest vertex id takes up [Google, 2019a]. In contrast with edgebreaker, this method preserves the vertex order within triangles. On the other hand the resulting encoding will not be as small.

2.3.5 Parallelogram prediction

Parallelogram prediction compresses meshes by predicting the missing vertex of a triangle, knowing the two vertices from the edge shared with a previously decoded triangle [Touma and Gotsman, 1998; Google, 2018]. A basic prediction can be made with the so-called parallelogram rule, as defined by formula 2.1, with the variables referring to the vertices in Figure 2.4.

$$r^p = v + w - u \quad (2.1)$$

However, this prediction can be improved by taking an estimated crease angle between the two triangles (the "crease" being the shared edge—the angles are depicted on Figure 2.4 with the lines crossing edges (s, v) and (v, u)) into account, as opposed to assuming that they are co-planar. This angle is estimated based on at least one set of adjacent triangles, that thus has to be decoded already. If there are more adjacent triangles available, it is taken as the average of the crease angles between the two sets of triangles that are closest to parallel to the shared edge of the current triangle and the predicted one [Touma and Gotsman, 1998]. As example, r_{cap} on Figure 2.4 is the predicted vertex of the next triangle by taking the crease angles between triangles (s, v, t) and (s, v, w) .

Ultimately, the mesh is compressed by storing the error of the predicted vertex to the actual vertex (as a vector) rather than the vertex itself. It is therefore a completely lossless technique.

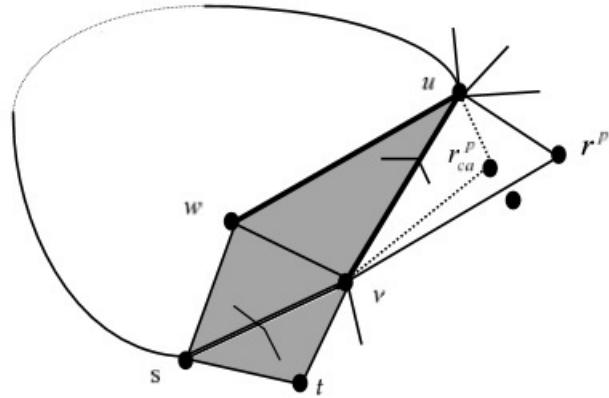


Figure 2.4: Illustrated how a new vertex can be predicted using parallelogram prediction. r^p is the vertex as predicted with the parallelogram rule, r_{ca}^p is as predicted with the crease angle taken into account. [Figure adapted](#) from Touma and Gotsman [1998]

3 Related work

3.1 CityJSON

CityJSON [CityJSON, 2019] is an encoding for storing digital twins, containing geometries and attributes of real world features such as buildings, terrains, roads, and waterbodies. It is based on the data model of CityGML [Open Geospatial Consortium], but uses JSON as encoding for it rather than GML which poses advantages in file compactness and ease of use. The file format has four main characteristics that can be exploited for compression purposes: the JSON encoding, geometries, attributes, and textures. As mentioned in Section 2.1, prior knowledge on the input data allows for more suitable compression techniques to be chosen. Furthermore, the explanation helps in understanding the difference with other file formats that are introduced in Sections 3.4.1 and 3.5.

The JSON format is easy for humans to read, while at the same time relatively lightweight (in comparison to for example GML) [Introducing JSON, 2019]. In its specification several data types are defined: boolean values, numbers, strings, arrays (ordered lists of elements), and objects (containing key-value pairs). These data types can be combined and nested [Ledoux et al., 2019]. In Listing 3.2 it is shown how this looks like with example snippet of a CityJSON dataset. However, the readability comes at the cost of verbosity. This is exemplified by the binary JSON-like format CBOR which is introduced in Section 3.6.1.

A CityJSON file is a JSON object contains 4 mandatory keys:

1. "type" (which is always of value "CityJSON")
2. "version"
3. "CityObjects" (a JSON object containing objects that represent geographical features)
4. "vertices" (an array containing the coordinates of all vertices)

A feature in CityJSON is called a city object and can be of a variety of types (shown in Table 3.1) as derived from CityGML's data model [Ledoux et al., 2019]. A city object is not limited to these types, as extensions can add others. A 2nd-level city object can be a child of a 1st-level city object parent, and thus linked together in this way.

It always has a "geometry" key which is an array containing more than 0 geometry objects, which can represent several LoDs. A geometry object in turn can be of several types of 3D geometric primitives. The vertex coordinates of a 3D model are stored in one array as shown in the above list, but the connectivity information is stored in geometry objects. Here, references to vertices are stored. This is inspired on the Wavefront OBJ [Reddy] format, and these parts are depicted in the green in listings 3.2, and 3.3. Additionally, a city object can have an "attributes" object as member. Here all attributes can be stored as defined in CityGML's data model, with which extension is possible as well [Ledoux et al., 2019; CityJSON, 2019].

In addition, there are optional keys for extensions, metadata, coordinate transform, appearance, and geometry-templates. Metadata can contain for example the geographical extent of the data, and coordinate transform gives the option to define a coordinate offset, which enables compression through the conversion of coordinates to integers and moving the origin to (0, 0, 0) [Ledoux et al., 2019; CityJSON, 2019]. This actually is a form of quantisation (Section 2.3.1). The basic structure of a CityJSON file is shown in listing 3.1, including "metadata" and "transform".

3 Related work

1st-level city objects	2nd-level city objects
Building	BuildingPart BuildingInstallation
Bridge	BridgePart BridgeInstallation BridgeConstructionElement
CityObjectGroup CityFurniture GenericCityObject LandUse PlantCover Railway Road SolitaryVegetationObject TINRelief TransportSquare	
Tunnel	TunnelPart TunnelInstallation
WaterBody	

Table 3.1: Types of city objects (features) that CityJSON natively supports (from Ledoux et al. [2019])

Lastly, it can contain materials, textures, and semantic surfaces belonging to city objects. These are respectively stored following the X3D specifications and with additional COLLADA files [Ledoux et al., 2019; CityJSON, 2019]. This is however out of the thesis scope (see 1.3). 

Listing 3.2: Snippet of basic structure of CityObjects. Attributes in orange, geometries in green

Listing 3.1: Snippet of CityJSON structure with mandatory members in orange

```
{
  "type": "CityJSON",
  "version": "1.0",
  "CityObjects": {...},
  "vertices": {...},
  "metadata": {
    "geographicalExtent": [
      84616.468,
      447422.999,
      -0.47,
      85140.839,
      447750.636,
      13.8
    ],
    "referenceSystem": "urn:ogc:def:crs:EPSG::7415"
  },
  "transform": {
    "scale": [
      0.001,
      0.001,
      0.001
    ],
    "translate": [
      84616.468,
      447422.999,
      -0.47
    ]
  }
}
```

```
{
  "type": "CityJSON",
  "version": "1.0",
  "CityObjects": {
    "b0a8da4cc-2d2a-11e6-9a38": {
      "type": "BuildingPart",
      "attributes": {
        "creationdate": "2014-07-09",
        "terminationdate": "",
        "function": "garage"
      },
      "geometry": [
        {
          "type": "Solid",
          "boundaries": [
            [
              [
                [
                  [0, 1, 2],
                  [1, 2, 3],
                  [1, 3, 4],
                  [0, 1, 4]
                ],
                ...
              ],
              "lod": 2,
              "parents": ["b0a8da4cc-2d2a-11e6-9a39"]
            ],
            ...
          ],
          "vertices": [...]
        }
      ]
    }
  }
},
```

Listing 3.3: Snippet of vertex array of CityJSON, highlighted in green

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {...},
  "CityObjects": {...},
  "vertices": [
    [
      85012.343,
      447455.577,
      0.27
    ],
    [
      85010.804,
      447448.808,
      -0.27
    ],
    [
      85013.832,
      447447.447,
      0.27
    ],
    ...
  ]
}
```

3.2 Draco

Draco, an open-source library by Google for compression and decompression of geometric meshes [Google, 2019b], can be used to compress the geometrical part of the CityJSON structure. Generally, it works with the Edgebreaker algorithm (which is explained in Section 2.3.3), their own sequential connectivity method (Subsection 2.3.4), parallelogram prediction (Subsection 2.3.5), and the quantisation of coordinates (Subsection 2.3.1). It can compress both 3D geometrical meshes and point clouds, but only the former is relevant for this application. On average, it compresses meshes (that are in Wavefront OBJ format) to about 2 percent of their original file size [Google, 2019c]. It has encoders and decoders written in C++ and JavaScript (and exclusively a decoder in WebAssembly), enabling it to be suitable for both offline (in this case preparation of compressed CityJSON files) and online (decoding the compressed geometries in the browser) purposes.

Not just the savings on file storage itself are beneficial, but since files are transferred over the web for browser use, time is saved as the data is in a slimmer format. The network speed can form a bottleneck, and it is exemplified by Cesium that using Draco can potentially improve the rendering speed as it is decoded asynchronously, meaning that parts are already rendered while the complete dataset is still being downloaded [Cesium, 2018].

There are three main configurations that can be set when encoding a mesh. First, there is the compression level parameter which goes from 0 to 10. Based on this parameter, different compression



features that the library encompasses are used, with 10 giving the most compressed result but at the same time having the worst decompression time. Second, a parameter that specifies the amount of bits to which the input values should be quantised can be chosen (which would make the compression lossy if used—see Section 2.3.1), and lastly, there is the option to retain the separation between objects Google [2019c].

Draco compressed meshes are stored in their native .drc binary format. The format generally consists of four parts: a header, metadata (object id for example), connectivity data, and attributes, as shown on Figure 3.1.



Figure 3.1: The general structure of the Draco file format (.drc). Figure taken from Galligan [2019]

3.3 glTF

glTF is a file format that is designed for efficient loading of 3D graphics and their dissemination over the internet [Khronos Group, 2019]. It is formatted in JSON and keeps file sizes small, can be processed fastly and is an extensible format. The elements that it contains conceptually have a hierarchical relationship as shown in Figure 3.2. Roughly explained, a scene is one model, and a file can contain multiple ones. A scene in turn contains nodes (which are features—for example a building). These do not have a specified coordinate references system (CRS), but can have their own translation, rotation, and scale properties, similar to CityJSON's "transform" (see Section 3.1). Their geometries are stored in a combination of meshes, buffers, bufferViews, and accessors, and it is possible to include skins, animations, cameras, materials, and textures. Moreover, it is possible to include for example buffer geometries and textures in binary format.

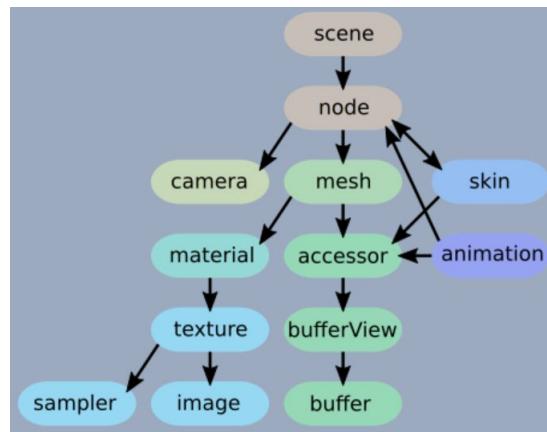


Figure 3.2: The conceptual structure of the glTF format (cropped from Khronos Group [2019])

3.4 Cesium 3D Tiles

Despite this part being out of scope for this thesis (as stated in the introduction), it is still important to notice its existence as it can make a contribution to the efficient streaming of 3D geoinformation.

Besides that, it is implemented by the b3dm format (see Section 3.4.1), which is an alternative to CityJSON for web visualisation.

Created for the sharing, visualisation, and analysis of 3D geoinformation, the 3D Tiles specification [Analytical Graphics, Inc., 2018] is a JSON object in which the data is separated into tiles. Every tile is described by a bounding volume and the tiles are organised into a tree, which could be a quadtree, octree, k-d tree, or grid. Tiles can have children, creating a hierarchical structure with which different LoD can be represented. The tile parent would contain the 3D model in a low LoD, and every child would contain either more details that can be rendered when needed, or again a full 3D model but with a higher LoD [Cesium, 2020].

Its implementation can thus benefit the streaming of massive datasets as it allows the loading of smaller parts (tiles) as they are needed, e.g. when they are zoomed into, and in the LoD that they are wanted. In Section 3.4.1 an implementation of 3D Tiles is explained, giving more details about the structure.

3.4.1 Batched 3D Model

b3dm is a binary file format that can store heterogeneous 3D models and incorporates the 3D Tiles specification [Analytical Graphics, Inc., 2019]. It roughly consists of three parts: a feature table, batch table, and glTF geometry. In the specification a feature is a 3D model and the feature table can thus be seen as a collection of tiles. Every feature may have a position property indicating the coordinates of its centre.

In the batch table the properties of the contents of features are contained. These properties are stored in JSON format, but in a different fashion than CityJSON's CityObjects. It is an array value with the name of the property as key, with the array containing values for the ordered objects of the feature. This means that if there is an n amount of objects, the array for a property will contain n elements. An example is as follows:

```
"height" : [10.0, 20.0, 15.0]
```

Alternatively, the properties can be stored in a binary body to which a reference is placed within the aforementioned JSON structure. Since this way of storing attributes is inefficient when there are properties that only apply to a subset of objects (for example "amount of leafs" for trees, which buildings would not have), it is possible to define object types which have their own set of property names.

Lastly, the geometries of the objects from all features are stored as glTF, which allows for the use of Draco compression as well since it can be embedded therein.

The format therefore is similar to the aimed end-result of this thesis, with the additional benefit of having the option to structure the data into tiles. On the contrary, it lacks the benefits of the CityGML data structure, the tools that exist for CityJSON, and (further) compression of attributes [Analytical Graphics, Inc., 2019].

3.5 Esri I3S

I3S [Esri, 2019b] is an open file format by Esri. The scene layer that it encodes in turn contains large volume heterogeneous geographic data (e.g. 3D objects, point clouds, and buildings). It is formatted in JSON and specifically designed for the efficient streaming of massive 3D datasets. Similarly to CityJSON, Scene Layers encompass the geometry, texture, and attributes of features. Features are grouped into nodes with the nodes forming a regular or irregular tree by their bounding sphere (or

3 Related work

oriented bounding box), akin to respectively a quadtree or an R-tree. Parent nodes in a tree contain the features with a lower LoD. Because of this structure, specific nodes can be found faster and loaded more efficiently as well as simpler representations of geometries can be shown when they are looked at from afar [Esri, 2019a,b].

A node in itself does not contain complete features, but rather IDs which point to geometry buffers (encoded with Draco), textures, and attributes. The standard allows any geodetic coordinate reference system to be used, with the node having the same CRS as the vertices that it encompasses. The vertex position are stored with the centre of the node's bounding area as offset, allowing for compression and easier visualisation in for example three.js. LoD are not the same in I3S as in CityJSON—it uses thinning, clustering, and generalisation of 3D objects or meshes, as well as the downsampling of textures. The to-be visualised LoD (thus which leaf of the tree is chosen by the viewer) is determined by screen size, resolution, bandwidth, and memory [Esri, 2019b,a].

3.6 Binary JSON

Binary file formats can be more concise and processed (reading or writing) faster, as explained in Section 2.2.1. It is therefore relevant for the improvement of CityJSON's efficiency. Additionally, it can be beneficial to store binary data in JSON format, such as encryption keys or graphics like Draco geometries (see 3.2). JSON natively does not allow for this and requires such data to be encoded in Base64 format. This is a family of binary-to-text encoding schemes that is utilised when binary data needs to be stored, while working with media that are created to handle ASCII rather than binary data [CBOR, 2019]. The downside of this is that a string in Base64 encoding can be 33% larger than its binary form [Mozilla, 2020].

There are several different binary formats that encode data in a JSON-like manner. In a JSON library for C++ by Lohmann [2020], the following ones are implemented: BSON, MessagePack, UBJSON, and CBOR. I have chosen to use CBOR for the reasons below, as derived from a comparison with CBOR from CBOR [2013], but this does not necessarily mean it is the best-performing option for the use cases of this thesis. BSON is specifically developed for MongoDB [MongoDB, 2020]. This means that it is adapted for that purpose and this could make it less efficient for other purposes. It does have the possibility for updating values, but this comes at the cost of less concise storage. MessagePack is similar to CBOR, but has caveats that have to do with backwards compatibility and there is not much space for extension. UBJSON exactly mimics the JSON data model, which means that it can not store binary values. However, this functionality is necessary for storing Draco geometries within a CityJSON-like (compressed) file.

3.6.1 CBOR

CBOR has two clear purposes that were stated prior to its development, which are storing binary strings into a JSON-like format and compression of data [CBOR, 2019]. Like JSON, it follows a key-value structure. Roughly explained, every data item starts with an initial byte that contains information about the major type of the item and supplementary information. The major type can be one of the following [CBOR, 2013]:

- Unsigned integer
- Negative integer
- Byte string
- Text string
- Array

- Map (JSON-style object)
- Optional semantic tagging of other types (such as date/time string, URI, base64)
- Floating-point number

When reading the data, knowing the type and length of a value increases the performance as mentioned in [Section 2.2.1](#).

4 Methodology

The network speed is always a bottleneck when transmitting complete CityJSON datasets over the web and doing something with the data, as the receiver can only do something with it once the data has arrived. The assumption we thus make is that compressed datasets will have a better performance when the time that is gained in network transmission speed is larger than the time that is lost decoding the data after arrival. This Chapter lays out the plan to benchmark the performance of several types of compressed CityJSON compared to uncompressed CityJSON, to find the advantages and disadvantages of techniques and give recommendations on what the best compression types would be given specific scenarios.

Shown in Figure 4.1 is the structure of the methodology, divided into four general parts:

- The implementation of compression techniques for CityJSON
- The preparation of datasets to test compression techniques with
- Preparation of a testing platform that can work with both uncompressed and compressed CityJSON files
- Testing and benchmarking the performance of different file types on the performance indicators of Table 4.1, which were introduced in Section 1.2

The first three are finished prior to the benchmarking, in order to avoid having to redo work for part 2 when something had been changed in part 1. The tasks in part 1 are grouped together because there is no necessary order and are implemented simultaneously. All four are explained further in the next sections.

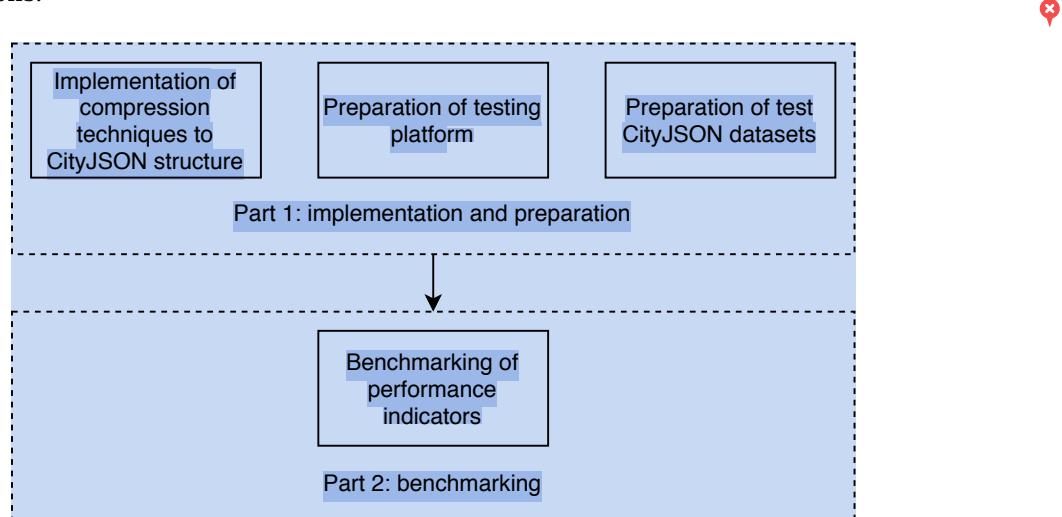


Figure 4.1: The general research workflow

Compression type performance indicators
Visualisation time
Querying time
Spatial analysis time
Editing time
File size compression
Lossiness

Table 4.1: The six performance indicators on which variants of compressed CityJSON are assessed

4.1 Implementation of compression techniques

In Chapter 2 a multitude of compression techniques has been introduced. Different combinations of these techniques are applied to the CityJSON structure and be compared on the six defined performance indicators (see Table 4.1), with the performance of uncompressed CityJSON as the baseline. A possible outcome is that certain combinations of techniques are more suitable for specific purposes, making it difficult to choose one that is the best. Chapter 6 shows which compression types For this reason, there is a depthly discussion about it in Chapter 6.

Table 4.2 shows which combinations of compression techniques I implement on CityJSON and benchmark on time performance. The table has a divide between compression methods where the geometries are kept as is, and where they have been replaced by Draco-compressed geometry. All other names besides "original" and "draco" indicate either the inclusion of a binary format or a compression technique.

The only compression technique that has not been explained in Chapter 2 is the so-called replace technique, which I have created. It means that all keys and values of "CityObjects" from the json structure are copied to a separate array. The elements of the array are sorted by frequency, and the keys and values in the json are replaced with a reference to the array. This has three supposed advantages:

1. The json structure is retained, of which parts could be reconstituted as wanted. It would for instance be possible to only decode the first City Object rather than all of them. When using a compression technique directly on the json structure, all data has to be decoded prior to use.
2. The array allows for greater compression as the keys and values can be compressed at once with for example zlib, as opposed to compressing all of them individually.
3. Redundancy is removed, as keys and values that appear more than once are only stored as one array element.

The array can be stored as a member of the CityJSON object and thus optionally be included in a CBOR representation. An example of how a dataset compressed in this way looks like is found in Section 5.4.3.

Ultimately there are 20 different types of CityJSON tested, of which 19 are compressed variants and 1 is original CityJSON. However, "original CityJSON" is defined as including the "transform" member (as stated in Section 1.2, with "transform" being explained in Section 3.1), which is done for the uniformity of datasets and because Draco needs coordinates to be stored in this manner before compression. It is however already a form of compression, but included in the CityJSON specification [CityJSON, 2019].

They tested CityJSON types are divided into two categories: having geometry compressed by Draco, or geometry as originally implemented in CityJSON. So there are actually 10 pairs, within which the

	Original geometry	Draco geometry
1	original	draco
2	original-zlib	draco-zlib
3	original-cbor	draco-cbor
4	original-cbor-zlib	draco-cbor-zlib
5	original-cbor-smaz	draco-cbor-smaz
6	original-replace	draco-replace
7	original-replace-zlib	draco-replace-zlib
8	original-cbor-replace	draco-cbor-replace
9	original-cbor-replace-zlib	draco-cbor-replace-zlib
10	original-cbor-replace-huffman	draco-cbor-replace-huffman

Table 4.2: Combinations of compression methods, separated by type of geometry.

difference is the use of Draco. Table 4.2 shows all CityJSON implementations, and in the following enumeration a brief description is given for all 10 combinations:

1. The dataset is in the style of regular CityJSON, but needs to include "transform". "original" is the baseline to which all other combinations are compared. As for "draco", the Draco Binary Large OBject (blob) is placed underneath the JSON structure, with a delimiter inbetween. It can not be added as a member of the CityJSON object as JSON does not allow binary values. This means that the result is a binary file.
2. The same as 1, but with the complete file compressed with zlib (see Section 2.2.5).
3. Similar to 1, but with the CityJSON object parsed into CBOR (see Section 3.6.1). In this case, the Draco blob is added as a member of the CityJSON object as it does allow for binary values. This would make it easier and quicker to read as well, since the file does not need to be split by a delimiter.
4. Encoded in CBOR like 3, but all keys and string values of the CityJSON object are compressed by smaz. This only works in combination with cbor, because strings are binary-encoded by smaz.
5. Same as 3, but with the resulting CBOR (Section 3.6.1) compressed with zlib (Subsection 2.2.5).
6. Similar to 1, but with all keys and values copied to an array and replaced by references.
7. Same as 6, but with the array compressed with zlib.
8. Same as 6, but with the CityJSON object parsed into CBOR.
9. Same as 8, but with the array compressed with zlib.
10. Same as 8, but with the array compressed with Huffman coding (see Section 2.2.4).

Decompression functions have to be written as well, both to be able to assess the loss of information due to the compression process and to allow for the data to be decompressed for use in the web application. Loss that is expected to see mainly concerns the precision (or even accuracy) of coordinates, the ordering of vertices, and the ordering of City Objects. The former is a large problem as the thesis concerns geographical data that can be used for analysis, as opposed to visualisation only. It can happen when Draco is not used in the correct way (see Section 3.2). The second one can happen with Draco as well, but since its geometries are triangulated it only means that vertices of triangles are ordered differently which is not problematic. Lastly, a different ordering of City Objects is acceptable because all information is retained, but it should still be noted as it is possible that a dataset is ordered with a purpose.

4 Methodology

Encoding JSON into CBOR is likely to give a better performance than other compression methods (such as directly compressing the datasets with zlib), as explained by UBJSON [2020]. While zlib should decrease the file sizes further than CBOR, the processing—encoding and decoding—of zlib-compressed data is less efficient. In addition, a binary JSON-like format enables a person to still inspect human readable parts (such as strings), which can help with debugging. Another option however is to zlib a CBOR file, which can be done if storage space is an important issue.

A possibility is to create an original binary format for CityJSON. However, since JSON can already be directly encoded into CBOR, this is not necessary as it would overcomplicate it. It could be done in a similar way as b3dm (see Section 3.4.1) by retaining the JSON structure and having a separate part for binary values. However, in this way only values are stored in binary and not keys, and when creating a dataset it needs to be chosen which values are stored in binary.

4.2 Performance of (compressed) datasets

The performance of the different compression method combinations are assessed following the six main performance indicators from Table 4.1. For each of these, specific data operations are defined that are tested, with original CityJSON datasets as the baseline to which the performance of compressed CityJSON files is compared. Below follows an explanation on all operations, and an overview can be seen on Figure 4.2.

The file size performance is defined as the compression factor. This is the size of a compressed dataset divided by the size of the uncompressed one. For all other operations, the performance is defined as the time that it takes to complete the operation with the compressed dataset divided by the time that it takes to finish the same operation with the original dataset.

1. Querying. Querying one City Object of a dataset (by its ID) and querying all City Objects (by an attribute field that all City Objects have in common). The query returns the matching City Objects including the vertices that their Geometry Objects refer to.
2. Visualisation. The rendering of all City Objects of a dataset.
3. Spatial analysis. Buffering one City Object and buffering all City Objects by 1 metre in 2D, so their geometries are first projected to 2D.
4. Editing of attributes. Adding "1" at the end of the ID of one City Object, and doing the same thing for all City Objects. If the file had to be decompressed before the operation, it has to be compressed again.
5. Editing of geometry. Raising the height of all vertices of one City Object by 1 metre, and doing the same thing for all objects. If the file had to be decompressed before the operation, it has to be compressed again.

These are performed with all datasets, which are compressed as described in Section 4.1. All combinations of operations, datasets, and compressions are thus assessed. The assessment is done by benchmarking the time that is needed to complete an operation, taken as the average over ten attempts to eliminate abnormalities.

Being secondary indicators, lossiness and possibility of asynchronous loading are described separately. The former can involve loss of object order, loss of vertex order, and triangulation of a file that was originally non-triangulated

Ultimately, by putting all results in a table, the best (combination of) methods can be found for every separate criterion. On each of the indicators, the compression method is evaluated and ranked. An overview of the benchmarking workflow can be found on Figure 4.2.

4.3 Testing platform and benchmarking

To test the performance of using CityJSON and its different compressed variants on the web, a web application has to be set up that enables this. While two options already exist (Boersma [2019] and CityJSON [2020]), it is better to create one that is of a simple structure so that it can easily be adapted for experimental functionalities for the research.

Therefore, I have set up a Flask application which acts like a server, which integrates functionalities of Draco [Google, 2019c], cgio [3D geoinformation research group at TU Delft, 2019], self-created functions, and harp.gl [HERE, 2019] through JavaScript. An explanation of the mentioned tools is given in [Chapter B](#). The platform enables to do the operations of [Figure 4.2](#) on all (compressed) CityJSON variants by visiting corresponding Flask routes in a browser. The operations have short names as shown in [Table 4.3](#) to easily identify them in the results.

Every dataset (8 of them, having different characteristics—see [Section 4.4](#)) is compressed by every CityJSON implementation type (20) and then benchmarked on the time that it takes to complete every individual operation (9) on them. This means that there are 1440 planned resulting time benchmarks. Every time benchmark is performed multiple times in order to retrieve more accurate results, as there are some uncertainties in the timing meaning that every iteration will have a slightly different time value. The time values are in milliseconds and outliers within the 10 iterations are removed based on having a z-score of 2. The z-score is the amount of standard deviations that a measurement differs from the average of the set.

The benchmarks are performed through a web testing library that visits the specific Flask route for a combination of dataset, operation, and compression type in a browser with a throttled network speed to simulate practical conditions.

Visualisation	visualise
Querying	queryone queryall
Spatial analysis	bufferone bufferall
Editing (attributes)	editattrone editattrall
Editing (geom)	editgeomone editgeomall

Table 4.3: Names for all operations in the results.

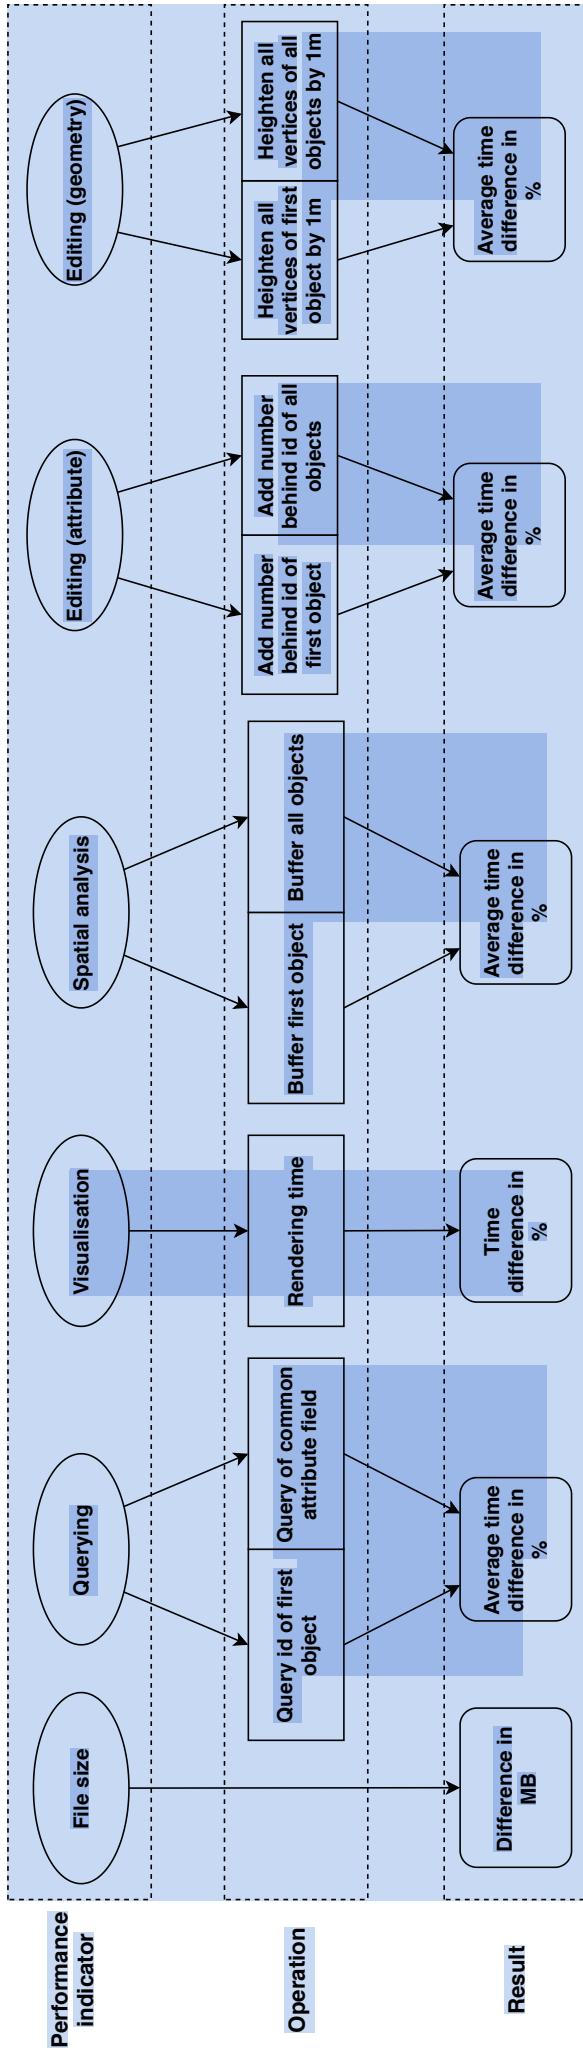


Figure 4.2: The workflow of the benchmarking. This is done for every compression method and all datasets.

4.4 Dataset preparation

Since different dataset characteristics can influence the compression factor and benchmarking performance of compression types, it would be good to use a varied selection of CityJSON datasets to test with. For example, a compression method that makes use of the repetition of data will work better when there are more repeated elements in the dataset that can be taken advantage of. The file size of the to be compressed dataset can be an indication of this, as larger files tend to have more repeated parts. For example the keys "type" and "boundaries" are there for every City Object, and possibly repeated attributes. For this reason it is expected that larger datasets will generally have a lower file size multiplier.

The LOD is interesting because it indicates a difference in the geometrical complexity of the features. It is possible that there is a difference to be seen in the impact of Draco compression depending on it. The variety of feature types is another characteristic because similarly to LOD, terrains can potentially have more complex geometries than buildings. As for attributes, the amount of them per feature could impact the performance of attribute compression, as it would be expected that more attributes are percentually compressed further. Lastly, the language they are in can potentially make a difference when the smaz compression technique is used.

The file size will however be the most important about datasets, since the bottleneck is assumed to be the network speed when using large datasets on the web. If a dataset is already small, for example 1 MB, compressing it to 0.7 MB would only have the file be downloaded 0.06 seconds faster on a 5 MB/s (or 40 Mbit/s) internet connection (which is the average in The Netherlands). Compressing a 100 MB file to 70 MB on the other hand will improve the transmission speed by 6 seconds.

An overview of important dataset characteristics can be seen below, and the chosen datasets can be found in Section 5.5.

- File size
- Repetition of data
- LOD
- Variety of feature types
- Size of attributes
- Attribute language

4.5 Overview

Concludingly, the following results are presented in Chapter 6:

1. The resulting compression factors of individual datasets as compressed by every compression type, and the average compression factor per compression type.
2. The encoding time of individual datasets with every compression type, and the average encoding time per compression type.
3. Time performance benchmarks on all compression type, dataset, and operation combinations with both server implementations. As conclusion, plots are made showing the average performance of a compression type per individual dataset, meaning that the operations are averaged. With this you can answer the question: "I have this dataset and want to do anything with it, which compression types are the most suitable?". Taking averages over all datasets is not done since they are not a representative selection, as opposed to the operations.

5 Implementation and experiments

This Chapter firstly describes how the testing platform is set up to assess the performance of different types of compressed CityJSON. Then, I explain the way in which the operations defined in Section 4.3 are performed on the testing platform, and subsequently give information on the performance benchmark,. After that comes the implementation of the compression types of Section 4.1 and a description of the datasets on which these are tested. Lastly I present the results of experiments that do not directly belong in Chapter 6 but analyse the lossiness of compression types, helped with making implementation choices, or are relevant in other ways. Tools that are used for the implementation are briefly described in Appendix B.

5.1 Flask server

The Flask server is the base of the testing platform introduced in Chapter 4. Flask is a lightweight Python web framework with which web applications can be built [Flask, 2020]. It enables combined use of Python and JavaScript, where the former handles the transmission of data (from datasets that are stored on the server) to the client, while the latter handles the processes that are performed in the client, such as the visualisation. In this way, CityJSON files (and compressed variants) can be used on the web, with operations performed on it on either the server or in the client.

The testing platform is set up in Representational state transfer (REST)ful [Fielding, 2000] style—with a clear separation between the client and the server, and having stateless communication between these two, which means that every request to the server needs to contain all necessary information to fulfill the request as the server does not remember previous interactions with the client. In Flask, routes (in the form of URLs) are defined in the server's code and coupled with a function that is executed when the specific route is visited, i.e. the client makes a request to the server. For the testing platform I have created routes and corresponding functions that enable to perform the operations laid out in Table 4.3. When such a function is finished, it returns an html template that is rendered in the client, a dataset (that is prepared on the server if applicable), and variables with information on the current operation so that the client knows how to process the data.

There are two different server implementations: one with the compression of datasets beforehand, and one with compression on the fly. As for the first one, all datasets are compressed by every compression types and these files can be requested from the server. Compressed files can not be manipulated before decompressing them (however, see Section 5.6.2). Since the presumed advantage of using compressed files on the web is the mitigation of the data transfer time from server to client, decompressing and performing an operation on a dataset on the server is unwanted because it would diminish this potential speed gain. Therefore, the server transmits the full compressed dataset to the client, which will proceed to decompress it and execute the specified operation. For a fair comparison, operations on original CityJSON data are also performed in the client, despite the possibility to process them before transmission.

With the second implementation, an original CityJSON dataset is always loaded by the server after a request has been made. The specified operation is then performed on it (except for visualisation, which can only be done in the client), and if necessary the processed data is then compressed before transmission. This means that it is possible to only transmit a selection of city objects rather than

5 Implementation and experiments

having to send them all. But because now the compression time is of importance for the time that is lapsed finishing an operation, this could hamper the performance of compression types, especially when they take long.

In Figure 5.1, an overview of the server implementation used in this thesis is shown. It is further explained in Section 5.1.4.

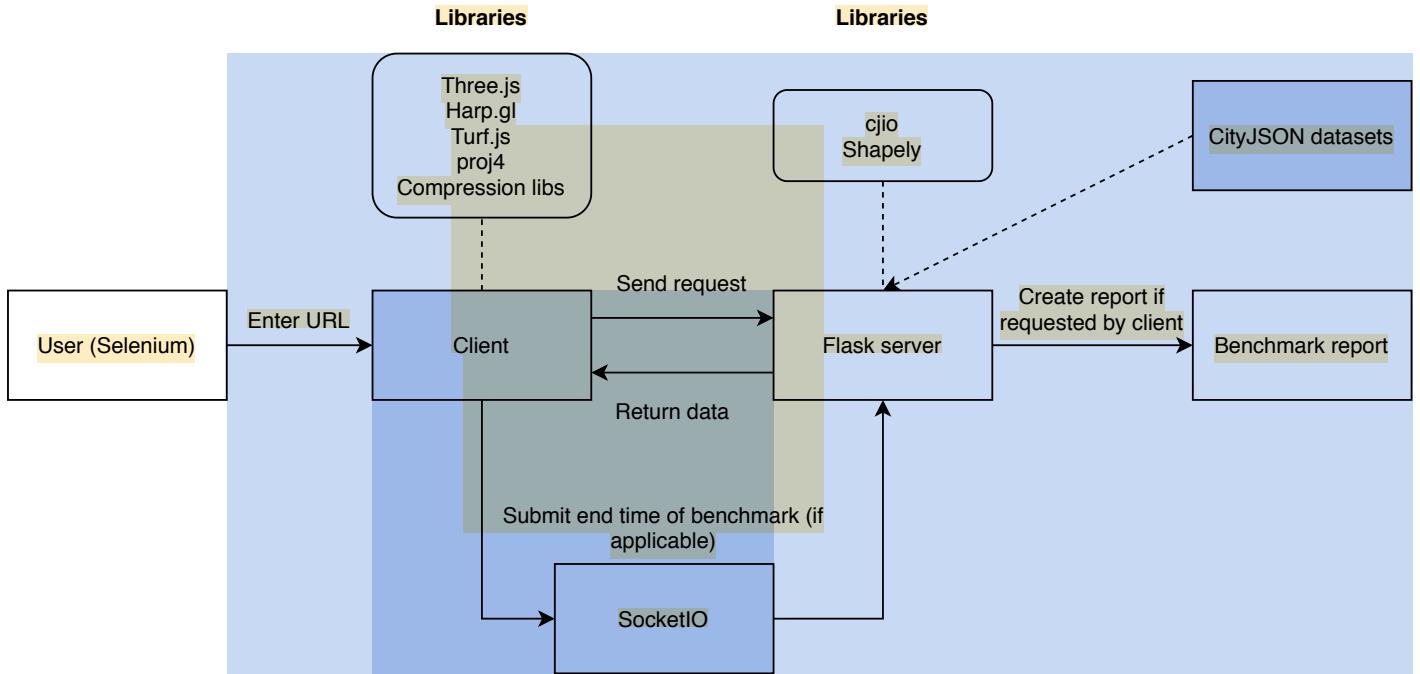


Figure 5.1: Diagram showing an overview of the structure of the testing platform.

5.1.1 Three.js

Three.js [Three.js, 2020a] is used for the visualisation of the datasets.

It is a higher level 3D library for JavaScript that helps visualising 3D geometries by providing functions that ultimately render graphics through the lower level WebGL. WebGL has the ability to visualise 3D graphics in the browser using hardware acceleration, allowing for greater performance compared to using the CPU only [WebGL, 2011]. In its coordinate system, geometries are being centered around the origin and their coordinates normalised, such that all fall within the range [-1, 1]. Objects can thus have negative coordinates and they are not georeferenced in Three.js [Three.js, 2020a].

The CityJSON viewer [Boersma, 2019] contains code that parses CityJSON geometries into a Three.js mesh. Using this the CityObjects can be visualised. Part of this code is used in the testing platform, to be used for data that has its geometries stored in CityJSON's native manner rather than as a Draco blob.

A Three.js mesh is always formed only of triangles and can contain one material [Three.js, 2020b]. Within a mesh object, the geometry can be stored in two manners. It can be stored in the mesh similarly to OBJ and CityJSON (see Section 3.1): as a combination of three-dimensional vectors as coordinates and triangular faces with index references to the vectors. Such an object is named a **Geometry** in the library and this is what the CityJSON viewer parses geometries of datasets into. The alternative is a **BufferGeometry** object [Three.js, 2020c], which is less user-friendly but processed more efficiently. It simply contains an array of coordinates as floats, which would represent a matrix.

This is one similarly for the faces. The Draco decoder for JavaScript decodes a Draco blob into a BufferGeometry. This difference can have an impact on the benchmarking results 

5.1.2 cpio

cpio [3D geoinformation research group at TU Delft, 2019] stands for CityJSON input/output. It is a Python command-line interface that has several functions that can be used to work with CityJSON files. It can be used to help preparing CityJSON datasets for the testing part, but more importantly, it can facilitate the querying and analysis of the datasets by integrating it with Flask and possibly Python libraries such as Shapely [3D geoinformation research group at TU Delft, 2019]. I chose to host files on the server and process queries with cpio and other tools, but another option is to use a Database Management System (DBMS) such as 3DCityDB [Technische Universität München, 2020].

5.1.3 Draco

As explained in Section 3.2, Google’s Draco library is an open-source library to compress and decompress 3D geometric meshes for the purpose of improved storage and transmission efficiency. It is used for the compression of the geometrical part of CityJSON’s file structure[Google, 2019b].

The Draco library includes an encoder and decoder for JavaScript, which can conveniently be used in the testing platform. Three.js has a function which decodes Draco using the previously mentioned decoder and loads it as a mesh that can subsequently be visualised. However, as opposed to Draco’s C++ decoder, the one in JavaScript lacks the functionality for separating sub-objects (as features or CityObjects are called in Draco)—it would create one large mesh out of the Draco geometry. This makes it impossible to query objects which is necessary for several of the operations (see Section 5.2) in this thesis.

Therefore, I have altered the inner workings of the JavaScript decoder. I have added a part that is supposed to enable the passing of one or more object IDs, making it only return these specified objects as a mesh. This is only possible after the Draco blob has been fully decoded, adding time to the process. This is done by iterating over all faces of the decoded Draco geometry, querying the sub-object ID from its first vertex from a metadata querying function of Draco, and checking if it equals one of the IDs that had been previously passed to the decoder.

Additionally, the three.js function does not allow passing a Draco blob directly but requires a link instead. This link is subsequently downloaded into a BytesArray after which it is decoded. For this reason, Draco files are downloaded separately by visiting a Flask route. The Draco parts are taken out of the compressed CityJSON files in order to not affect the transmission speed and thus the benchmarking results.

The encoder, which is used after editing a file, can be used unaltered. However, it is not possible to choose a compression level between 0 and 10 like the C++ encoder can, but instead a choice can only be made between using Edgebreaker (Subsection 2.3.3) and the sequential connectivity encoding method (Subsection 2.3.4).

5.1.4 Testing platform overview

The diagram of Figure 5.1 shows an overview of the testing platform. The user (who is simulated by Selenium for the benchmarking—see Section 5.3) sends a request through the client to the Flask server. Subsequently the server processes the request and data if needed (if compression is done on the fly), and starts a timer for the benchmarking. It then returns the requested data (the full

5 Implementation and experiments

CityJSON file or a compressed variant in the case of compression in advance) and information on the requested operation to the client. The client parses the data (decompressing it if necessary) and finishes the operation. Both the client and server are supported by several libraries. Once the operation is finished, it will emit a message with the end time of the benchmark to SocketIO [Socket.io, 2020], which allows communication between JavaScript and Python. Socket.io in turn sends the message to the Flask server. The latter calculates the time that has lapsed and stores this, and a report of the benchmark is created when requested by the user.

5.2 Operations

The way in which operations are implemented influences the time benchmarking results. Additionally, they can be different for varying data types, and understanding what is happening under the hood helps with reasoning why the results of compression types turned out as they did.

In every case, once the client makes a request to the server, the server stores the starting time of the process. It then processes uncompressed CityJSON (and compresses) if needed and stores the result as a file. Flask will then render a HyperText Markup Language (HTML), while submitting information to the client about the process—the name of the requested dataset, the compression type, and the operation name. The information is used to enable the client to parse the correct URL with which the needed file can be downloaded.

5.2.1 Visualisation

As the file that needs to be visualised is downloaded, it is decoded as described in Section 5.4 if necessary. If it contains original CityJSON geometry objects, these are converted into a three.js mesh using code based on the CityJSON viewer. Otherwise, the JavaScript Draco decoder is used to convert the Draco geometry to a three.js mesh. In that case also the object IDs to which every vertex belongs has to be queried using `MetadataQuerier()` of the Draco decoder. This is necessary because the Draco decoder creates one large mesh out of all objects, making it impossible to link attributes to separate objects. When visualising a 3D city model it is important that this can be done, and it makes the comparison to the visualisation of Geometry Objects fair.

When this is finished, the three.js mesh is added to the harp.gl map and visualised. The client then returns the current time to the server, which will calculate the time that has lapsed and add it to the benchmarking report.

5.2.2 Querying

Querying one object

The first object of the dataset is queried on its ID. With compression on the fly this is done on the server, otherwise in the client. The queried city object is taken by accessing the value of the ID as key in the "CityObjects" member of the CityJSON object. It is placed in a new CityJSON object, and its geometry is then processed by cjo's "subset.process_geometry" function, retaining only the vertices that belong to the object and updating the vertex indices in the geometry object to start with 0. The bounding box is updated, and the CityJSON object is then stored in a file and can be downloaded by the client. As the download is completed, the client will submit the end time to Flask. Since cjo does not exist for JavaScript I have translated the "subset.process_geometry" function to it, in case the querying is done in the client.

With Draco geometries, the difference is that the vertices corresponding to the object ID need to be found with "MetadataQuerier()" of the Draco decoder. As explained in Section 5.1.3 however this does not go well. Still, in any case all vertices would have to be traversed and this is being simulated.

Querying all objects

The functions do not know beforehand that all objects are queried, in order to simulate the querying process properly.

Uncompressed, all **city objects** are traversed and checked on containing the attribute field. If so, its ID is appended to an array. This array is passed to cgio's "get_subset_ids" function which will update the CityJSON object's City Objects, "transform", geometry, and bounding box. The resulting CityJSON object is stored as a file and can be downloaded by the client.

With compression types with original geometries, this is done in the same way. When using Draco, again all vertices of the Draco geometry have to be traversed.

5.2.3 Spatial analysis

Buffering one feature

With **compression on the fly**, the server will get the first City Object from the CityJSON object. The vertices have to be transformed back to original with cgio (thus removing "transform"), which means that they become floating points again and are not offset from the origin. This is necessary because the buffer result would be different if the coordinates are not represented in the same units (e.g. metres). The geometry of the object is converted to 2D Shapely polygons after which the library's buffering function can be used. The resulting buffer polygon is parsed into a CityJSON Geometry Object and replaces the geometry of the City Object. Accordingly the array of vertices is replaced, and thus contains the vertices of the buffer. The result is a CityJSON object containing the buffered **city object** which could then be visualised, but that is not included in the time benchmark.

As for the **server implementation with compression in advance**, it is done in similar fashion in the client. The difference is that the buffering is done by Turf.js rather than Shapely. This means that the coordinates have to be transformed to WGS84 first (using proj4), since Turf.js does not work with other coordinate systems. Another difference is that Draco geometries are not parsed into a Geometry Object but rather in the same way as Draco geometries are structured—a Three.js BufferGeometry (see 5.1.1).

Buffering all features

This is exactly the same as [Section 5.2.3](#), but all features are iterated and the results are parsed into one CityJSON object or Draco geometry.

5.2.4 Editing (attributes)

Editing one feature

With **compression on the fly**, the corresponding City Object is retrieved from the CityJSON object and added to the "CityObjects" array with the key being the ID with "1" added at the end. The **(compressed)** file is stored and can be downloaded by the client.

5 Implementation and experiments

As for compression in advance, again the complete file has to be downloaded. Normally the same thing as above is done, and the dataset is subsequently encoded again. For compression types using the "replace" method however I tried to do it more efficiently by directly finding the ID in the array of attributes and editing that as well. In this way the CityJSON skeleton would not have to be encoded again. However, I realised that this is not necessarily possible to do when it is something other than an object ID being edited. For example when a string is edited that actually occurred twice in the dataset, you can not just edit that string in the attributes array but would have to add a new string and update the correct reference in the JSON skeleton. It is therefore not a completely fair comparison as methods using "replace" are now compressed much quicker than they should.

When using Draco, the Draco geometry is not decoded as we know only attributes are edited. It can be handled the same as other compression types.

Editing all features

This is done in the same way as explained in Section 5.2.4, but for every feature.

5.2.5 Editing (geometry)

Editing one feature

With all compression types, the file is transmitted as a whole in order to simulate a user editing a complete file in the browser. This means that they need to be able to see the complete data.

When editing a geometry of uncompressed CityJSON, all vertices that appear in the geometry of the feature are increased in height by 1 unit. This is actually not a completely correct implementation, as the vertices could be shared with other features. Instead, the vertices should be copied before editing, and the original ones checked on being used by other features, otherwise removed.

With compressed datasets that have original geometries this is done in the same way. If using Draco, the vertices are edited within the decoded Draco geometry. This means they have the same problem as the geometry editing implementation of uncompressed CityJSON.

Editing all features

It is done in the same way as with one feature (see Section 5.2.5), but now all vertices of the dataset are edited. This means that the problem of editing shared vertices is not there as all features are moved up by 1 unit.

5.3 Testing and benchmarking with Selenium

The Selenium WebDriver [Selenium, 2020] enables the simulation of web browser use, which can be utilised to test web applications. It is used to iteratively visit the Flask routes that correspond to a test with a specific file type, dataset, and operation. By having it control Google Chrome, it is possible to configure network settings that simulate an actual network connection rather than letting the server run with local computer speed. This will mostly impact datasets that are larger than the set download speed limit. Performing an operation on them will take longer as their downloading is simulated.

The download speed is set at 40Mbit/s (or 5MB/s) as this is the average network speed in The Netherlands [Cable, 2019]. The network latency, which is the time it takes for a request to be

transmitted from a client to a server [KeyCDN, 2018], is set to 0ms as it is harder to determine an appropriate number for this and it would add up the same amount of time for every test iteration anyway.

During long processes on either the server or in the client a time out exception can occur, interrupting the test. This would especially happen while buffering large datasets. The time out could happen for multiple parts of the test platform: Flask-SocketIO, Chrome, or Selenium. When the problem occurs it is not always certain at which point the test iteration breaks.

Attempting to relieve this problem, I set Selenium to wait a long time for an alert to pop up with the message being that the current task has been finished before continuing with the next test iteration. Flask-SocketIO has a function to alter its time out settings as well, but this turned out to not be important anymore after changing the structure of the test platform, which is downloading a file after the html template starts rendering, rather than transmitting the data directly with `render_template()`. Chrome also has a time limit before it returns a time out. It however does not support the alteration of this parameter. Firefox on the other hand does support this, but in liaison with Selenium it can not be throttled.

Ultimately, after attempts to fix the problem, it only remained to pose a problem for buffering all geometries of the large datasets, which is seen in Section 6.6.2.

5.4 Compression and decompression

The compression techniques are implemented using Python and the Draco compression library. Decompression always takes place on the testing platform in JavaScript. This **Section** follows the order of Table 4.2.

5.4.1 original-zlib, original-cbor, and original-cbor-zlib

These are the simplest compression types. It simply involves to encode the full dataset with zlib, cbor, or first cbor and then zlib. For zlib, the input needs to be a Unicode-encoded string. It requires the zlib and flunn libraries for Python to compress the data, and the pako and CBOR libraries for JavaScript for decompression. The implementations are as follows: "cm.j" being the CityJSON data:

```
import zlib
encoded = json.dumps(cm.j).encode() # encoding the JSON into a Unicode string
compressed = zlib.compress(encoded) # compressing the encoded JSON with zlib
```

```
import flunn
compressed = flunn.dumps(cm.j) # parsing the JSON into CBOR
```

```
import zlib, flunn
cborcompressed = flunn.dumps(cm.j) # parsing the JSON into CBOR
finalcompressed = zlib.compress(cborcompressed) # compressing the CBOR with
zlib
```

Decompression is done in the browser client in the following ways:

```
import pako
decompressed = pako.inflate(compressed) # decompressing zlib-compressed data
```

```
import CBOR
jsonData = CBOR.decode(compressed) # parsing CBOR into JSON

import pako
import CBOR
decompressed = pako.inflate(compressed) # decompressing zlib-compressed data
jsonData = CBOR.decode(decompressed) # parsing decompressed data (which is
now CBOR) into JSON
```

5.4.2 original-cbor-smaz

This compression type is not used because it does not work out of the box. The idea is to recursively replace all keys and string values of the CityJSON object by its smaz-encoded equivalent. It can be done using the smaz library for Python, and it seemed to work well. However, decoding the strings in JavaScript with a port of the Python library gave wrong results that could not be used.

Firstly, the binary code (which would be of type "Uint8Array" in JavaScript) has to be converted to a string. It can then be given to the "decompress" function of the smaz library, which would retrieve the character code of every character of the string. The character code is then compared to the smaz codebook (containing fragments of words or bigrams, see 2.2.6) and converted. However, it seemed that the character codes of some characters of the string are much bigger than the size of the codebook, thus not returning a proper value.

5.4.3 original-replace, original-replace-zlib, original-cbor-replace, and original-cbor-replace-zlib

For the "replace" compression method, all keys and values are traversed with a recursive function, counting the amount of times that strings appear in the dataset. Numbers are ignored because they normally are small or unique. All strings are then placed in an array ordered by frequency of appearance. All keys and (string) values are subsequently replaced by the index of which they appear in the array. The index is stored as a string as well because numbers are not replaced and they would otherwise get mixed up with values that are intended as index. This means that the dictionary is now more like a JSON-skeleton with references to values in the array. The advantage is that strings that occur multiple times are now only stored once, resulting in a compressed representation of the data.

The result is a dictionary with three members: "cm", "str", and "geom", containing respectively the JSON-skeleton, the array of strings, and the original geometry of the dataset. An example is shown in Listing 5.1.

Listing 5.1: Example of dataset compressed with replace method

```
{
  "cm": {
    "type": "Compressed CityJSON",
    "CityObjects": {
      0: "1",
      2: {
        3: "4",
        5: "6",
        7: "8"
      }
    },
    "str": ["type", "building", "attributes", "bgt_status", "bestaand",
            "creationdate", "2014-07-09", "class", "dek"],
    "vertices": [...]
  }
}
```

For the zlib, cbor, and zlib+cbor versions, the same thing is done on the aforementioned dictionary as was done on the original datasets in Section 5.4.1. These would be decompressed in the same way as well. The created dictionary itself is decompressed by recursively traversing all keys and values in the JSON skeleton and replacing them all with the string in the array that the keys and values are referring to.

5.4.4 original-cbor-replace-huff

For this combination of compression methods, the array containing all keys and values of the JSON would be Huffman-encoded (see 2.2.4). It can be done by using the array as input for the creation of a frequency tree, and encoding the array based on this tree. Both the encoded array and the tree would have to be transmitted in order to be able to decode the array. However, the structure of the tree differs greatly between the Python and JavaScript implementations. In Python, it is a dictionary with every character and the corresponding binary code. In the JavaScript implementation on the other hand, a full tree is needed as input. It is a nested dictionary where every key (which is a leaf of the tree) contains its left and right child. Because a conversion is needed, this compression type is not implemented due to time constraints.

5.4.5 Draco-compressed equivalents

The difference with the counterparts with uncompressed geometry is that all vertices and Geometry Objects are stripped from the CityJSON object prior to compressing the attributes. The geometries are converted to OBJ using cgio, which in turn is compressed with Draco. This is done by calling the terminal with the appropriate command. The compression level is always 5 as it is in the middle, and metadata always has to be included since otherwise the separation between different objects is lost. Quantisation is turned off since it would result in the loss of coordinate precision (see Section 2.3.1).

The geometries are converted to OBJ with the following line of code:

```
from cgio import cityjson
obj = cm.export2obj()
```

Subsequently, the OBJ is compressed by Draco with a terminal command akin to:

```
draco_encoder.exe -i dataset.obj -o dataset.drc -cl 6 -qp 0 -metadata
```

Where `-cl` indicates the compression level, `-qp` the quantisation, and `-metadata` the separation between objects (rather than compressing them into one full mesh).

Since it is troublesome to load a Draco geometry directly from a blob, they are not embedded in the files as they are benchmarked (but still included in the compression time and file size results). These files are downloaded separately by passing a URL to the Flask server that will initiate the downloading of the Draco file belonging to the tested dataset. This means that Draco geometries are not compressed further when using zlib, at least when testing the compressed datasets on time performance. It has to be kept in mind when interpreting the results, because the files would actually have been even smaller. Compressing the Draco files separately would not give the same results either, as there will always be extra overhead when two files are compressed separately.



5.5 Dataset description

A selection of datasets that are available on the CityJSON and Random3DCity webpages has been made, with the addition of a city model of Rotterdam (see Table 5.1). As mentioned in Section 4.4, the data should have varying characteristics. There is a difference in the inclusion of terrains, language (Dutch and English), file size, and LOD. The four last ones have been chosen to make a fairer comparison between datasets of different LODs, as they do not contain attributes which can influence the results. This is unwanted, because it is the geometrical compression (Draco) only that should make a difference in that regard.

	File size (MB)	LoD	City Objects	Avg. attribute char	Avg. vertices	Other characteristics	Source
Delft	1.1	1	570	557	115	Dutch, contains terrain	CityJSON [2020]
Den Haag	2.9	2	2498	112	34	Dutch, contains terrain	CityJSON [2020]
Rotterdam	5.3	1 & 2	654	106	82	Dutch, multi-LOD	Municipality of Rotterdam [2020]
Montréal	5.4	2	294	0	64	No attributes	CityJSON [2020]
Singapore	52.8	1	10966	2978	106	English	Biljecki et al. [2017]
TU Delft campus	69	1	17654	59	273	Dutch	Ledoux [2019]
New York	105	2	23777	68	132	English	CityJSON [2020]
Zürich	292.9	2	198699	7	45	English	CityJSON [2020]

Table 5.1: Chosen datasets and their characteristics, sorted by file size.



For the analysis of the benchmarking results (see Chapter 6), I have separated the datasets into two categories: larger and smaller datasets. The first four of Table 5.1 are the smaller ones (Delft, Den Haag, Rotterdam, and Montréal), and the last four are the larger ones (Singapore, TU Delft campus, New York, and Zürch). In addition to that, I have stripped the datasets from their attributes for another category of datasets. This enables to make a simple distinction between dataset types, which can influence the workings of compression types—larger datasets, larger datasets without attributes, smaller datasets, and smaller datasets without attributes.

Snippets of the datasets can be found in Appendix A.

5.6 Experiments

Instead of transmitting (processed and optionally compressed) CityJSON datasets to the client, it is also possible to stream features. To do this, the dataset needs to be prepared such that every feature contains its own geometry, rather than having references to the vertices list at the end. The OGC API CityJSON prototype describes how such a dataset would look for CityJSON. To compress it, it would be possible to replace the geometries with a Draco object. However, testing this it turned

out that Draco does not perform as well when it has to encode every object separately - the file size would be bigger. Since Draco geometries would have to be decoded client-side, this means that it would only take longer to use it in this way. For this reason performance tests have not been completed with this streaming implementation.

5.6.1 Data loss

Zlib, CBOR, smaz, and Huffman encoding are explained in Chapter 2 and are completely lossless compression techniques.

Draco is mostly lossless when not using quantisation (see Sections 2.3.1 and 3.2). If you do use it, some vertices are lost and precision is decreased, of which the extent depends on the set quantisation level. **Draco is set to not use it, but the original datasets are prepared with the transform member, which means that all datasets are actually already quantised.** The compression levels that can be set in Draco (Section 3.2) are separate from this, and only determine the used compression techniques that the library encompasses (see Section 3.2).

I have compressed OBJS of every dataset (thus converted from CityJSON) into all Draco compression levels, and decoded the Draco geometries again. Comparing the original OBJS to the decoded OBJS shows that they are completely identical, so Draco compression itself does not induce any loss of information. However, geometries are triangulated when converting CityJSON to OBJ with cgio [3D geoinformation research group at TU Delft, 2019], so if the input file dataset is not triangulated the decoded result is not the same. **Draco compresses triangulated meshes, so this would always be the case.**

The replace method (explained in Section 5.4.3) is supposed to be lossless. Since I have created it myself, it needs testing to be sure. I have done this using the DeepDiff library [Dehpour, 2020] which can thoroughly compare Python dictionaries, and as such also two JSONs. It turned out that there are some cases of loss: in the attributes, floats that have 0 as decimal are converted into integers in the end. This can be significant in some cases, because it can be confusing because for example 5.0 indicates that the number is not rounded to an integer, whereas 5 might have been.

5.6.2 Querying compressed CityJSON

The possibility of querying compressed datasets gives an advantage when not all features are requested. With the current implementation of the testing platform, a compressed dataset has to be transferred as a whole, whereas original CityJSON can be prepared on the server before transmitting the data. This means that when for example only one City Object is queried, the transmitted data is smaller when using original CityJSON as opposed to a compressed variant, losing the benefits of compression.

It is not possible to query a Draco geometry directly, at least such a function does not exist in its JavaScript decoder nor its C++ one (see Google [2020a,b]. It has to be decoded fully before separate features can be retrieved. Additionally, zlib is a general purpose compression library (see Section 2.2.5) that has the purpose of compressing files completely and provides no possibility to partially decompress files.

CBOR (see Section 3.6.1) on the other hand might have the potential to be queried in similar fashion to JSON, as it has a key-value structure as well. The libraries that I use for CBOR [Yura, 2016; Hildebrand, 2020] in this thesis do not provide clear functions to do this however. This needs to be investigated further and can be future work. It could also be done with other binary JSON formats that are mentioned in Section 2.2.1.

5.6.3 Draco compression level

Draco has several compression levels that can be set when encoding an OBJ. In order to avoid having too many variables to test in the benchmarking chapter, one compression level is chosen for every compression type that includes Draco compression. For this purpose I have run tests using Pytest to pick one that is the most suitable on average. The variables are encoding time, file size, and decoding time.

Theoretically a higher compression level will result in a lower file size, at the cost of a longer encoding and decoding time. The encoding time is the least important in this case, because it is mostly part of the dataset preparation. The most ideal option is the compression level that has a good ratio between file size and decoding time, because this will increase the performance on the web the most.

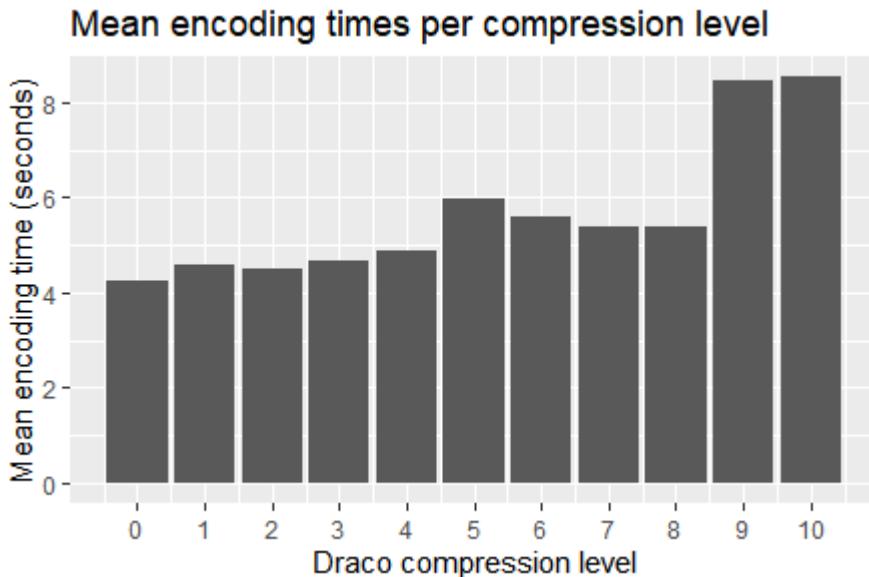


Figure 5.2: Benchmark of mean encoding times per different Draco compression level

Figure 5.2 shows that a higher compression level is not always slower than its predecessor. This is counterintuitive, but the reason for it is likely that individual objects are being encoded, rather than one big mesh. The different levels are probably determined based on their performance on larger meshes. Small meshes are less complex and compression techniques will therefore behave differently on them. Especially level 5 seems to be a bad choice if the encoding time is important, as it is even higher than 8 and 9. It is not known why level 5 in specific performs badly, as it is not specified in the Draco specifications which techniques are used per level.

As for the file sizes, Figure 5.3 shows the average file size multiplier of the Draco geometry compared to the original OBJ. There are three groups: 0, 1 to 4, and 5 to 10. Between the second two groups, the difference is probably because of the first one using the sequential connectivity method, and the second one using edgebreaker as this is said to be a main difference between compression levels. Compression level 0 might compress the vertices only and not the connectivity. Interestingly, despite level 5 taking relatively long to encode, its file size multiplier is not lower than several higher compression levels that encoded faster.

Figure 5.4 visualises the mean decoding times. It is tested with the Draco C++ library however, as opposed to the decoder for JavaScript which might perform differently. Between compression levels 1-6 there is not much of a difference, but 5 is again worse than 6. Despite it being the middle option,

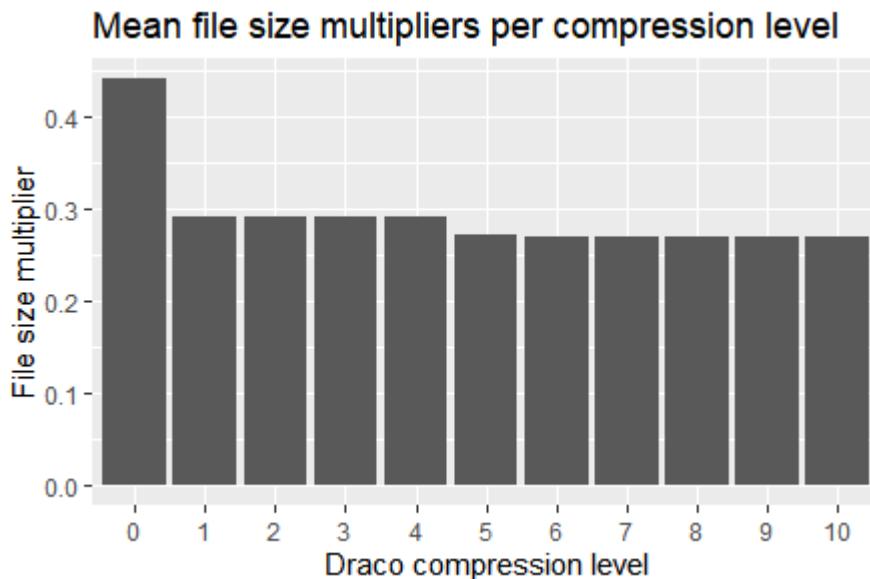


Figure 5.3: Mean file size multipliers per Draco compression level

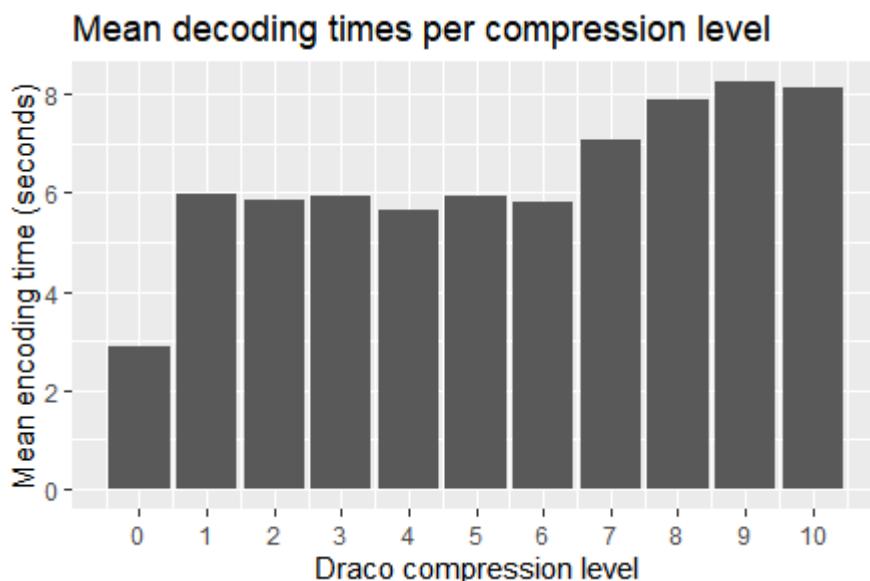


Figure 5.4: Benchmark of mean decode times per different Draco compression level

it turns out to not be an optimal choice. Level 0 is fast but it yields a relatively high file size. Levels 7-10 take longer to decode than 6 despite resulting in similar file sizes.

Compression level 6 has just a slightly longer decoding time than 4, but it is smaller in size. Conclusively, by having a low decoding time, small file size, and reasonable encoding time, it is the best option to use.

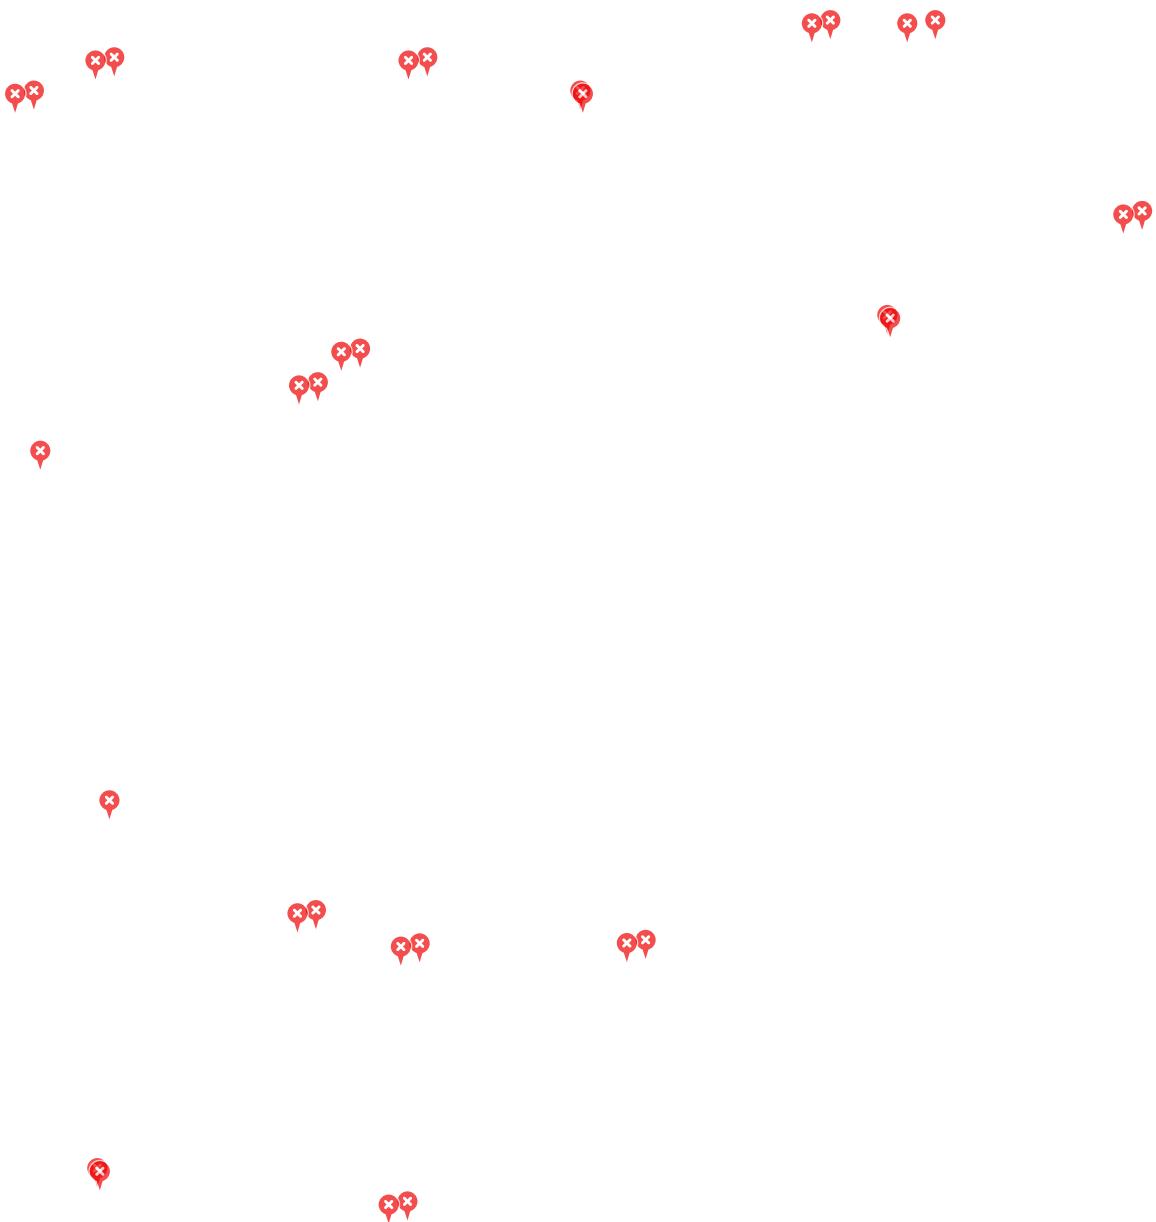


5.7 Computer specifications

The specifications of the used computer are listed in Table 5.2 below as the time benchmark results are influenced by it.

Model	Lenovo Thinkpad P1 Gen 2
CPU	i7-9750H@2.6Ghz
GPU	Nvidia Quadro T1000
OS	Windows 10 Home 64-bits
RAM	16 GB DDR4

Table 5.2: Specifications of computer used for benchmarking



6 Benchmarking results and analysis

The encoding, decoding, and file size performance are presented first in this Chapter as these results can help explaining the results in other parts, but they are naturally factored within the benchmarks of the operations (see Section 5.2) as well. After this the time benchmarks of the remaining performance indicators (see Table 1.1) are shown in separate sections. These Sections are subsequently divided per operation and by server implementation (with compression in advance or compression on the fly). Lastly there is a conclusion including a plot where all operations are averaged.

The datasets (see Section 5.5) are divided in two classes: bigger (Singapore, TU Delft, New York, and Zürich) and small datasets (Delft, Den Haag, Rotterdam, and Montréal). I arbitrarily did this because showing the results per dataset makes it too complex, while the file size is the clearest divide between the datasets and also a factor that is expected to be significant for the performance. In addition to that, for every dataset there is a variant from which the attributes are removed, with which it can be shown how the presence of attributes impacts compression.

Since the results are averaged by dataset group, box plots are presented. These show some statistics as well on the variety of the datasets within a dataset group. In the charts that show the performance of an operation on a compressed dataset, the baseline (uncompressed CityJSON) is shown with a dotted line. The time multiplier on the x-axis shows how much time it took to complete the specific task as compared to the baseline—a time multiplier of 0.5 means that it took half the time of doing the same thing with uncompressed CityJSON, and a time multiplier of 2 means that it took twice as long. The scales of the plots are not always fixed, which means that the four plots (large, large without attributes, small, and small without attributes) within one figure can have different scales. This requires extra attention when interpreting the results, but it is necessary to make them readable when the time multipliers between the two plots are very different.

Besides these plots, for all operations a table is added that shows the median performance of every compression type per dataset type. It is sorted ascendingly so that it can easily be seen which compression types are the best for the corresponding operation. 

The conclusion (Section 6.9) gives an overview of the results, indicating the advantages and disadvantages of every compression type. It includes two plots that show an overview of the performance of the compression types (Figures 6.22 and 6.23). Again per dataset group, but with all operations averaged. Through this it can more easily be noticed how different compression types compare when the wanted operation is unknown. Thus answering the question: "when you want to compress a dataset of this size and use it for any kind of purpose, which compression type would be advisable?".

Tables of the full results that are used to create the charts of the operations are found in Appendix D. For some combinations of use cases and compression types there were errors in the benchmarking. These are shown in those tables with NA and it results in missing compression types or strange looking boxplots in the charts.

6.1 Encoding performance

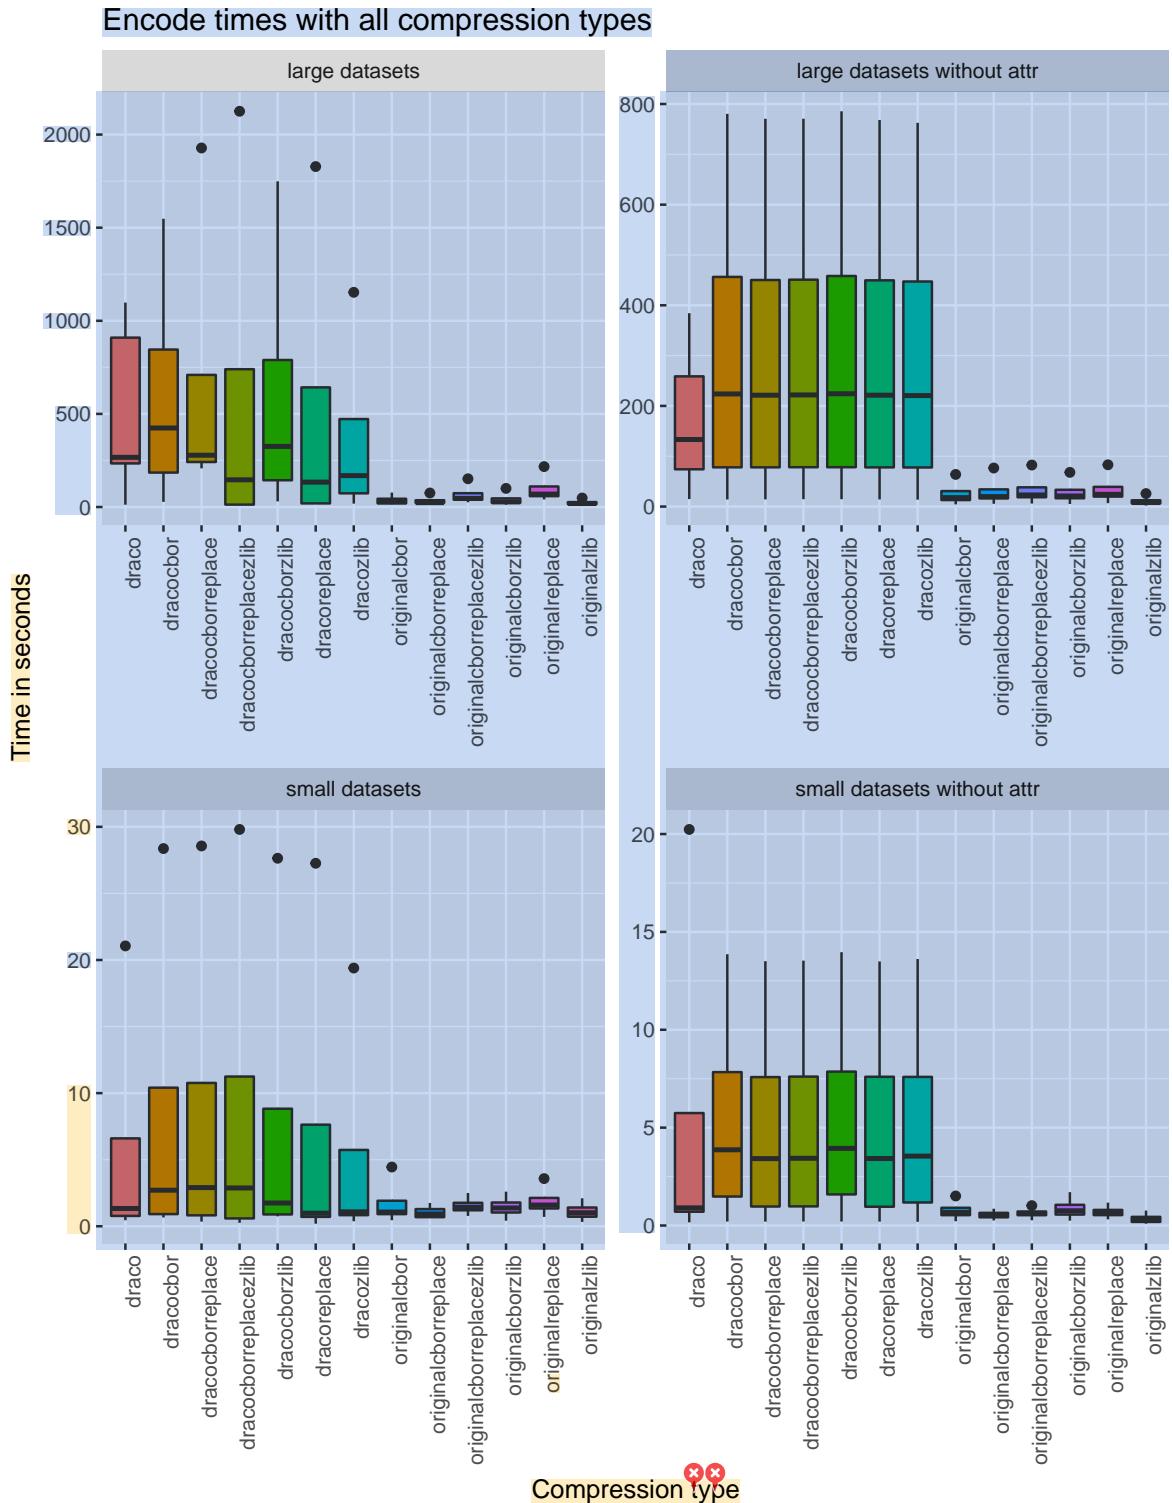


Figure 6.1: Boxplots of the encoding time performance of different compression types.

The encoding times shown here only apply to the preparation of datasets in Python (which is linked to C++ for Draco compression). Therefore they are not an influence on the other performance

results. After editing datasets they also need to be encoded, but this is done in the client and not in the server in the case of compression in advance. Encoding times in the client are incorporated in the editing benchmark results for that use case.

For datasets with attributes, compression types that include Draco always take considerably longer to encode, which is true for both bigger and small datasets. That the impact is this large is likely because of a CityJSON dataset having to be converted to OBJ first before it can be encoded by Draco. For large datasets, Draco compression itself already takes four minutes. It could become a large share of the execution time of the preparation of datasets. For small datasets this is not a problem because it is still fast in absolute terms, but it is still more expensive.

Draco-cbor on average takes longest for bigger datasets. Combining it in draco-cbor-replace-zlib it is faster. Possibly this is because CBOR can encode the keys and values of a dataset faster in this structure, where all strings are contained in one array and encoded in zlib which results in a binary value. For small datasets these two compression types have similar performance, so the difference in JSON-structure (which is parsed into CBOR) is only noticed with a large dataset. Draco-cbor-replace-zlib is the fastest of the Draco ones together with draco-replace.

I had expected that the replace method would be expensive, but there is no clear pattern in it. Draco-replace is faster than the other Draco types for bigger datasets, but original-replace is the slowest one without Draco. The other types that use replace but also are in CBOR are encoded faster, despite having more steps in the compression. It could possibly explained by CBOR being written to a file faster than JSON, or it is just simply written faster due to its size being smaller.

Original-zlib is however the fastest on the bigger datasets. Zlib thus seems to be the fastest compression technique to use on bigger datasets, but with smaller ones it is not significantly faster than others. Based on both plots, original-cbor, original-cbor-replace, or original-zlib would be the best choices if the encoding time is important.

As for the datasets from which the attributes are stripped, the results are not very much different. For both the large and the small datasets the encoding time is lower, which is to be expected since there is less data to compress. With the large datasets you also see that the median times for the compression types seem more uniform, which is because the geometry compression is now relatively more important, and this is generally divided into Draco compression versus (more or less) retaining the structure of the original geometries (but they are most of the time still compressed by zlib or CBOR, see Section 5.4).

6.2 Decoding performance

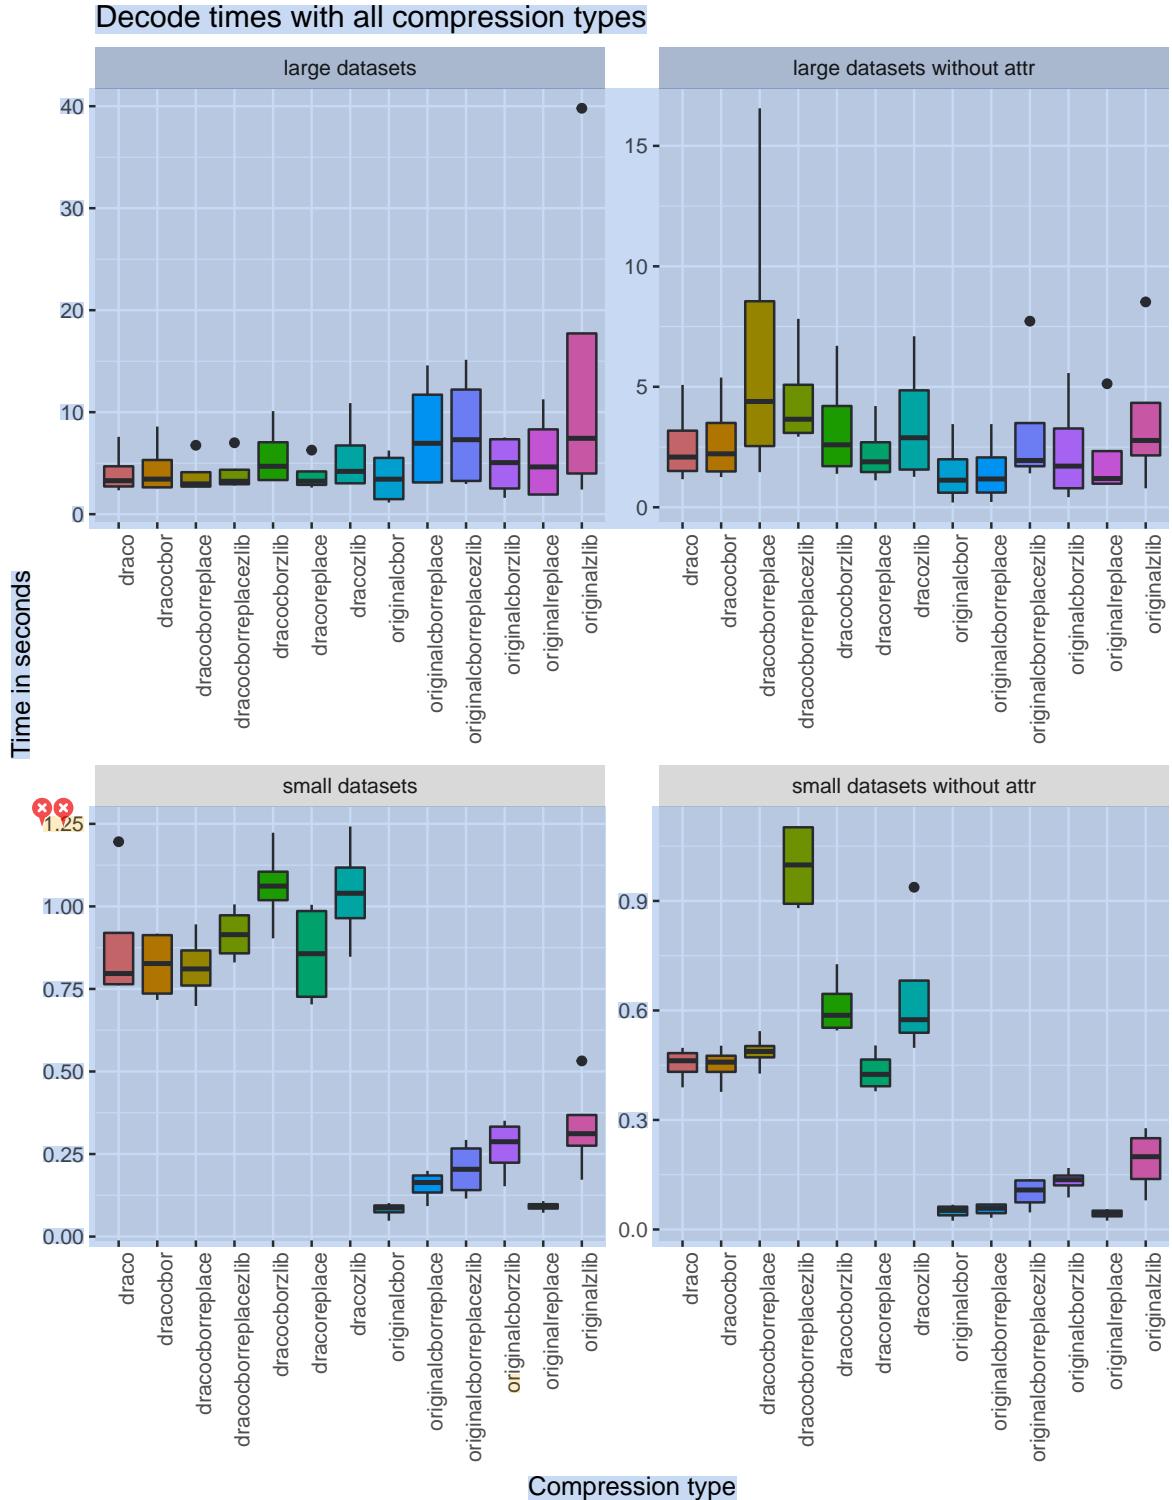


Figure 6.2: Boxplots of the decoding time performance of different compression types.

The decoding times are shown with a normal scale as the results do not have a large spread as there was with the encoding performance. These statistics should be a significant explaining factor

throughout the results chapter as it is a main property that explains the difference in performance between varying compression types.

While encoding Draco took a comparatively long time, it is generally decoded quicker with the bigger datasets ([with attributes](#)) than the other compression types. This is not the case for the [small](#) datasets, suggesting that the larger a dataset is, the more suitable it is to use Draco.

Original-cbor on the other hand has a low mean decoding time as well, for both dataset types. Combining CBOR in any way with replace will decrease the performance, probably because it is expensive to traverse all keys and values and replace them.

Original-zlib takes long to decode, contrasting its encoding results. Its outlier on the left plot is the Zürich dataset, which took a considerably longer time than decoding the same dataset compressed in another way. The extent of the corresponding box plot also indicates that the larger a dataset is, the longer it takes to decode it as compared to the other compression types. This can potentially pose to be a disadvantage of original-zlib.

The datasets without attributes again show faster performance since there is less data to decode. Because the replace method has (almost) no attributes to decode, it now adds less time to the decompression, although it is probably also not useful to use it when compared to original CityJSON. However, draco-cbor-replace (large datasets without attributes) and draco-cbor-replace-zlib (small datasets without attributes) do have notably higher decoding times, but I do not see plausible explanations for this.

6.3 File size

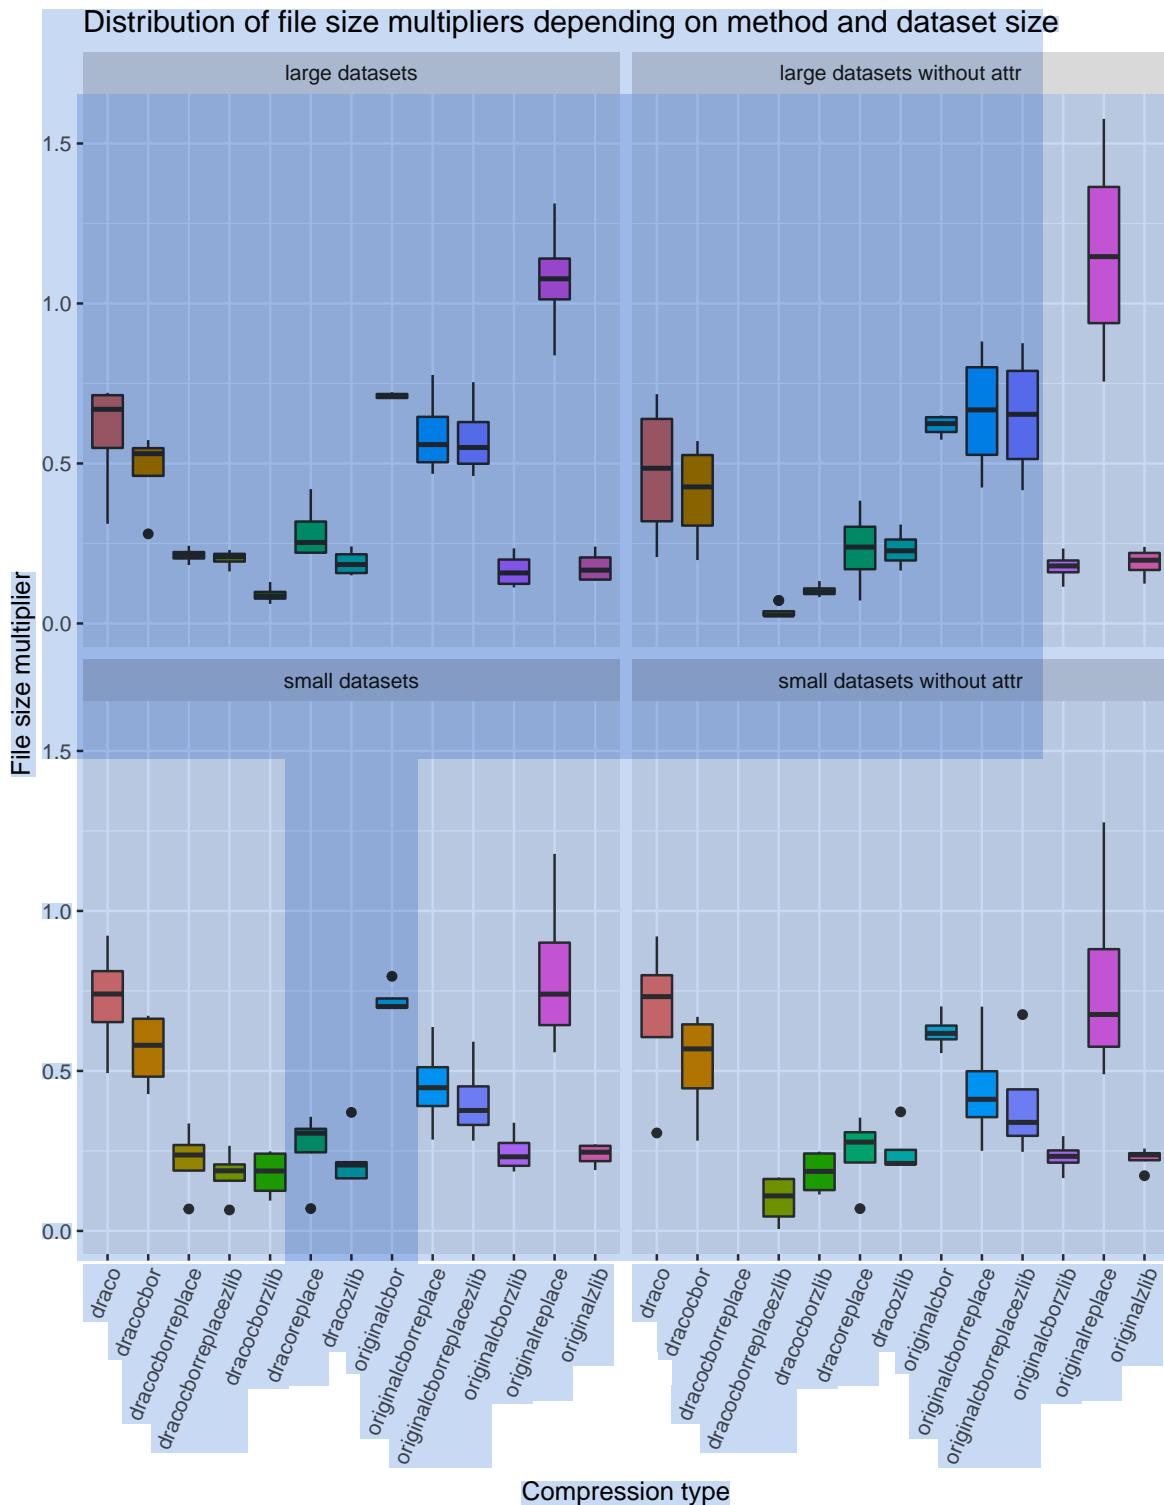


Figure 6.3: Boxplots of file size multipliers indicating to which extent compression types compress the data.

The file size multipliers are similar between the bigger [and small datasets, and datasets with or without attributes](#).

out attributes. The latter observation indicates that the compression of actual city object attributes does not contribute as much to compression as a whole, and that the compression of geometries, the JSON structure (CBOR), and general purpose compression (zlib) are more beneficial, with the latter two also naturally compressing the attributes of city objects when they are used. It is again shown that compressing geometry with Draco is mostly suitable for bigger datasets. These datasets have a lower file size multiplier than the small datasets compressed with Draco.

Original-replace is the only compression type that in some cases makes a file larger than the original. The array of attributes is then larger than the removed redundancy. The replace method in general does not yield very favourable results, even when the array of attributes is encoded.

Combining CBOR and zlib or using zlib only gives the smallest file sizes when compressing bigger datasets, for both Draco compression types and the ones with original geometry. This trend is seen for the small datasets as well. So compression types with zlib are both encoded quickly (Section 6.1) and yield small files. On the other hand zlib takes long to decode as shown in Section 6.2. Further performance benchmarks will thus demonstrate whether or not the decreased file sized outweighs the increased decoding time.

6.4 Visualisation

6.4.1 Compression in advance

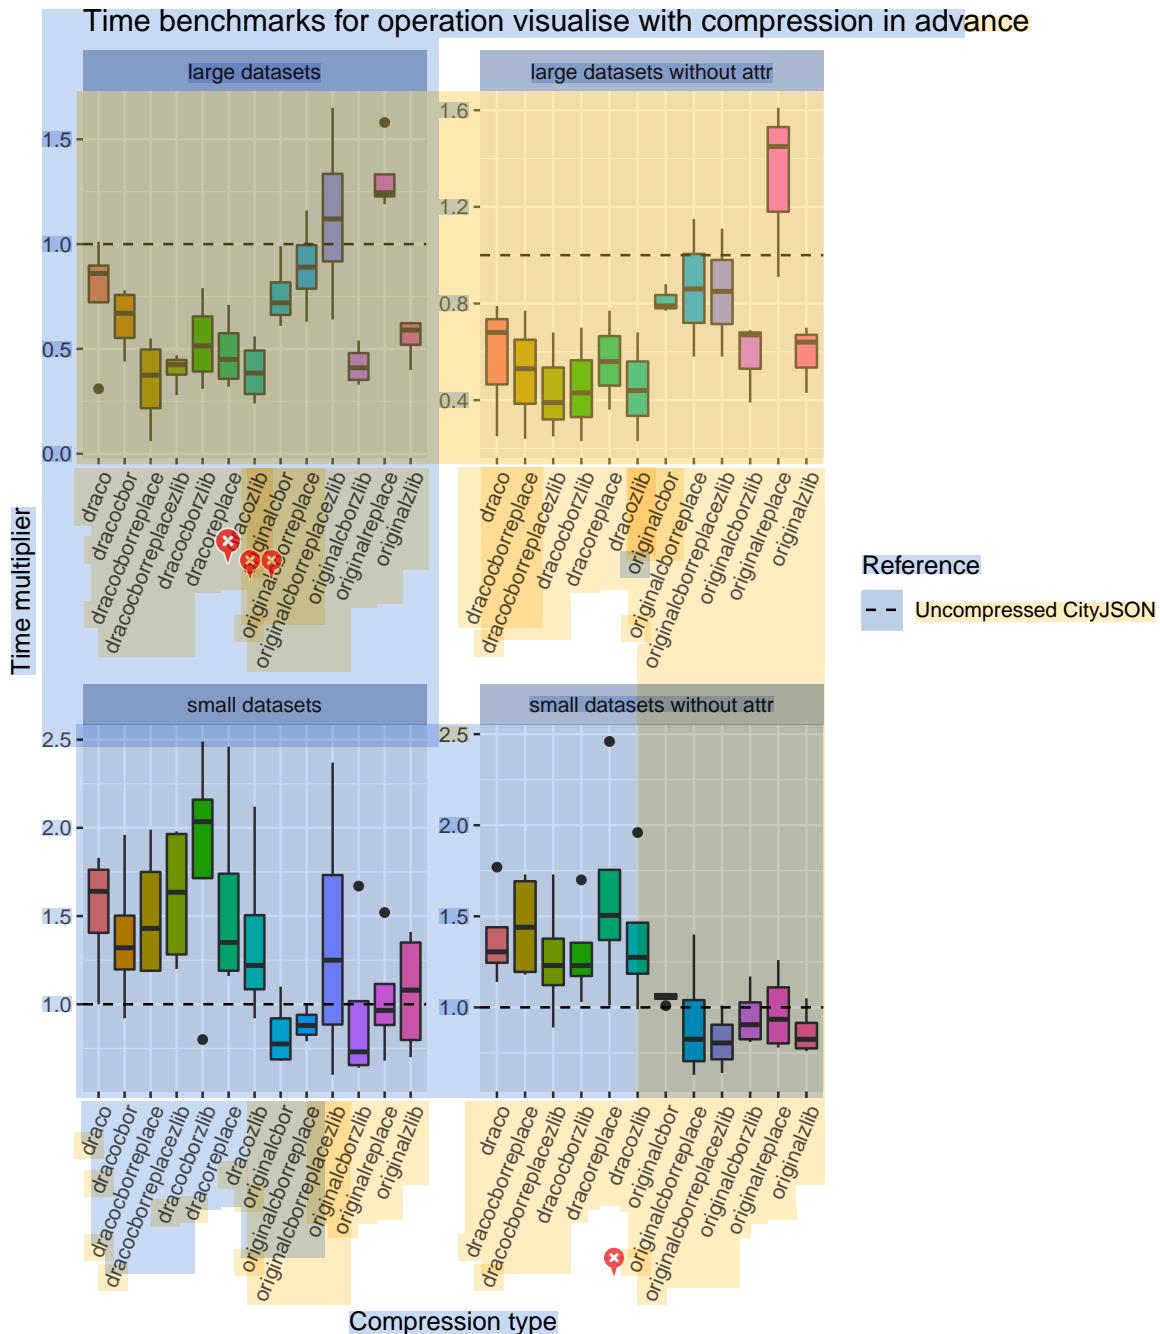


Figure 6.4: Boxplots of time performance of (in advance) compressed datasets on visualisation compared to uncompressed CityJSON.

Visualising compressed small datasets often yields a performance that is worse than the baseline, but this is especially the case when using Draco. It is not worth using it for small data, which can be explained by the overhead of Draco compression contributing a relatively large part in performance time. This is the same for small datasets that do not contain attributes. On the other hand,

original-cbor[✖] and original-cbor-zlib have a good median performance. Orginal-cbor-replace is an acceptable option as well. When not having to compress attributes, all compression types with original geometries retained do pretty well, except for orginal-cbor that suddenly has its median above the baseline, which is unexpected as this is not the case when attributes are included. It could be an abnormality in the benchmarking.

But for large datasets it is even more beneficial to use compression in general. With this use case it turns out that it can be worth it to use Draco if performance time is most important. But some compression types without it also do well, especially original-cbor-zlib and original-zlib. Even when attributes are not included, additional compression besides Draco proves to be useful as using Draco only comparatively performs worst amongst the Draco compression types.

Compression can improve the visualisation time of CityJSON datasets to 37%, 39%, 73%, or 80% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompresssed CityJSON, ^{✖✖✖} the case of compressing the dataset in advance.

Table 6.1: Median performance with visualise on large datasets, compression in advance

Compression type	median
dracocborreplace	0.375
draco zlib	0.385
originalcbor zlib	0.410
dracocborreplacezlib	0.425
dracoreplace	0.450
dracocborzlib	0.515
originalzlib	0.590
dracocbor	0.670
originalcbor	0.720
draco	0.860
originalcborreplace	0.890
originalcborreplacezlib	1.120
originalreplace	1.245

Table 6.2: Median performance with visualise on large datasets without attributes, compression in advance

Compression type	median
dracocborreplacezlib	0.39
dracocborzlib	0.43
draco zlib	0.44
dracocborreplace	0.53
dracoreplace	0.56
originalzlib	0.64
originalcbor zlib	0.67
draco	0.68
originalcbor	0.79
originalcborreplacezlib	0.85
originalcborreplace	0.86
originalreplace	1.45

Table 6.3: Median performance with visualise on small datasets, compression in advance

Compression type	median
originalcbor zlib	0.730
originalcbor	0.775
originalcborreplace	0.880
originalreplace	0.965
originalzlib	1.080
draco zlib	1.220
originalcborreplacezlib	1.250
dracocbor	1.320
dracoreplace	1.350
dracocborreplace	1.430
dracocborreplacezlib	1.635
draco	1.640
dracocborzlib	2.035

Table 6.4: Median performance with visualise on small datasets without attributes, compression in advance ^{✖✖✖}

Compression type	median
originalcborreplacezlib	0.805
originalcborreplace	0.825
originalzlib	0.825
originalcbor zlib	0.905
originalreplace	0.935
originalcbor	1.065
dracocborreplacezlib	1.230
dracocborzlib	1.230
draco zlib	1.275
draco	1.305
dracocborreplace	1.440
dracoreplace	1.505

6.4.2 Compression on the fly

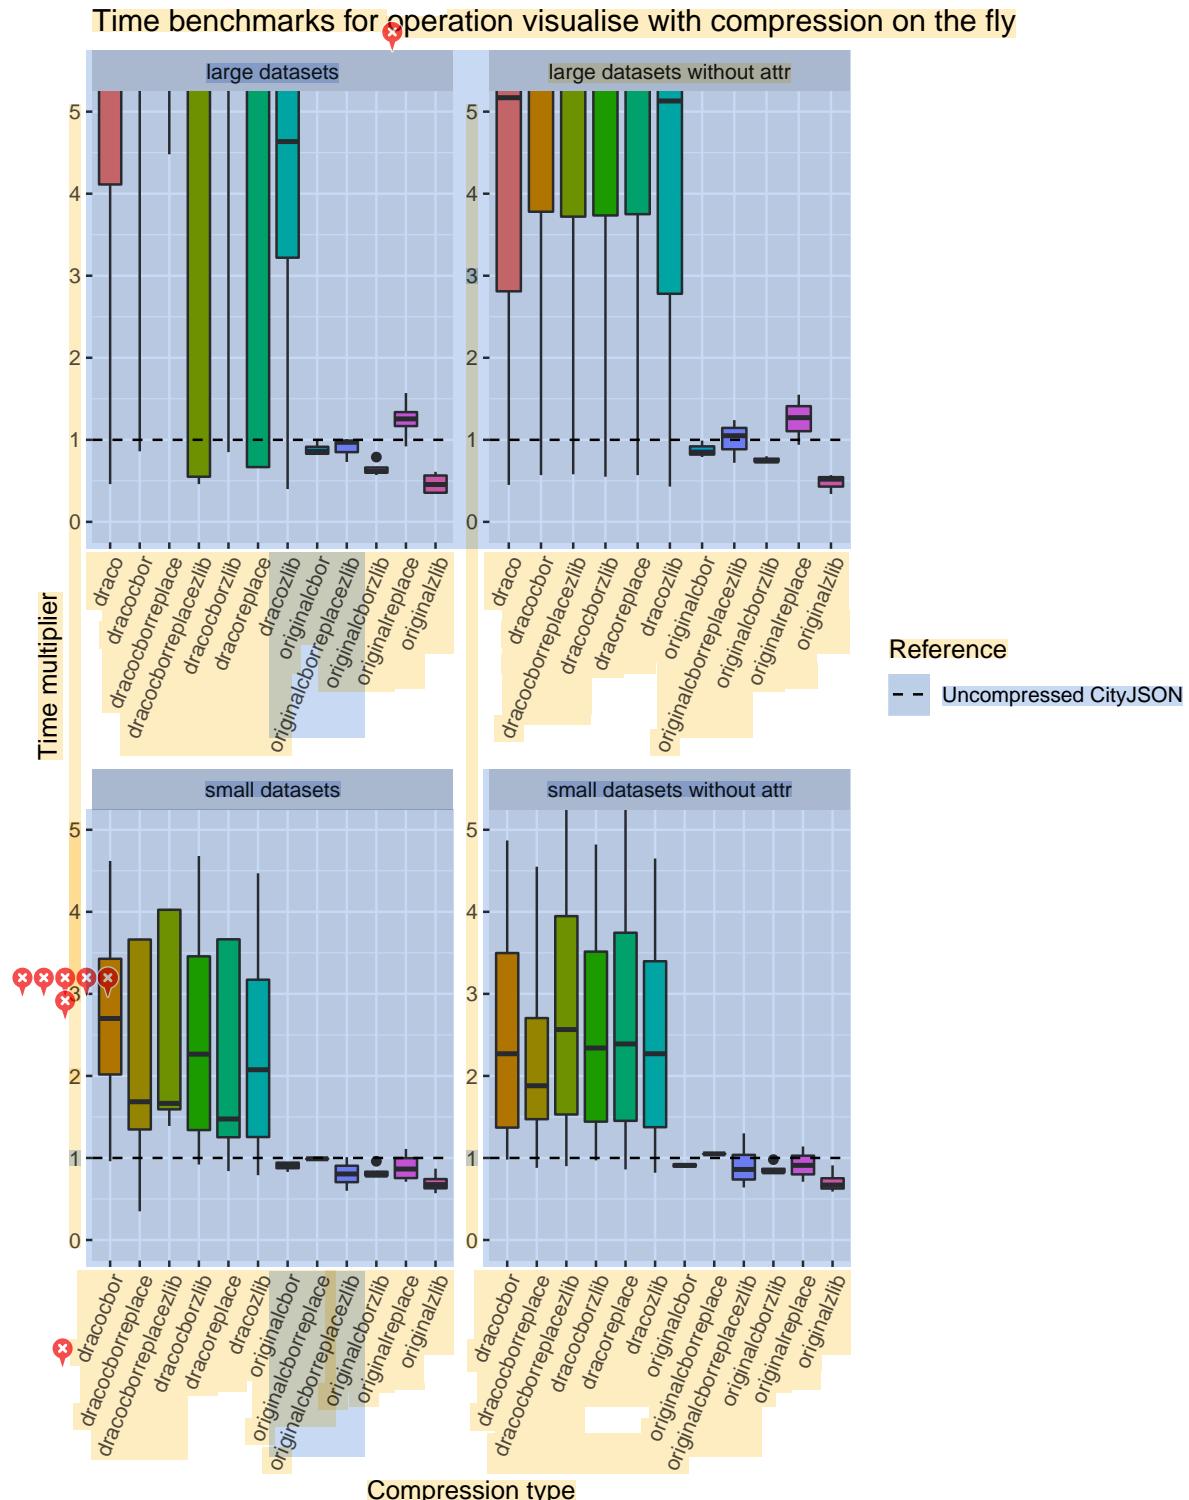


Figure 6.5: Boxplots of time performance of (on the fly) compressed datasets on visualisation compared to uncompressed CityJSON.

When compressing data after a request for the visualisation of a full dataset is made, the performance is slower since now the compression time of the compression type is of influence on the

results. This means that none of the compression types using Draco are suitable for use, since it takes relatively long to compress (which is also influenced by having to convert the geometries to OBJ first, see Section 5.4.5). They are all way above the performance baseline. The y-scale of the graphs is limited for this reason.

Compression types that do not use Draco are generally beneficial. Original-cbor, original-cborzlib, and original-zlib are best to use, without distinction between dataset type. The median performance gain is however slightly larger for large datasets.

On the fly compression can improve the visualisation time of CityJSON datasets to 45%, 52%, 67%, or 67% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON, in the case of compressing the dataset on the fly.

Table 6.5: Median performance with visualise on large datasets, compression on the fly

Compression type	median
originalzlib	0.455
originalcborzlib	0.615
originalcbor	0.860
originalcborreplacedzlib	0.970
originalreplace	1.255
dracozlib	4.635
draco	5.580
dracoreplace	6.015
dracocborreplacedzlib	6.460
dracocborreplace	8.670
dracocborzlib	8.705
dracocbor	11.465

Table 6.6: Median performance with visualise on large datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.52
originalcborzlib	0.74
originalcbor	0.85
originalcborreplacedzlib	1.05
originalreplace	1.27
dracozlib	5.13
draco	5.17
dracocborreplacedzlib	6.86
dracocborzlib	6.92
dracoreplace	6.93
dracocbor	6.99

Table 6.7: Median performance with visualise on small datasets, compression on the fly

Compression type	median
originalzlib	0.675
originalcborzlib	0.790
originalcborreplacedzlib	0.805
originalreplace	0.865
originalcbor	0.915
originalcborreplacedzlib	0.990
dracoreplace	1.475
dracocborreplacedzlib	1.665
dracocborreplace	1.685
dracozlib	2.075
dracocborzlib	2.265
dracocbor	2.700

Table 6.8: Median performance with visualise on small datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.670
originalcborzlib	0.830
originalcborreplacedzlib	0.860
originalreplace	0.910
originalcbor	0.910
originalcborreplacedzlib	1.050
dracocborreplace	1.880
dracocbor	2.270
dracozlib	2.270
dracocborzlib	2.340
dracoreplace	2.390
dracocborreplacedzlib	2.565

6.5 Querying

6.5.1 Querying one feature

Compression in advance

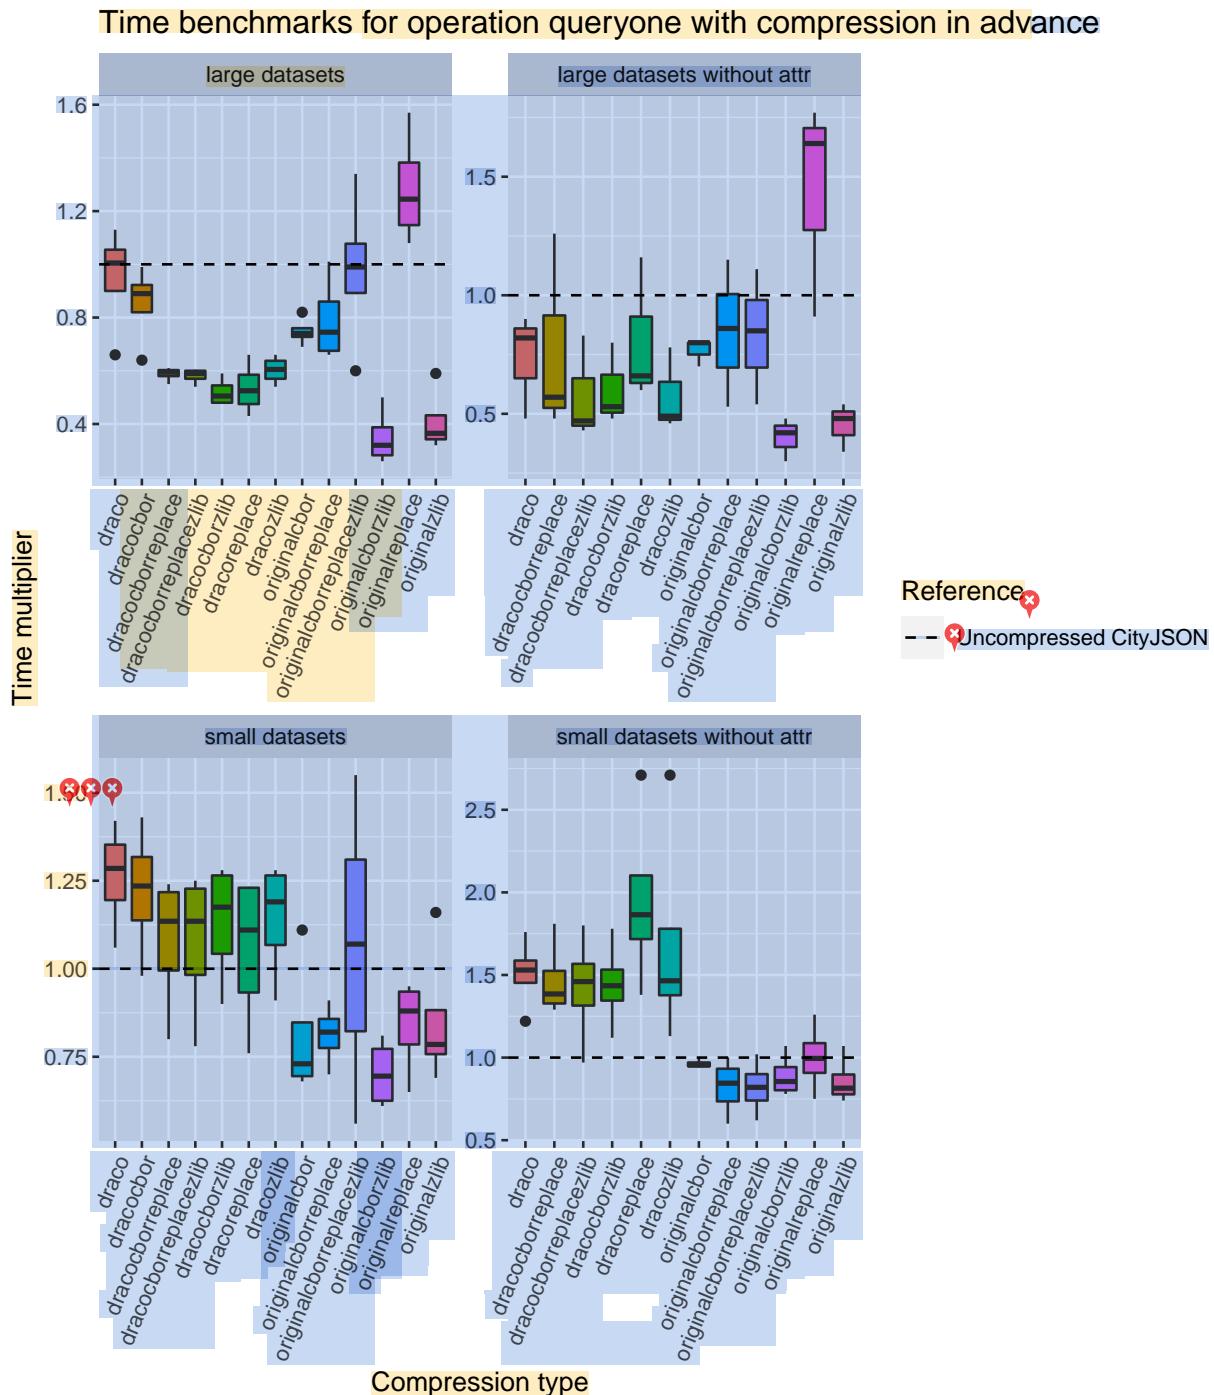


Figure 6.6: Boxplots of time performance of (in advance) compressed datasets on querying one feature compared to uncompressed CityJSON.

With the large datasets, most compression types have good performance, especially when datasets do have attributes. Compression types with original geometries and using the replace method do not do well, and this would be due to only one city object having to be compressed, which does not have a lot of repeated parts in it. Draco compression is not clearly better or worse. Original-cborzlib and original-zlib do especially well with this use case, indicating that they have a good balance between file size compression and decoding time.

With the small datasets using Draco yields worse performance. This is not expected when comparing it to the results of the large datasets, as in both cases only one city object has to be compressed. Without using Draco compression beneficial results can be achieved. There are no clear differences between having attributes or not, except for the results of original-cbor-replace-zlib having a large spread which I can not explain.

Compression in advance can improve the querying time of CityJSON datasets on one city object to 32%, 42%, 69%, or 81% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.9: Median performance with queryone on large datasets, compression in advance

Compression type	median
originalcborzlib	0.320
originalzlib	0.365
dracocborzlib	0.505
dracoreplace	0.525
dracocborreplacezlib	0.590
dracocborreplace	0.595
dracozlib	0.605
originalcbor	0.740
originalcborreplace	0.745
dracocbor	0.890
originalcborreplacedzlib	0.990
draco	1.005
originalreplace	1.245

Table 6.10: Median performance with queryone on large datasets without attributes, compression in advance

Compression type	median
originalcborzlib	0.42
dracocborreplacezlib	0.47
originalzlib	0.48
dracozlib	0.49
dracocborzlib	0.53
dracocborreplace	0.57
dracoreplace	0.66
originalcbor	0.80
draco	0.82
originalcborreplacedzlib	0.85
originalcborreplace	0.86
originalreplace	1.64

Table 6.11: Median performance with queryone on small datasets, compression in advance

Compression type	median
originalcborzlib	0.695
originalcbor	0.730
originalzlib	0.785
originalcborreplace	0.820
originalreplace	0.880
originalcborreplacedzlib	1.070
dracoreplace	1.110
dracocborreplace	1.135
dracocborreplacezlib	1.135
dracocborzlib	1.175
dracozlib	1.190
dracocbor	1.235
draco	1.285

Table 6.12: Median performance with queryone on small datasets without attributes, compression in advance

Compression type	median
originalzlib	0.815
originalcborreplacedzlib	0.820
originalcborreplace	0.845
originalcborzlib	0.855
originalcbor	0.955
originalreplace	0.995
dracocborreplace	1.385
dracocborzlib	1.435
dracocborreplacezlib	1.460
dracozlib	1.465
draco	1.530
dracoreplace	1.865

Compression on the fly

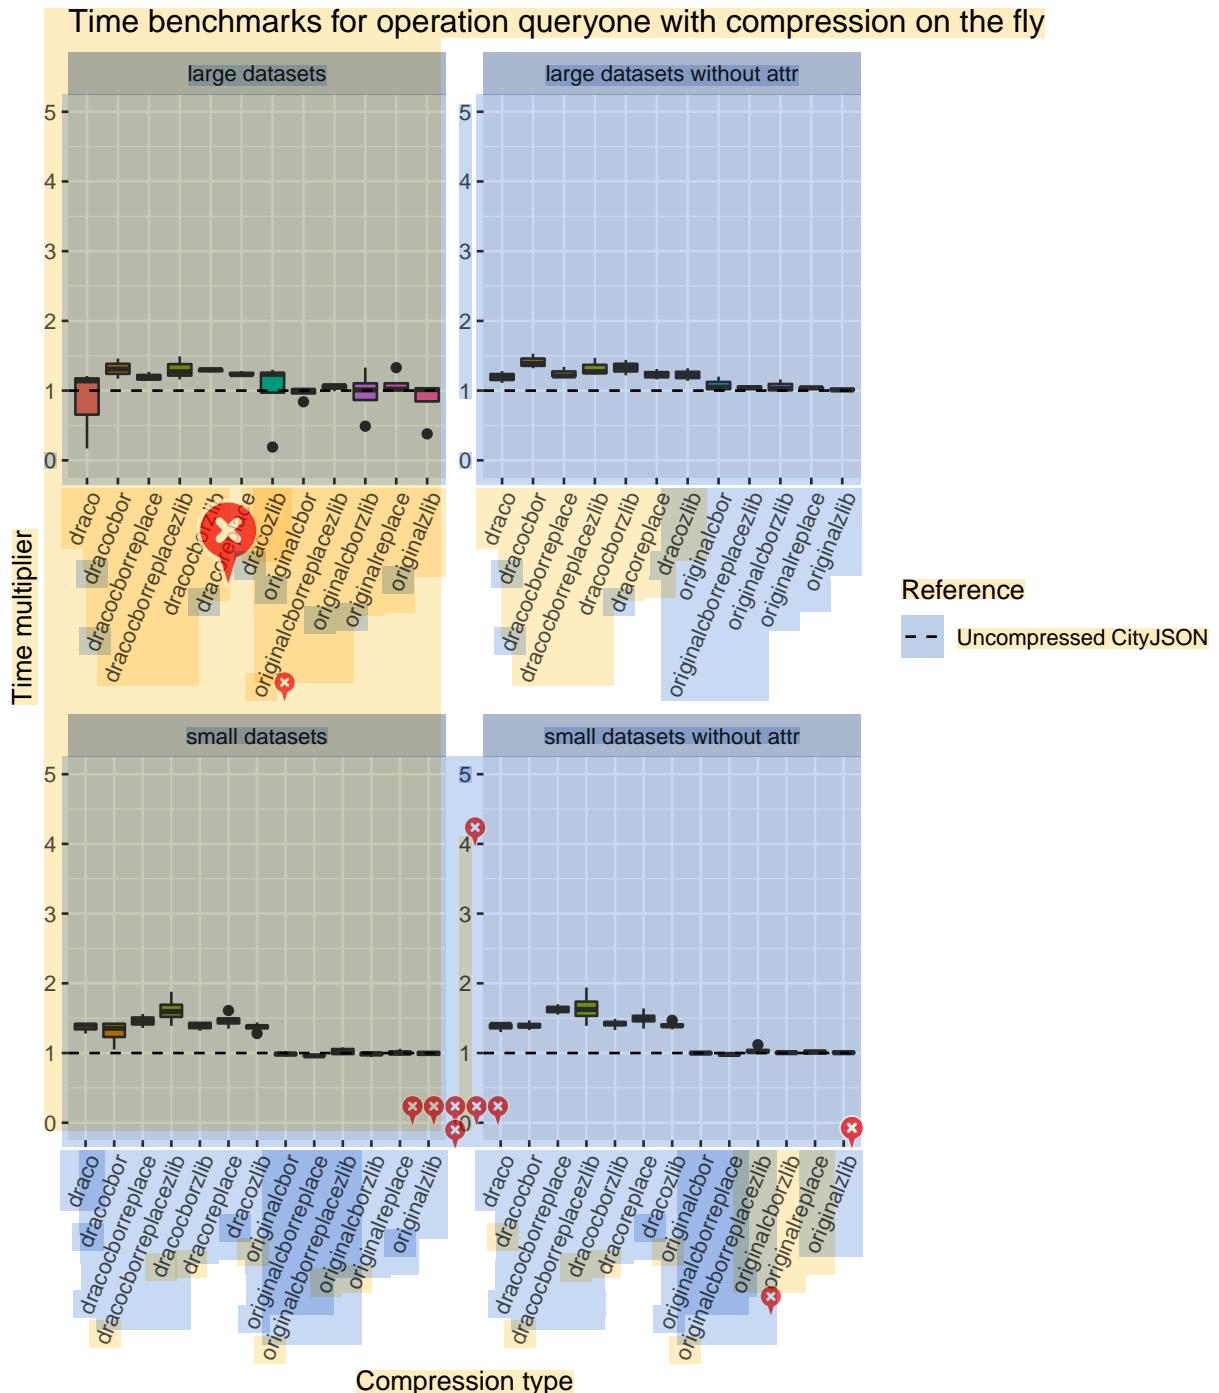


Figure 6.7: Boxplots of time performance of (on the fly) compressed datasets on querying one feature compared to uncompressed CityJSON.

For this use case, Draco compression is detrimental for performance. Compression time is added into it and this is not worth it when compressing only one city object. Other compression types are also often not worth it, especially on the large datasets, regardless of attributes. With the small datasets, only original-cbor-replace works well with both dataset types. One city object would not have much redundancy to compress and is already slim, so compression is not useful here and might

even give slower performance.

On the fly compression can improve the querying time of CityJSON datasets on one city object to 100%, 100%, 96%, or 98% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time^{xx} that it takes to perform the same operation with uncompressed CityJSON.

Table 6.13: Median performance with queryone on large datasets, compression on the fly

Compression type	median
originalcbor	1.005
originalcborzlib	1.010
originalzlib	1.015
originalreplace	1.020
originalcborreplacezlib	1.070
draco	1.140
dracocborreplace	1.180
dracoreplace	1.230
dracozlib	1.235
dracocborreplacezlib	1.270
dracocborzlib	1.300
dracocbor	1.310

Table 6.15: Median performance with queryone on small datasets, compression on the fly

Compression type	median
originalcborreplace	0.960
originalcbor	0.985
originalcborzlib	0.990
originalzlib	0.995
originalreplace	1.000
originalcborreplacezlib	1.020
dracocbor	1.355
dracozlib	1.380
draco	1.400
dracocborzlib	1.400
dracoreplace	1.455
dracocborreplace	1.460
dracocborreplacezlib	1.595

Table 6.14: Median performance with queryone on large datasets without attributes, compression on the fly

Compression type	median
originalzlib	1.010
originalcborzlib	1.030
originalcborreplacezlib	1.040
originalreplace	1.040
originalcbor	1.050
draco	1.195
dracocborreplace	1.210
dracoreplace	1.220
dracozlib	1.220
dracocborreplacezlib	1.270
dracocborzlib	1.330
dracocbor	1.390

Table 6.16: Median performance with queryone on small datasets without attributes, compression on the fly^{xx}

Compression type	median
originalcborreplace	0.980
originalcbor	1.000
originalcborzlib	1.005
originalzlib	1.005
originalreplace	1.015
originalcborreplacezlib	1.020
dracocbor	1.390
dracozlib	1.390
draco	1.410
dracocborzlib	1.425
dracoreplace	1.495
dracocborreplace	1.625
dracocborreplacezlib	1.625

6.5.2 Querying all features

Compression in advance

When querying all city objects, the whole dataset is traversed but the compression is performed on the complete file in the end. With large datasets, using Draco is beneficial compared to compression types that do not use it. This is likely due to Draco geometries being queried in a different way, as

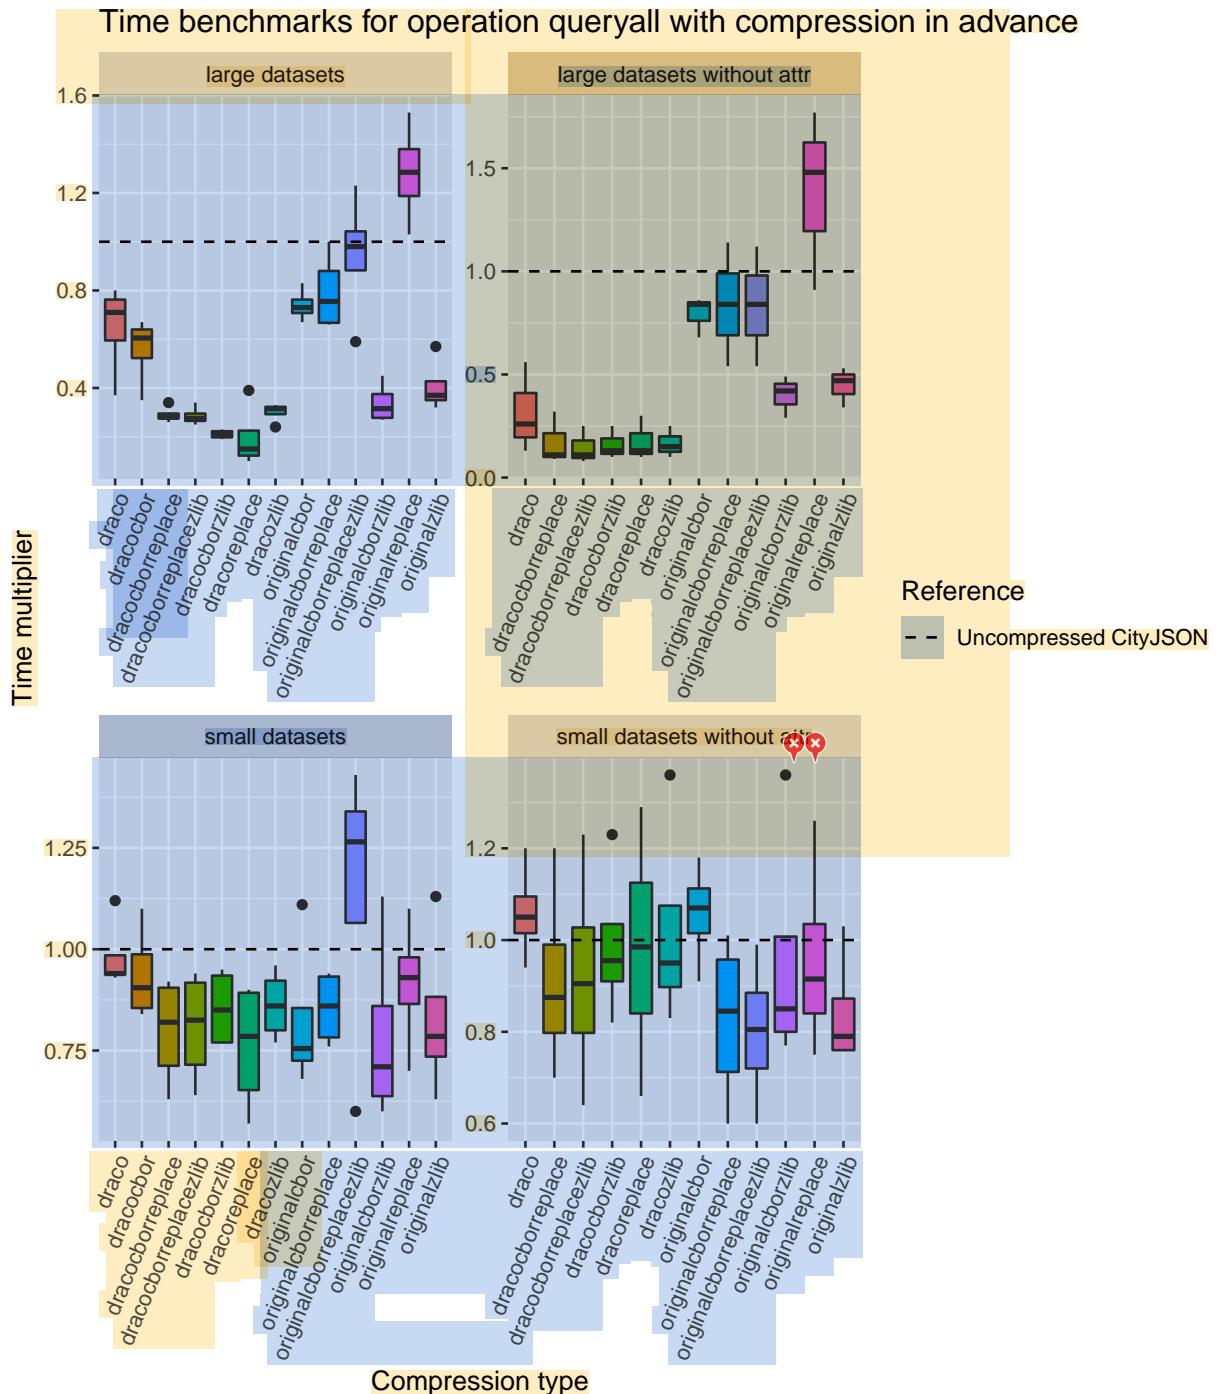


Figure 6.8: Boxplots of time performance of (in advance) compressed datasets on querying all features compared to uncompressed CityJSON.

explained in Section 5.2.2. However, original-cbor-zlib and original-zlib are also good, and are better than using draco alone or draco-cbor with datasets that have attributes.

As for the small datasets, Draco again provides less added benefits as is generally the case when datasets are compressed beforehand. There is no clear pattern to be noted when comparing compression types to each other. But as seen before as well, using cbor, zlib, or combinations of these prove to do well.

Compression in advance can improve the querying time of CityJSON datasets on all city objects to 15%, 11%, 71%, or 79% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.17: Median performance with queryall on large datasets, compression in advance

Compression type	median
dracoreplace	0.150
dracocborzlib	0.210
dracocborreplacezlib	0.275
dracocborreplace	0.280
dracoziplib	0.315
originalborzlib	0.315
originalzlib	0.370
dracocbor	0.605
draco	0.710
originalcbor	0.730
originalcborreplace	0.755
originalcborreplacezlib	0.980
originalreplace	1.285

Table 6.18: Median performance with queryall on large datasets without attributes, compression in advance

Compression type	median
dracocborreplace	0.11
dracocborreplacezlib	0.11
dracocborzlib	0.13
dracoreplace	0.13
dracoziplib	0.15
draco	0.26
originalborzlib	0.42
originalzlib	0.47
originalcbor	0.84
originalcborreplace	0.84
originalcborreplacezlib	0.84
originalreplace	1.48

Table 6.19: Median performance with queryall on small datasets, compression in advance

Compression type	median
originalborzlib	0.710
originalcbor	0.755
dracoreplace	0.785
originalzlib	0.785
dracocborreplace	0.820
dracocborreplacezlib	0.825
dracocborzlib	0.850
dracoziplib	0.860
originalcborreplace	0.860
dracocbor	0.905
originalreplace	0.930
draco	0.940
originalcborreplacezlib	1.265

Table 6.20: Median performance with queryall on small datasets without attributes, compression in advance

Compression type	median
originalzlib	0.790
originalcborreplacezlib	0.805
originalcborreplace	0.845
originalborzlib	0.850
dracocborreplace	0.875
dracocborreplacezlib	0.905
originalreplace	0.915
dracoziplib	0.950
dracocborzlib	0.955
dracoreplace	0.985
draco	1.050
originalcbor	1.070

6 Benchmarking results and analysis

Compression on the fly

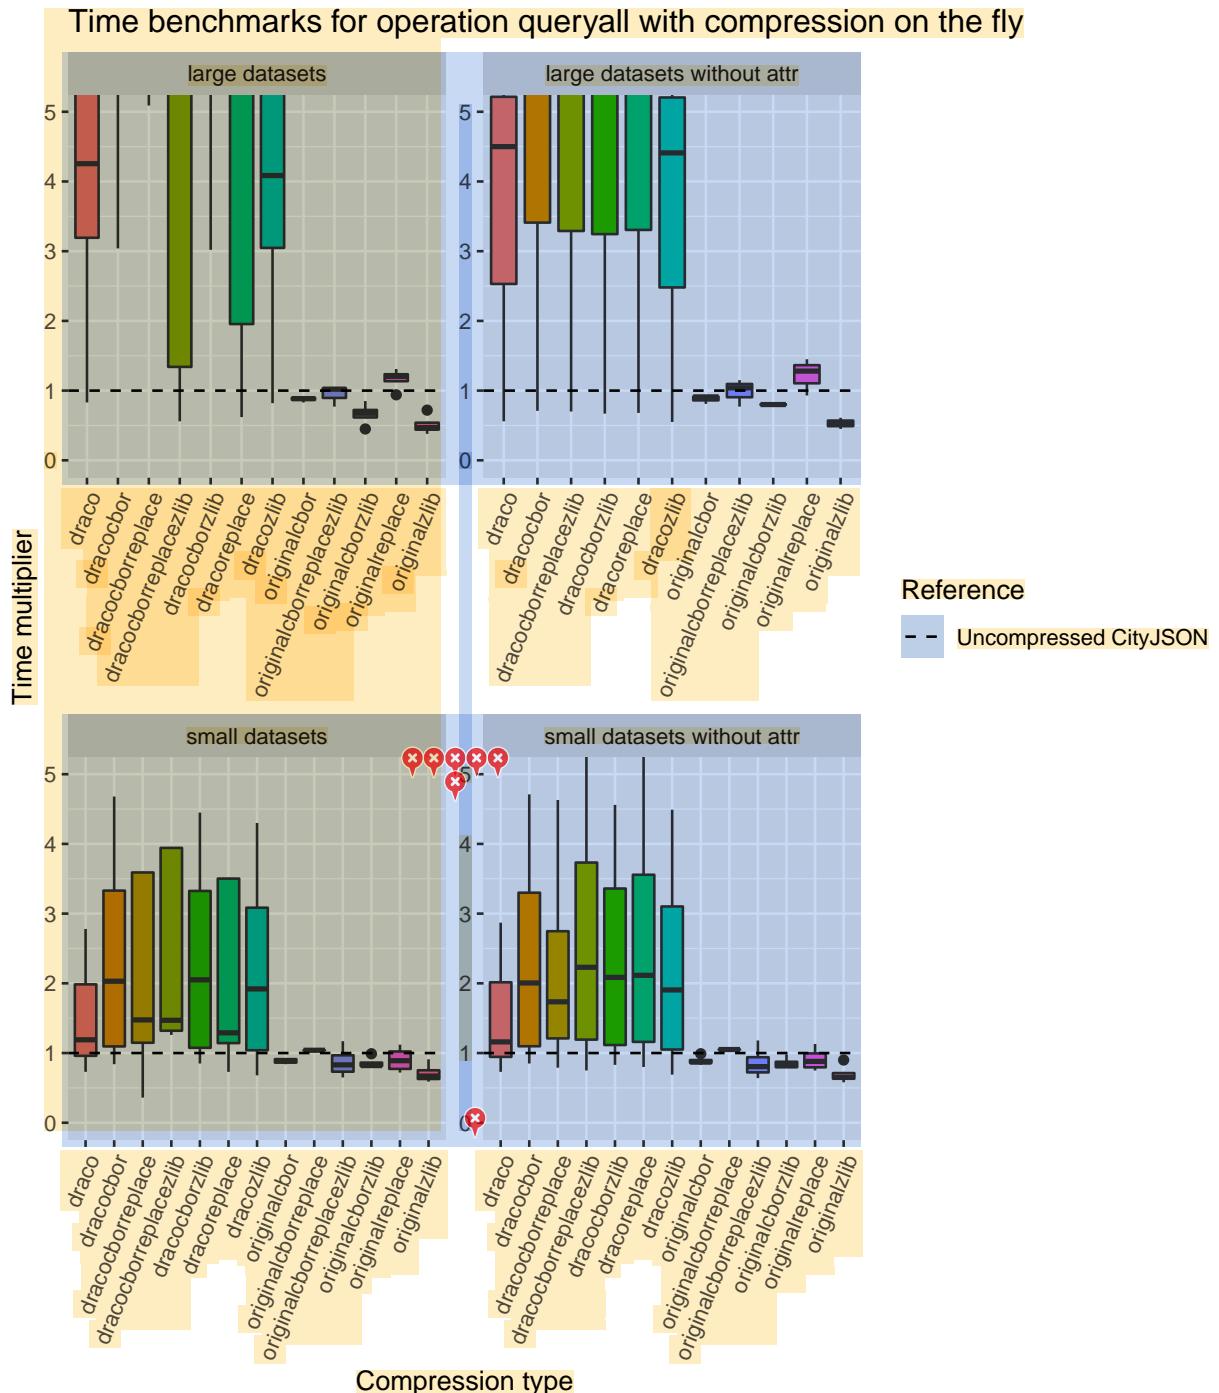


Figure 6.9: Boxplots of time performance of (on the fly) compressed datasets on querying all features compared to uncompressed CityJSON.

The graphs are limited by a time multiplier of 5. Again, using Draco is not a good idea when compressing datasets on the fly because of the relatively high compression time. Original-replace is unsuitable with large datasets, and original-cbor-replace-zlib also does not do very well. Original-zlib and original-cbor-zlib are the best ones, both with and without attributes. As for the small datasets these two are also good, but using the replace method is now not noticeably worse.

On the fly compression can improve the querying time of CityJSON datasets on all city objects to 47%, 53%, 67%, or 65% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.21: Median performance with queryall on large datasets, compression on the fly

Compression type	median
originalzlib	0.470
originalcborzlib	0.675
originalcbor	0.885
originalcborreplacedlib	1.020
originalreplace	1.205
dracozlib	4.085
draco	4.255
dracoreplace	6.125
dracocborreplacedlib	6.340
dracocborzlib	7.465
dracocbor	9.835
dracocborreplace	13.260

Table 6.22: Median performance with queryall on large datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.53
originalcborzlib	0.80
originalcbor	0.91
originalcborreplacedlib	1.04
originalreplace	1.28
dracozlib	4.41
draco	4.50
dracocborzlib	5.82
dracocborreplacedlib	5.88
dracoreplace	5.93
dracocbor	6.11

Table 6.23: Median performance with queryall on small datasets, compression on the fly

Compression type	median
originalzlib	0.670
originalcborzlib	0.820
originalcborreplacedlib	0.830
originalcbor	0.890
originalreplace	0.890
originalcborreplacedlib	1.040
draco	1.190
dracoreplace	1.290
dracocborreplacedlib	1.470
dracocborreplace	1.475
dracozlib	1.920
dracocbor	2.030
dracocborzlib	2.050

Table 6.24: Median performance with queryall on small datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.650
originalcborreplacedlib	0.805
originalcborzlib	0.820
originalcbor	0.870
originalreplace	0.880
originalcborreplacedlib	1.050
draco	1.160
dracocborreplace	1.735
dracozlib	1.905
dracocbor	2.005
dracocborzlib	2.085
dracoreplace	2.115
dracocborreplacedlib	2.230

6.6 Spatial analysis

6.6.1 Buffering one feature

Compression in advance

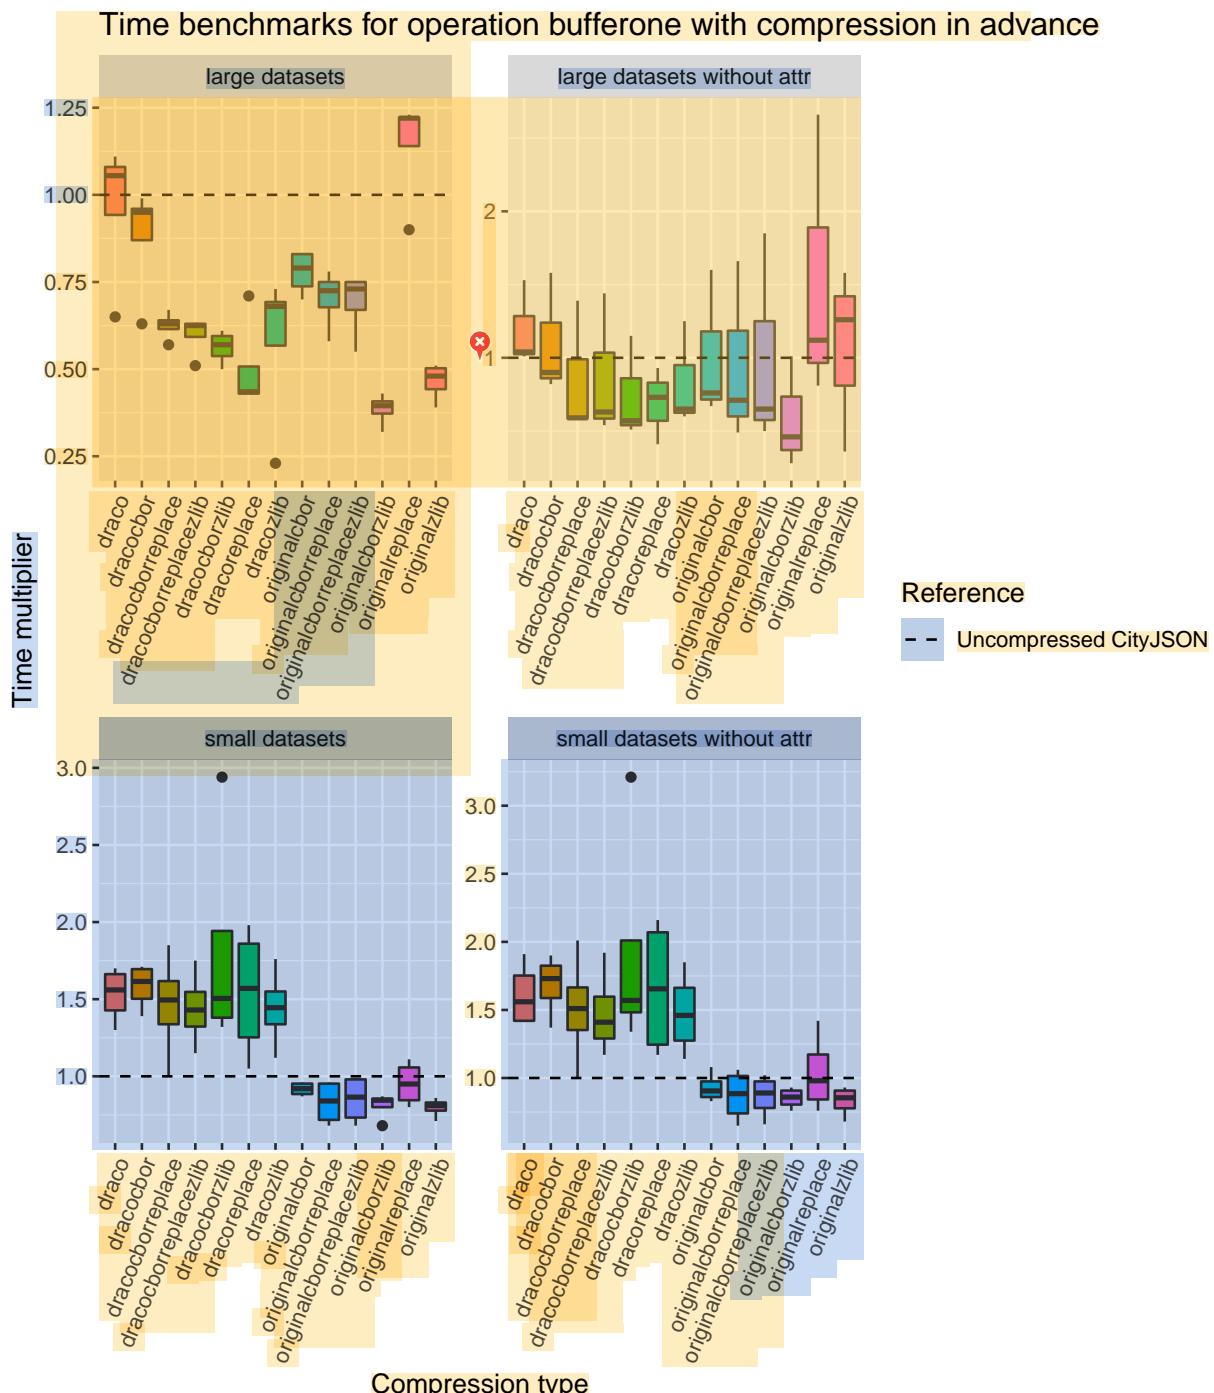


Figure 6.10: Boxplots of time performance of (in advance) compressed datasets on buffering one feature compared to uncompressed CityJSON.

Draco geometries are handled in a different way than original geometries as explained in Section 5.2.3. Despite that, for the large datasets it is not easy to see whether it is better to use it or not. There are also no clear rends to be seen, except for that original-replace is bad both with and without attributes, and original-cbor-zlib is the best option for both of these dataset types. In general, most compression types work well with large datasets with attributes but not as good when attributes are not included. In the latter case both attribute and JSON compression techniques have less data to compress, which might make them less effective. Regarding the small datasets, the use of Draco gives worse results than not using it due to overhead. The other compression types general do well as compared to the baseline, except for original-replace.

Compression in advance can improve the buffering time of CityJSON datasets on one city object to 39%, 46%, 81%, or 85% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.25: Median performance with bufferone on large datasets, compression in advance

Compression type	median
originalcborZlib	0.395
dracoreplace	0.435
originalzlib	0.480
dracocborZlib	0.570
dracocborreplaceZlib	0.625
dracocborreplace	0.630
dracoZlib	0.680
originalcborreplace	0.725
originalcborreplaceZlib	0.730
originalcbor	0.790
dracocbor	0.950
draco	1.055
originalreplace	1.220

Table 6.26: Median performance with bufferone on large datasets without attributes, compression in advance

Compression type	median
originalcborZlib	0.46
dracocborZlib	0.57
dracocborreplace	0.59
dracocborreplaceZlib	0.63
dracoZlib	0.65
originalcborreplaceZlib	0.65
originalcborreplace	0.71
dracoreplace	0.73
originalcbor	0.76
dracocbor	0.90
draco	1.04
originalreplace	1.12
originalzlib	1.26

Table 6.27: Median performance with bufferone on small datasets, compression in advance

Compression type	median
originalzlib	0.810
originalcborreplace	0.840
originalcborZlib	0.845
originalcborreplaceZlib	0.865
originalcbor	0.920
originalreplace	0.950
dracocborreplaceZlib	1.430
dracoZlib	1.445
dracocborreplace	1.495
dracocborZlib	1.505
draco	1.560
dracoreplace	1.570
dracocbor	1.615

Table 6.28: Median performance with bufferone on small datasets without attributes, compression in advance

Compression type	median
originalzlib	0.855
originalcborZlib	0.860
originalcborreplace	0.885
originalcborreplaceZlib	0.890
originalcbor	0.905
originalreplace	0.980
dracocborreplaceZlib	1.410
dracoZlib	1.460
dracocborreplace	1.510
draco	1.560
dracocborZlib	1.570
dracoreplace	1.655
dracocbor	1.730

Compression on the fly

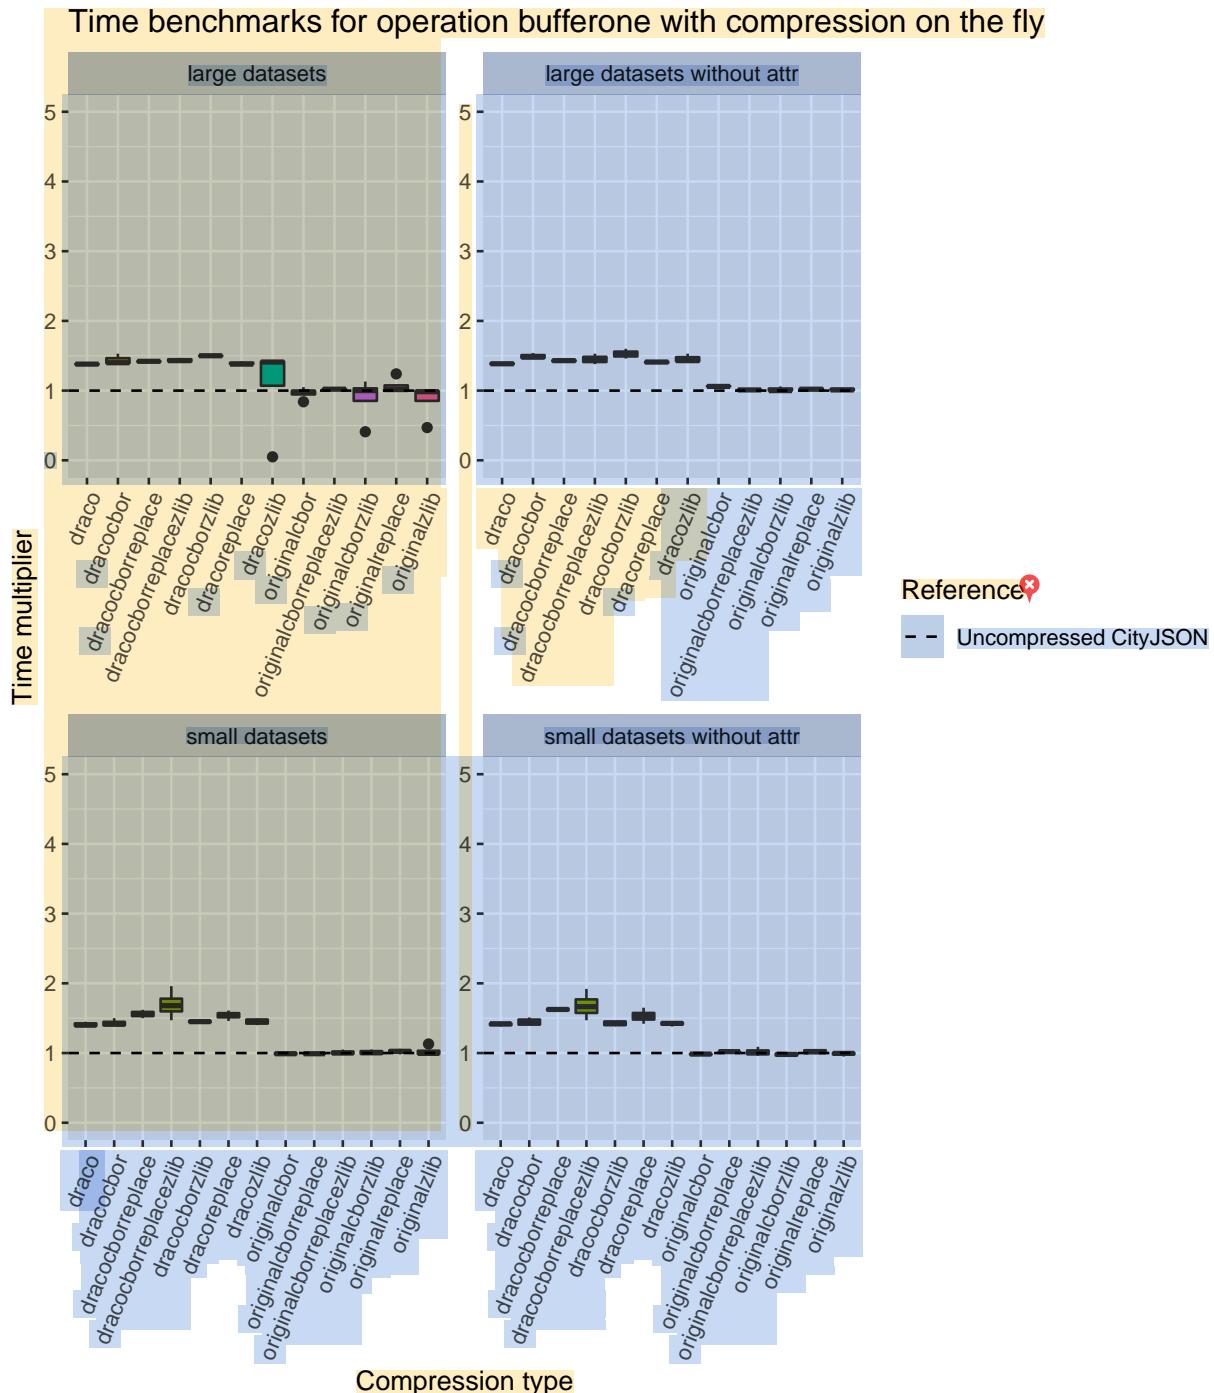


Figure 6.11: Boxplots of time performance of (on the fly) compressed datasets on buffering one feature compared to uncompressed CityJSON.

As opposed to compressing in advance, it is almost never worth it to compress buffer data on the fly since the medians of the boxplots are above the baseline or very slightly underneath it. This means that the added compression time is not outweighed by the faster transmission of the data. Especially when using Draco this is true. When using a compression type that retains original geometries, the results can be neutral so it is at least still fine to implement such kind of compression.

On the fly compression can improve the buffering time of CityJSON datasets on one city object to 98%, 100%, 99%, or 98% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.29: Median performance with bufferone on large datasets, compression on the fly

Compression type	median
originalcbor	0.980
originalzlib	0.985
originalcborzlib	1.000
originalreplace	1.015
originalcborreplacedzlib	1.020
draco	1.380
dracoreplace	1.380
dracocbor	1.400
dracozlib	1.410
dracocborreplace	1.420
dracocborreplacedzlib	1.430
dracocborzlib	1.500

Table 6.30: Median performance with bufferone on large datasets without attributes, compression on the fly

Compression type	median
originalcborzlib	1.000
originalcborreplacedzlib	1.010
originalzlib	1.010
originalreplace	1.020
originalcbor	1.060
draco	1.385
dracoreplace	1.410
dracocborreplace	1.430
dracozlib	1.430
dracocborreplacedzlib	1.440
dracocbor	1.480
dracocborzlib	1.520

Table 6.31: Median performance with bufferone on small datasets, compression on the fly

Compression type	median
originalcbor	0.990
originalcborreplacedzlib	0.990
originalzlib	0.995
originalcborreplacedzlib	1.000
originalcborzlib	1.005
originalreplace	1.025
draco	1.400
dracocbor	1.415
dracocborzlib	1.450
dracozlib	1.455
dracoreplace	1.545
dracocborreplace	1.560
dracocborreplacedzlib	1.680

Table 6.32: Median performance with bufferone on small datasets without attributes, compression on the fly

Compression type	median
originalcborzlib	0.980
originalcbor	0.985
originalzlib	0.995
originalcborreplacedzlib	1.005
originalcborreplacedzlib	1.005
originalcborreplacedzlib	1.005
originalreplace	1.020
draco	1.420
dracocborzlib	1.425
dracozlib	1.425
dracocbor	1.440
dracoreplace	1.525
dracocborreplace	1.625
dracocborreplacedzlib	1.665

6.6.2 Buffering all features

Compression in advance

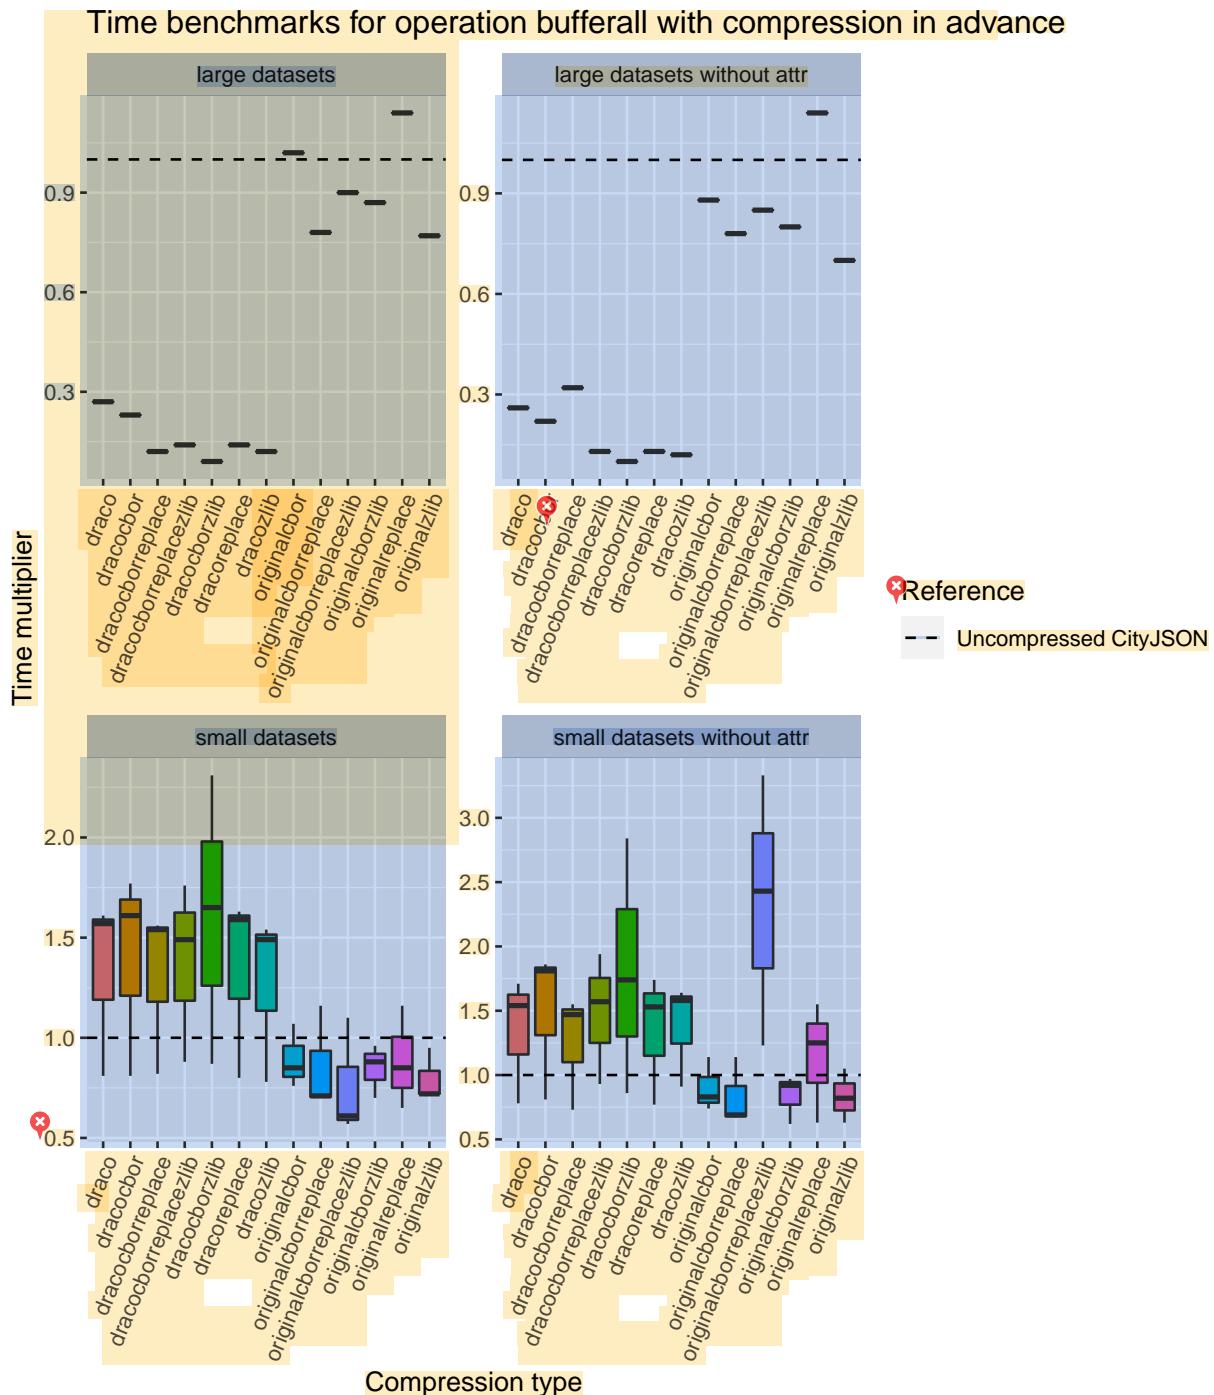


Figure 6.12: Boxplots of time performance of (in advance) compressed datasets on buffering all features compared to uncompressed CityJSON.

When buffering all features in the client, using Draco is clearly a good idea for the large datasets. Draco geometries are buffered in a different way (see Section 5.2.3), and this shows a more apparent advantage when all geometries are buffered instead of one. Compression types without it generally

still improve the performance in comparison to the baseline, but not as much. For small datasets these perform better however, which is also seen with the other use cases with small datasets. Draco is again unsuitable to use with those. But ultimately, it is questionable whether or not it is wise to use compression at all for this use case, since the boxplots show that most compression types that have a median below the baseline still have their boxplots touching or above the baseline.

Compression in advance can improve the buffering time of CityJSON datasets on all city objects to 9%, 10%, 61%, or 69% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.33: Median performance with bufferall on large datasets, compression in advance

Compression type	median
dracocborzlib	0.09
dracocborreplace	0.12
dracozlib	0.12
dracocborreplacezlib	0.14
dracoreplace	0.14
dracocbor	0.23
draco	0.27
originalzlib	0.77
originalcborreplace	0.78
originalcborzlib	0.87
originalcborreplacezlib	0.90
originalcbor	1.02
originalreplace	1.14

Table 6.34: Median performance with bufferall on large datasets without attributes, compression in advance

Compression type	median
dracocborzlib	0.10
dracozlib	0.12
dracocborreplacezlib	0.13
dracoreplace	0.13
dracocbor	0.22
draco	0.26
dracocborreplace	0.32
originalzlib	0.70
originalcborreplace	0.78
originalcborzlib	0.80
originalcborreplacezlib	0.85
originalcbor	0.88
originalreplace	1.14

Table 6.35: Median performance with bufferall on small datasets, compression in advance

Compression type	median
originalcborreplacezlib	0.61
originalcborreplace	0.71
originalzlib	0.72
originalcbor	0.85
originalreplace	0.85
originalcborzlib	0.88
dracocborreplacezlib	1.49
dracozlib	1.49
dracocborreplace	1.54
draco	1.57
dracoreplace	1.59
dracocbor	1.61
dracocborzlib	1.65

Table 6.36: Median performance with bufferall on small datasets without attributes, compression in advance

Compression type	median
originalcborreplace	0.69
originalzlib	0.82
originalcbor	0.83
originalcborzlib	0.92
originalreplace	1.25
dracocborreplace	1.47
dracoreplace	1.53
draco	1.54
dracocborreplacezlib	1.57
dracozlib	1.58
dracocborzlib	1.74
dracocbor	1.81
originalcborreplacezlib	2.43

Compression on the fly

Time benchmarks for operation bufferall with compression on the fly

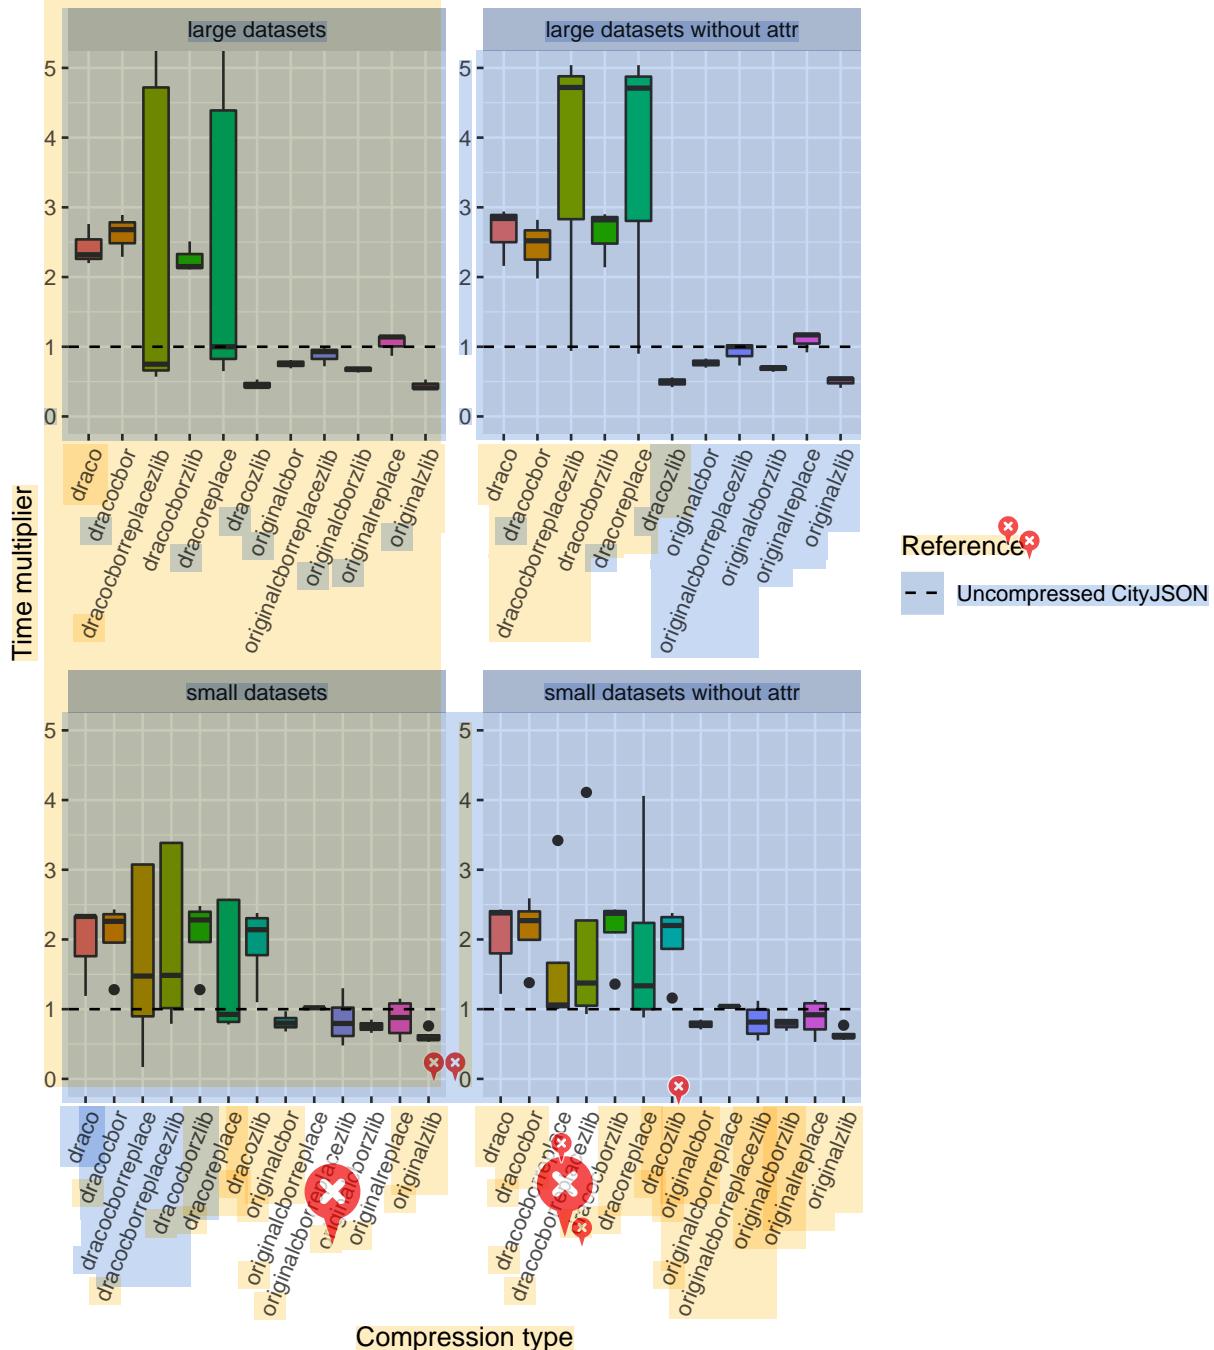


Figure 6.13: Boxplots of time performance of (on the fly) compressed datasets on buffering all features compared to uncompressed CityJSON.

Using Draco compression on the fly is not useful. Other compression types on the other hand do provide benefits, except when the replace method is used. This is true for all dataset types, and there are no other interesting patterns to see in the graphs.

On the fly compression can improve the buffering time of CityJSON datasets on all city objects to

41%, 50%, 57%, or 60% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.37: Median performance with bufferall on large datasets, compression on the fly

Compression type	median
originalzlib	0.41
dracozlib	0.43
originalcborzlib	0.68
dracocborreplacezlib	0.75
originalcbor	0.76
originalcborreplacezlib	0.93
dracoreplace	1.00
originalreplace	1.14
dracocborzlib	2.15
draco	2.32
dracocbor	2.68

Table 6.38: Median performance with bufferall on large datasets without attributes, compression on the fly

Compression type	median
dracozlib	0.50
originalzlib	0.54
originalcborzlib	0.70
originalcbor	0.77
originalcborreplacezlib	1.00
originalreplace	1.17
dracocbor	2.52
dracocborzlib	2.82
draco	2.84
dracoreplace	4.71
dracocborreplacezlib	4.72

Table 6.39: Median performance with bufferall on small datasets, compression on the fly

Compression type	median
originalzlib	0.575
originalcborzlib	0.750
originalcborreplacezlib	0.795
originalcbor	0.800
originalreplace	0.880
dracoreplace	0.925
originalcborreplace	1.020
dracocborreplace	1.475
dracocborreplacezlib	1.485
dracozlib	2.140
dracocbor	2.260
dracocborzlib	2.280
draco	2.330

Table 6.40: Median performance with bufferall on small datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.600
originalcbor	0.785
originalcborzlib	0.800
originalcborreplacezlib	0.815
originalreplace	0.920
originalcborreplace	1.040
dracocborreplace	1.060
dracoreplace	1.335
dracocborreplacezlib	1.375
dracozlib	2.200
dracocbor	2.270
dracocborzlib	2.375
draco	2.380

6.7 Editing (attributes)

6.7.1 Editing attributes one feature

Compression in advance

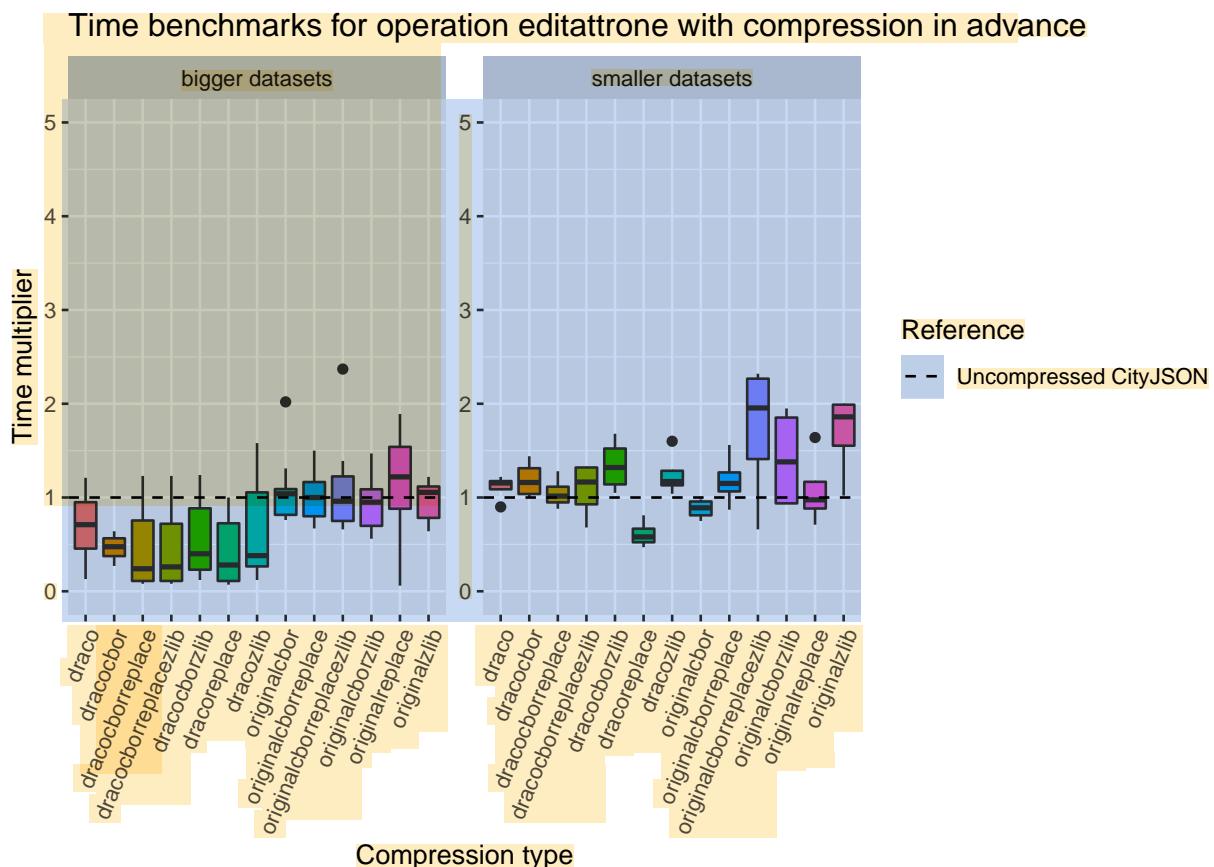


Figure 6.14: Boxplots of time performance of (in advance) compressed datasets on editing attributes of one feature compared to uncompressed CityJSON.

For this use case the compression time becomes relevant, as opposed to the other use cases with compression in advance. Despite that, the results seem unexpectedly good. When using Draco compression in combination with large datasets, the performance is rather well both compared to the baseline and to compression types without Draco. That is because the Draco geometries are not touched as they are not edited, so they also do not need to be compressed again and they make for small file sizes, benefitting the transmission time. The addition of other compression techniques is beneficial as well since using Draco only does not compare very well to the other Draco compression types. When not using Draco original-cbor, original-cbor-zlib, and original-zlib do well.

With the small datasets most compression types are not suitable. Draco-replace does work well which could be explained by the different way in which attributes can be edited with the replace method and the advantage of Draco as mentioned in the paragraph above, but other compression types using this method do not do particularly well. Original-cbor is also slightly better than the baseline.

Compression in advance can^{fix} improve the attribute editing time of CityJSON datasets on one city object to 10% or 58% (when respectively using large datasets and small datasets) of the time that it

takes to perform the same operation with uncompressed CityJSON.

Table 6.41: Median performance with editattroneTable 6.42: Median performance with editattrone
on large datasets, compression in advance on small datasets, compression in advance

Compression type	median
dracoreplace	0.100
dracocborreplace	0.135
dracocborreplacezlib	0.140
dracocborzlib	0.335
dracoilib	0.360
dracocbor	0.475
originalcborzlib	0.600
draco	0.655
originalzlib	0.730
originalcbor	0.810
originalcborreplace	1.015
originalcborreplacezlib	1.135
originalreplace	1.330

Compression type	median
dracoreplace	0.580
originalcbor	0.890
originalreplace	0.975
dracocborreplace	1.015
originalcborreplace	1.150
draco	1.155
dracocbor	1.160
dracocborreplacezlib	1.165
dracoilib	1.170
dracocborzlib	1.320
originalcborzlib	1.380
originalzlib	1.860
originalcborreplacezlib	1.955

Compression on the fly

Time benchmarks for operation editattrone with compression on the fly

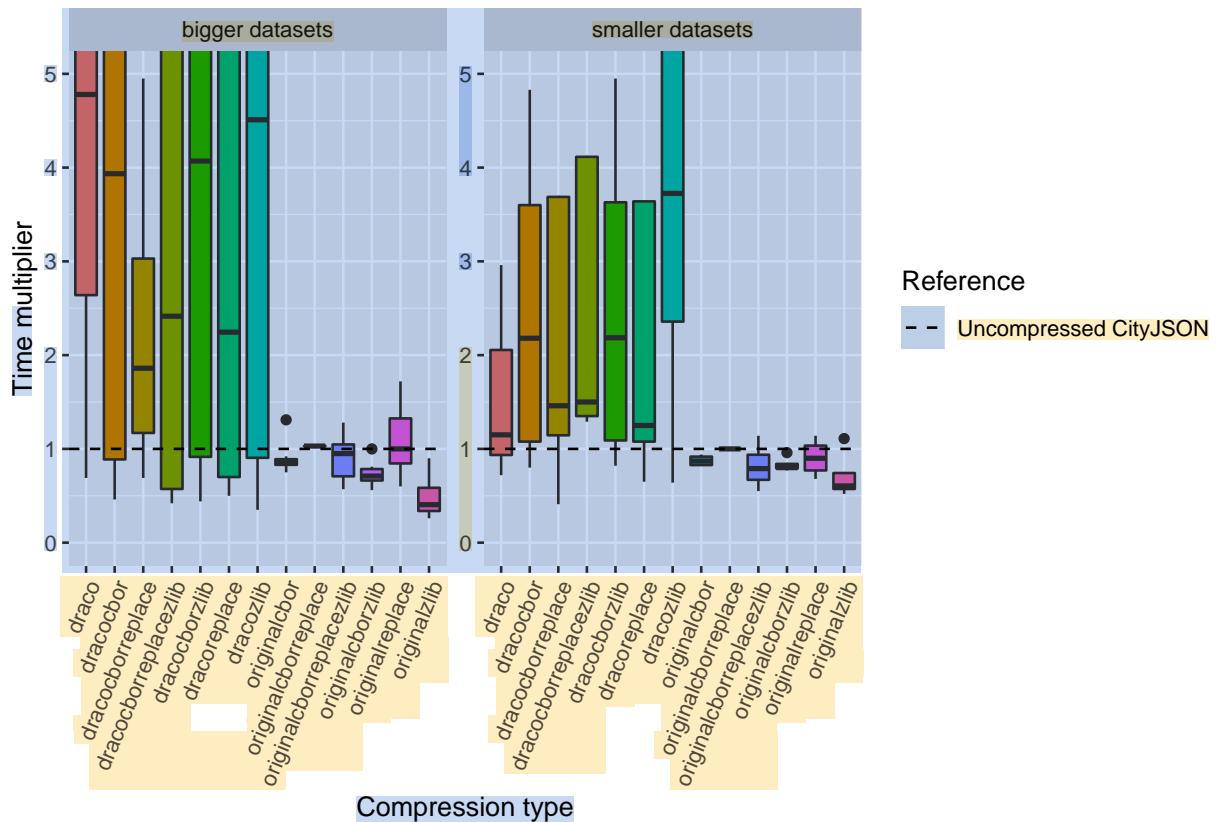


Figure 6.15: Boxplots of time performance of (on the fly) compressed datasets on editing attributes of one feature compared to uncompressed CityJSON.

6 Benchmarking results and analysis

The data is edited before being compressed with this use case. Draco is again not suitable to use when compressing data on the fly with neither large nor small datasets. As for the other compression types, original-zlib and original-cbor-zlib are beneficial for performance with the larger ones, but with the smaller ones only original-cbor and original-cbor-zlib have a (slight) advantage in every case. All other types stretch above the baseline.

On-the-fly compression can improve the attribute editing time of CityJSON datasets on one city object to 33% or 60% (when respectively using large datasets and small datasets) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.44: Median performance with editattrone

Table 6.43: Median performance with editattrone on small datasets, compression on the fly

on large datasets, compression on the fly

Compression type	median
originalzlib	0.330
dracocborreplacezlib	0.420
originalcborzlib	0.570
dracoreplace	0.760
originalcbor	0.870
originalcborreplacezlib	0.960
originalreplace	1.075
dracocbor	4.600
draco	4.990
draco	5.480
dracocborzlib	5.610

Compression type	median
originalzlib	0.605
originalcborreplacezlib	0.790
originalcborzlib	0.800
originalcbor	0.870
originalreplace	0.900
originalcborreplace	1.000
draco	1.150
dracoreplace	1.250
dracocborreplace	1.460
dracocborreplacezlib	1.500
dracocbor	2.180
dracocborzlib	2.185
draco	3.725

6.7.2 Editing attributes all features

Compression in advance

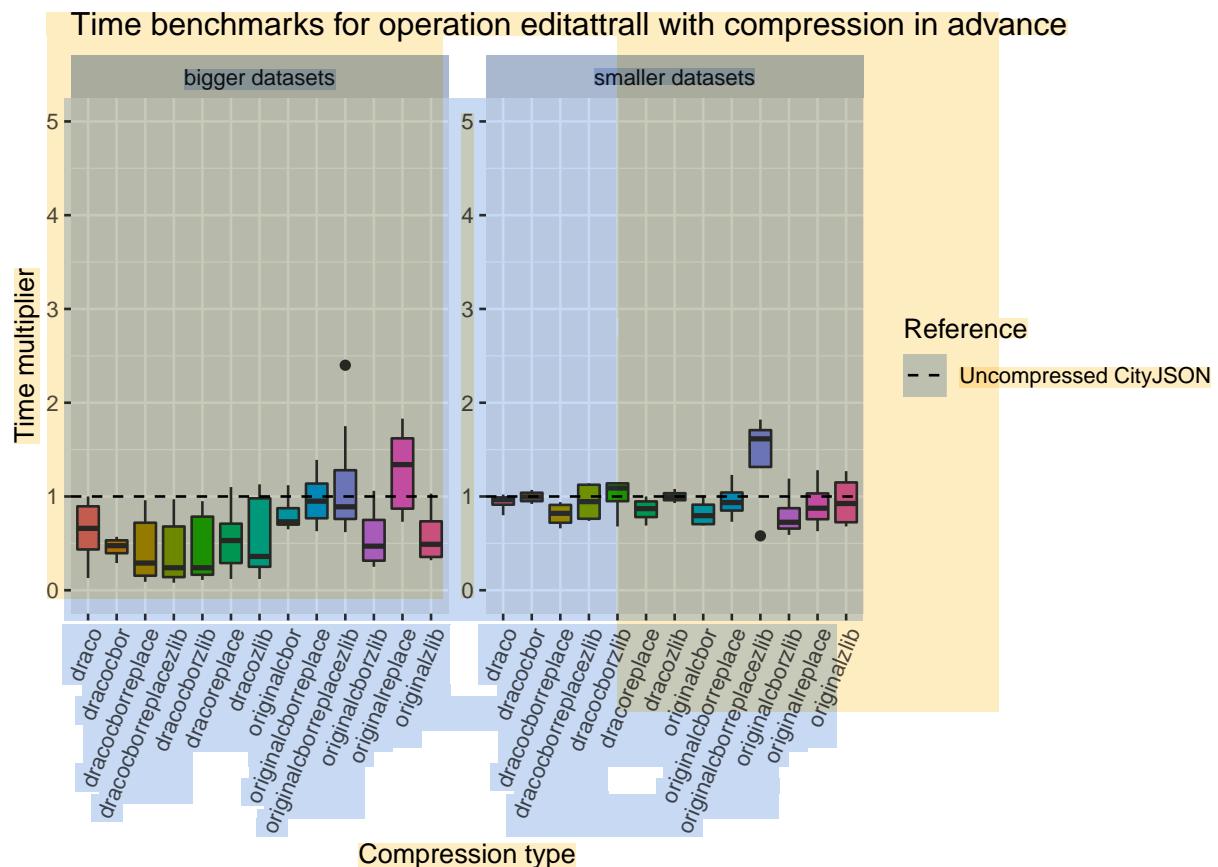


Figure 6.16: Boxplots of time performance of (in advance) compressed datasets on editing attributes of all features compared to uncompressed CityJSON.

When editing attributes of a large dataset it is not a good idea to use the replace method, regardless of the seemingly more efficient way to edit attributes with it (see Section 5.2.4). But apart from that, compression is beneficial for this use case. Using Draco works well, which would be because it is untouched as only attributes are edited and Draco geometries are small, benefitting the transmission time. Original-cbor, original-cbor-zlib, and original-zlib also perform well.

With the small datasets most compression types do not always perform well—only draco-cbor-replace has a boxplot that remains under the baseline. This means that the overhead when compressing the data again does not outweigh the benefits of compression as much as with large datasets.

Compression in advance can improve the attribute editing time of CityJSON datasets on all city objects to 14% or 72% (when respectively using large datasets and small datasets) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.45: Median performance with editattrall
on large datasets, compression in advance

Compression type	median
dracocborreplace	0.140
dracocborreplacezlib	0.170
dracocborzlib	0.205
originalcborzlib	0.300
draco zlib	0.350
originalzlib	0.355
dracoreplace	0.410
dracocbor	0.475
draco	0.630
originalcbor	0.715
originalborreplace	1.130
originalreplace	1.490
originalborreplacezlib	1.510

Table 6.46: Median performance with editattrall
on small datasets, compression in advance

Compression type	median
originalcborzlib	0.725
originalcbor	0.795
dracocborreplace	0.820
dracoreplace	0.870
originalreplace	0.875
originalzlib	0.925
originalborreplace	0.935
dracocborreplacezlib	0.945
draco	0.965
dracocbor	0.995
draco zlib	0.995
dracocborzlib	1.090
originalborreplacezlib	1.615

Compression on the fly

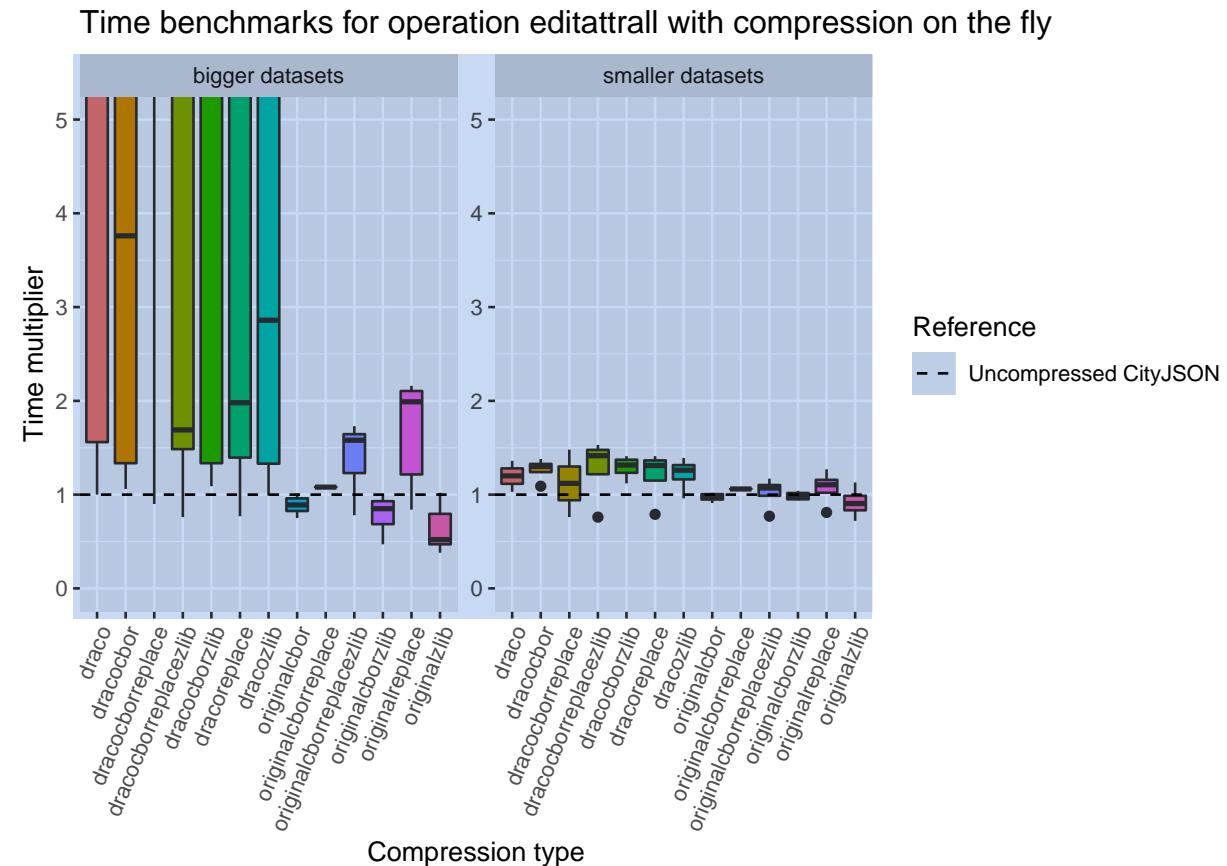


Figure 6.17: Boxplots of time performance of (on the fly) compressed datasets on editing attributes of all features compared to uncompressed CityJSON.

As is always the case with compression on the fly, the use of Draco compression is detrimental for performance. With the large datasets original-cbor, original-cbor-zlib, and original-zlib yield a proper performance gain, which would be the result of a good balance between file size and (de)compression time. As for the small datasets, the differences between the compression types are smaller and there is no clear best one to use as the better performing ones are around the baseline. Thus in that case you might be better off not using compression at all.

On the fly compression can improve the attribute editing time of CityJSON datasets on all city objects to 47% or 90% (when respectively using large datasets and small datasets) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.47: Median performance with editattrall
on large datasets, compression on the fly

Compression type	median
originalzlib	0.475
originalcborzlib	0.680
originalcbor	0.825
originalcborreplacedlib	1.660
originalreplace	2.105
dracoreplace	15.180
draco	16.910
dracocborreplacedlib	16.950
draco	18.210
dracocborreplace	27.220
dracocborzlib	28.735
dracocbor	32.745

Table 6.48: Median performance with editattrall
on small datasets, compression on the fly

Compression type	median
originalzlib	0.905
originalcbor	0.980
originalcborzlib	0.980
originalcborreplaced	1.060
originalcborreplacedlib	1.070
originalreplace	1.105
dracocborreplace	1.120
draco	1.200
draco	1.260
dracocbor	1.300
dracoreplace	1.310
dracocborzlib	1.315
dracocborreplacedlib	1.415

6.8 Editing (geometry)

6.8.1 Editing geometry one feature

Compression in advance

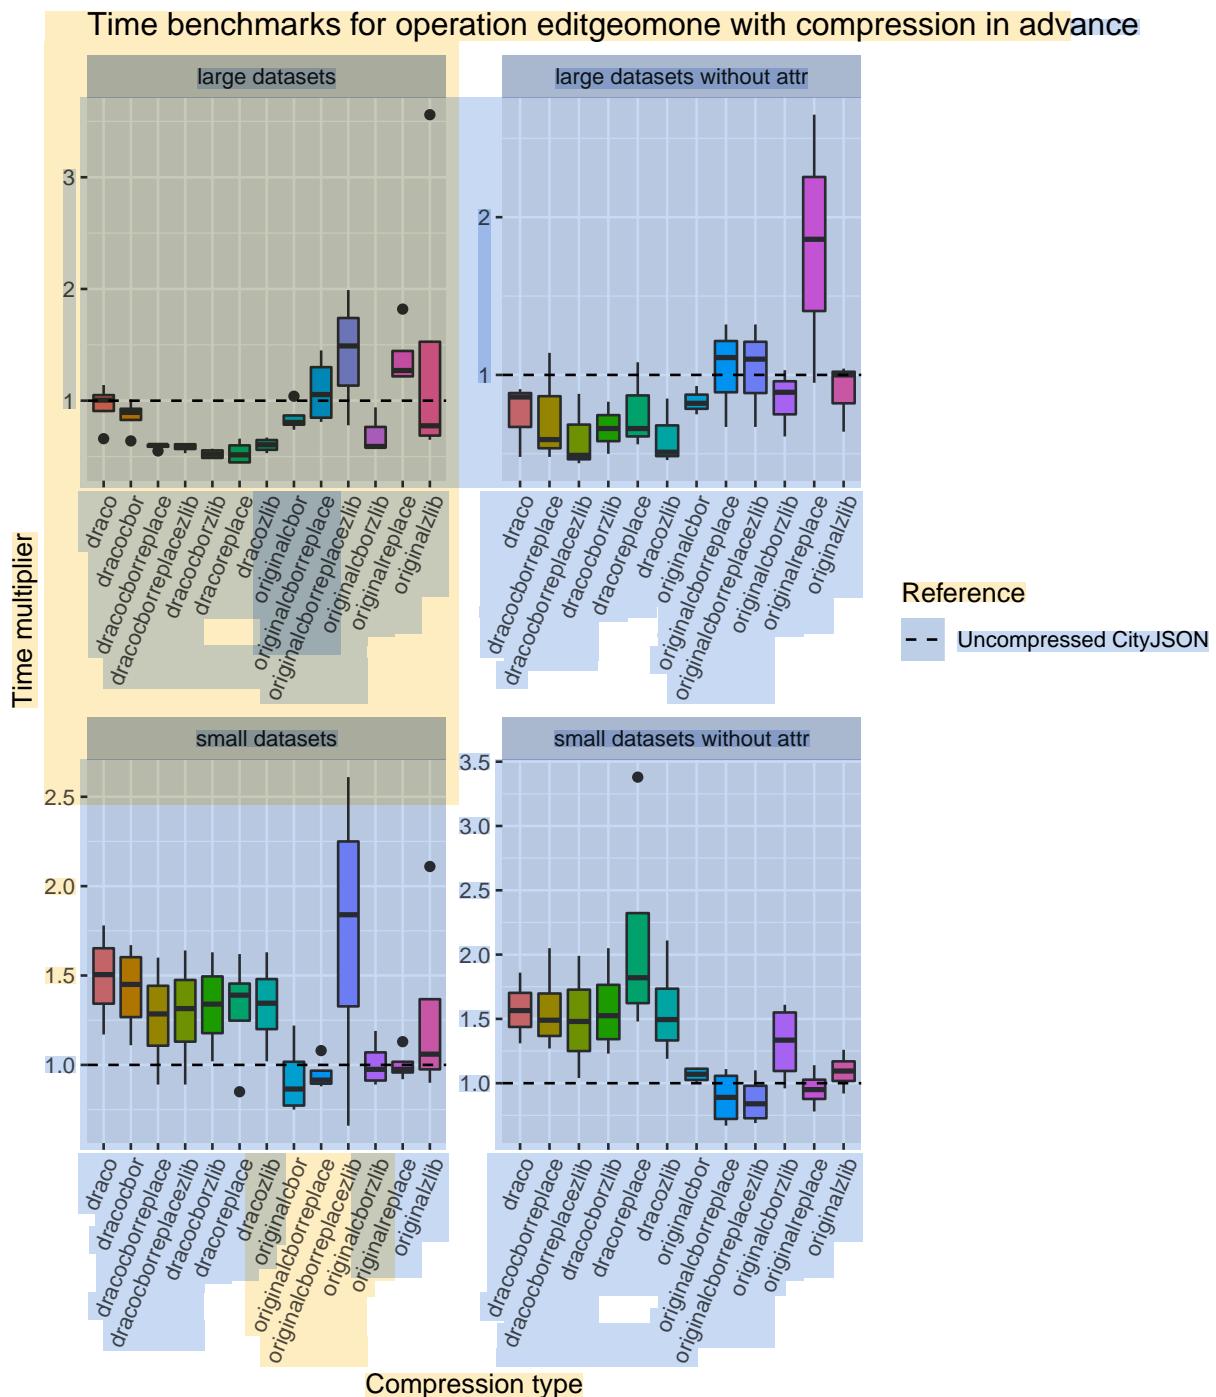


Figure 6.18: Boxplots of time performance of (in advance) compressed datasets on editing geometry of one feature compared to uncompressed CityJSON.

6 Benchmarking results and analysis

Draco is suitable for this use case when editing large datasets. This is despite Draco methods having high compression times as shown in Section 6.1. But the results in that Section were created by compressing CityJSON geometries into Draco geometries with OBJ conversion as intermediate step (see Section 5.4.5). In this case however the Draco geometries are edited directly as they are decoded in the client. It can then be compressed again without additional steps. The use of additional compression techniques is beneficial as well. If not using Draco, original-cbor and original-cbor-zlib are good choices. This is the same for datasets that do not contain attributes.

With the small datasets none of the compression types perform particularly well. Draco is now not beneficial which would be due to the overhead of Draco compression that is more significant with small data. Original-cbor has the lowest median with attributes, but is above the baseline with datasets that do not have attributes. In that case there is less data for CBOR to compress which could be the explaining factor.

Compression in advance can improve the geometry editing time of CityJSON datasets on one city object to 51%, 49%, 86%, or 84% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.49: Median performance with editgeomone on large datasets, compression in advance

Compression type	median
dracoreplace	0.515
dracocborzlib	0.520
originalcborzlib	0.590
dracocborreplacezlib	0.595
dracocborreplace	0.600
dracozlib	0.605
originalzlib	0.775
originalcbor	0.805
dracocbor	0.895
draco	1.005
originalcborreplace	1.055
originalreplace	1.270
originalcborreplacezlib	1.490

Table 6.50: Median performance with editgeomone on large datasets without attributes, compression in advance

Compression type	median
dracocborreplacezlib	0.49
dracozlib	0.51
dracocborreplace	0.59
dracocborzlib	0.66
dracoreplace	0.66
originalcbor	0.82
draco	0.86
originalcborzlib	0.89
originalzlib	1.00
originalcborreplacezlib	1.10
originalcborreplace	1.11
originalreplace	1.86

Table 6.51: Median performance with editgeomone on small datasets, compression in advance

Compression type	median
originalcbor	0.865
originalcborreplaced	0.915
originalreplace	0.975
originalcborzlib	0.975
originalzlib	1.060
dracocborreplace	1.285
dracocborreplacezlib	1.315
dracocborzlib	1.340
dracozb	1.345
dracoreplace	1.390
dracocbor	1.450
draco	1.505
originalcborreplacedzlib	1.840

Table 6.52: Median performance with editgeomone on small datasets without attributes, compression in advance

Compression type	median
originalcborreplacedzlib	0.840
originalcborreplaced	0.890
originalreplace	0.950
originalcbor	1.070
originalzlib	1.095
originalcborzlib	1.335
dracocborreplacezlib	1.480
dracocborreplace	1.490
dracozb	1.495
dracocborzlib	1.525
draco	1.565
dracoreplace	1.820

Compression on the fly

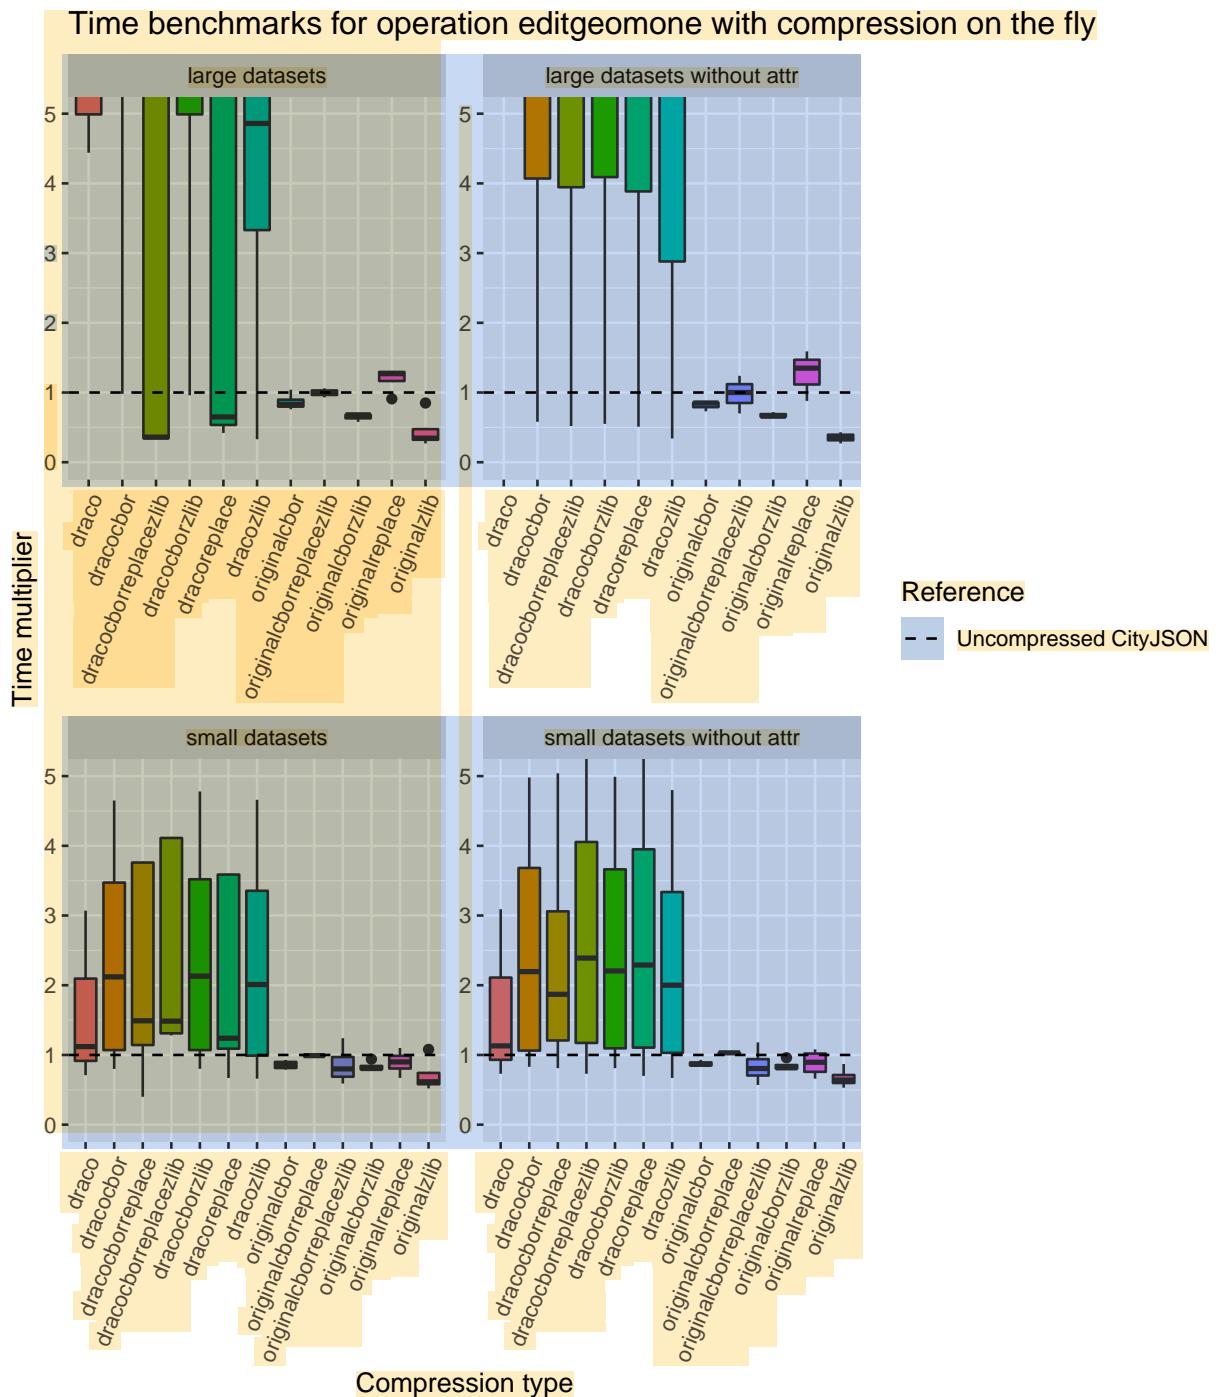


Figure 6.19: Boxplots of time performance of (on the fly) compressed datasets on editing geometry of one feature compared to uncompressed CityJSON.

Ignoring the compression types using Draco, it can be beneficial to use compression on the fly for this use case. With large datasets, original-cbor, original-zlib, and original-cbor-zlib do well. They are also fine to use with small datasets, but here the compression types have similar performances. There are no other notable observations to be made.

On the fly compression can improve the geometry editing time of CityJSON datasets on one city object to 34%, 36%, 61%, or 64% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.53: Median performance with editgeomone on large datasets, compression on the fly

Compression type	median
originalzlib	0.345
dracocborreplacezlib	0.360
dracoreplace	0.650
originalcborzlib	0.680
originalcbor	0.830
originalcborreplacezlib	1.000
originalreplace	1.270
dracozlib	4.860
draco	5.540
dracocborzlib	9.020
dracocbor	11.700

Table 6.54: Median performance with editgeomone on large datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.36
originalcborzlib	0.66
originalcbor	0.85
originalcborreplacezlib	1.00
originalreplace	1.35
draco	5.42
draco	6.60
dracoreplace	7.26
dracocborreplacezlib	7.37
dracocbor	7.56
dracocborzlib	7.63

Table 6.55: Median performance with editgeomone on small datasets, compression on the fly

Compression type	median
originalzlib	0.615
originalcborreplacezlib	0.800
originalcborzlib	0.805
originalcbor	0.855
originalreplace	0.900
originalcborreplace	0.990
draco	1.120
dracoreplace	1.240
dracocborreplacezlib	1.485
dracocborreplace	1.490
draco	2.010
dracocbor	2.120
dracocborzlib	2.130

Table 6.56: Median performance with editgeomone on small datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.640
originalcborreplacezlib	0.805
originalcborzlib	0.815
originalcbor	0.865
originalreplace	0.895
originalcborreplace	1.030
draco	1.130
dracocborreplace	1.870
draco	2.000
dracocbor	2.195
dracocborzlib	2.205
dracoreplace	2.290
dracocborreplacezlib	2.390

6.8.2 Editing geometry all features

Compression in advance

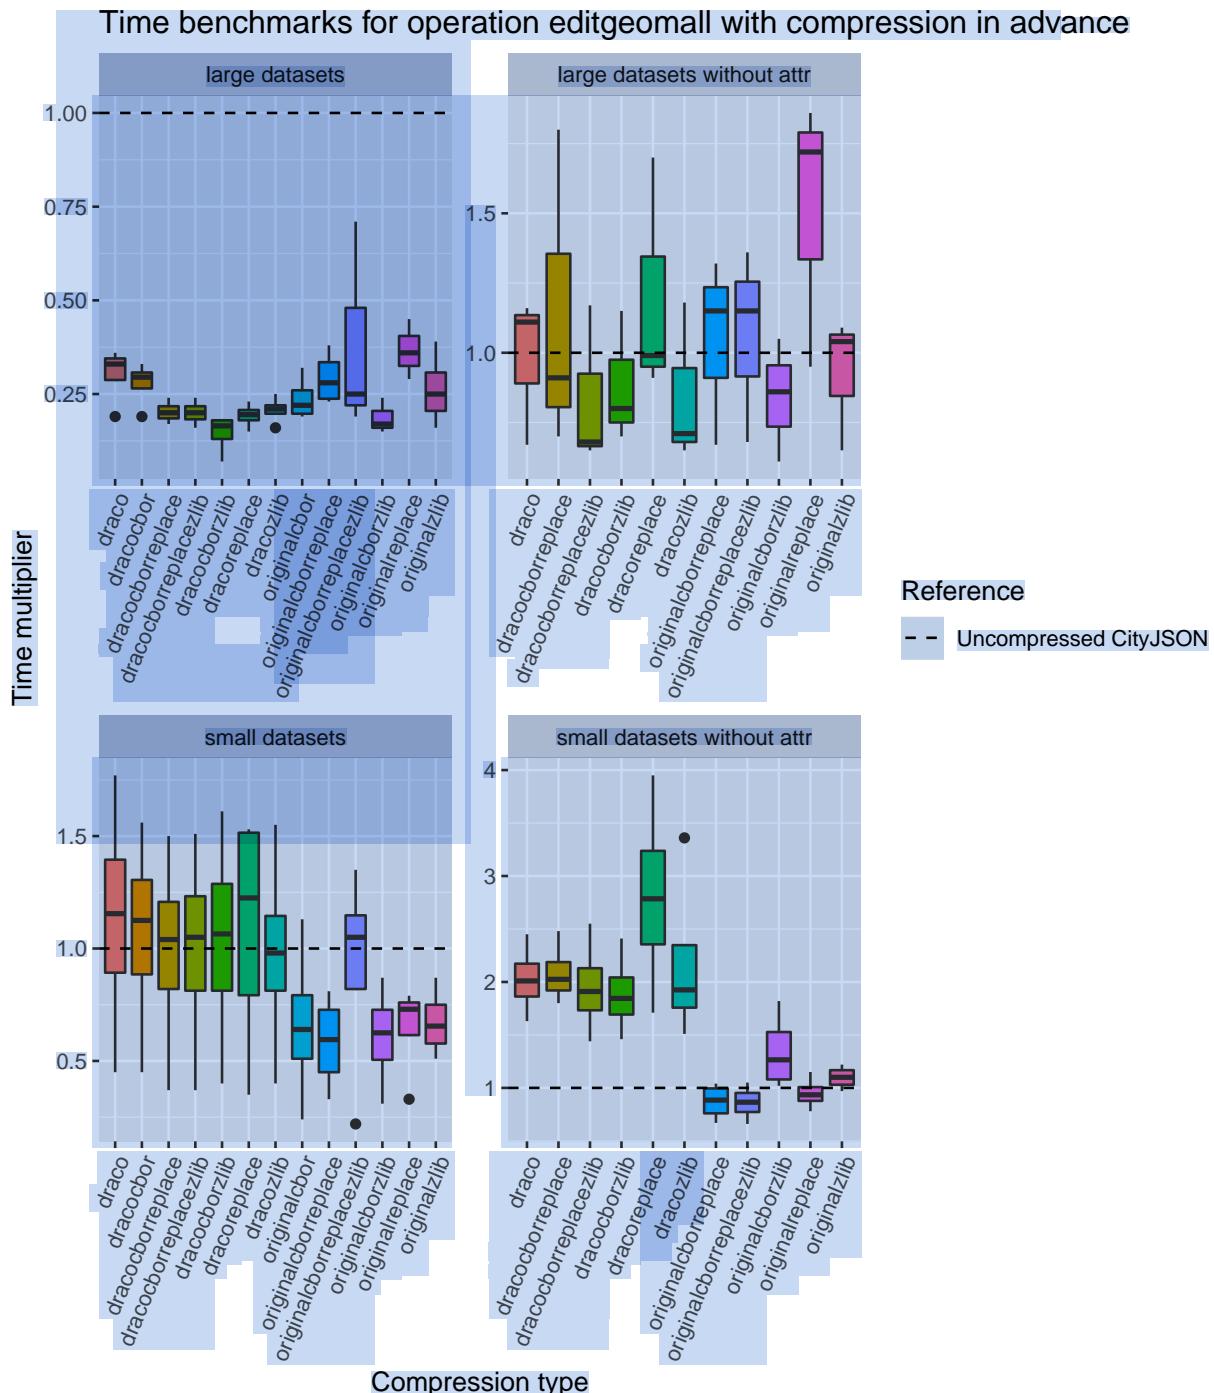


Figure 6.20: Boxplots of time performance of (in advance) compressed datasets on editing geometry of all features compared to uncompressed CityJSON.

With the large datasets that include attributes all compression types perform well. This is not seen when only editing the geometry of one city object, when performance was only good with Draco. It is surprising because now all city objects have to be traversed for compression types without Draco,

but with Draco all vertices already had to be iterated over, even when just wanting to edit one feature (see Section 5.2.5). The results are also quite different for large datasets without attributes, and generally worse—none of the compression types is good to use with every of these datasets as the boxplots indicate. This is surprising as well since there is less data to compress again after the editing in that case.

As for the small datasets, Draco compresses negatively influences the performance, presumably because of the overhead. Other compression types generally work well, except for draco-cbor-replace-zlib when attributes are included and original-cbor-zlib when they are not. The latter one tends to perform well in other use cases so it could be an anomaly. The former type includes the replace method that does not do well so far, but other compression types that use it do not do as bad. This can indicate that it is not worth it to compress the array of attributes with zlib (see Section 5.4.3).

Compression in advance can improve the geometry editing time of CityJSON datasets on all city objects to 16%, 68%, 59%, or 86% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.57: Median performance with editgeomall
on large datasets, compression in advance

Compression type	median
dracocborzlib	0.165
originalcborzlib	0.170
dracoreplace	0.195
dracocborreplace	0.200
dracocborreplacezlib	0.200
dracozlib	0.210
originalcbor	0.220
originalcborreplacezlib	0.250
originalzlib	0.250
originalcborreplace	0.280
dracocbor	0.295
draco	0.330
originalreplace	0.360

Table 6.58: Median performance with editgeomall
on large datasets without attributes, compres-
sion in advance

Compression type	median
dracocborreplacezlib	0.68
dracozlib	0.71
dracocborzlib	0.80
originalcborzlib	0.86
dracocborreplace	0.91
dracoreplace	0.99
originalzlib	1.04
draco	1.11
originalcborreplace	1.15
originalcborreplacezlib	1.15
originalreplace	1.72

Table 6.59: Median performance with editgeomall on small datasets, compression in advance

Compression type	median
originalcborreplace	0.595
originalcborzlib	0.625
originalcbor	0.640
originalzlib	0.655
originalreplace	0.730
dracozlib	0.980
dracocborreplace	1.040
dracocborreplacezlib	1.050
originalcborreplacezlib	1.050
dracocborzlib	1.065
dracocbor	1.125
draco	1.155
dracoreplace	1.225

Table 6.60: Median performance with editgeomall on small datasets without attributes, compression in advance

Compression type	median
originalcborreplacezlib	0.865
originalcborreplace	0.885
originalreplace	0.935
originalzlib	1.100
originalcborzlib	1.265
dracocborzlib	1.845
dracocborreplacezlib	1.910
dracozlib	1.925
draco	2.010
dracocborreplace	2.025
dracoreplace	2.785

Compression on the fly

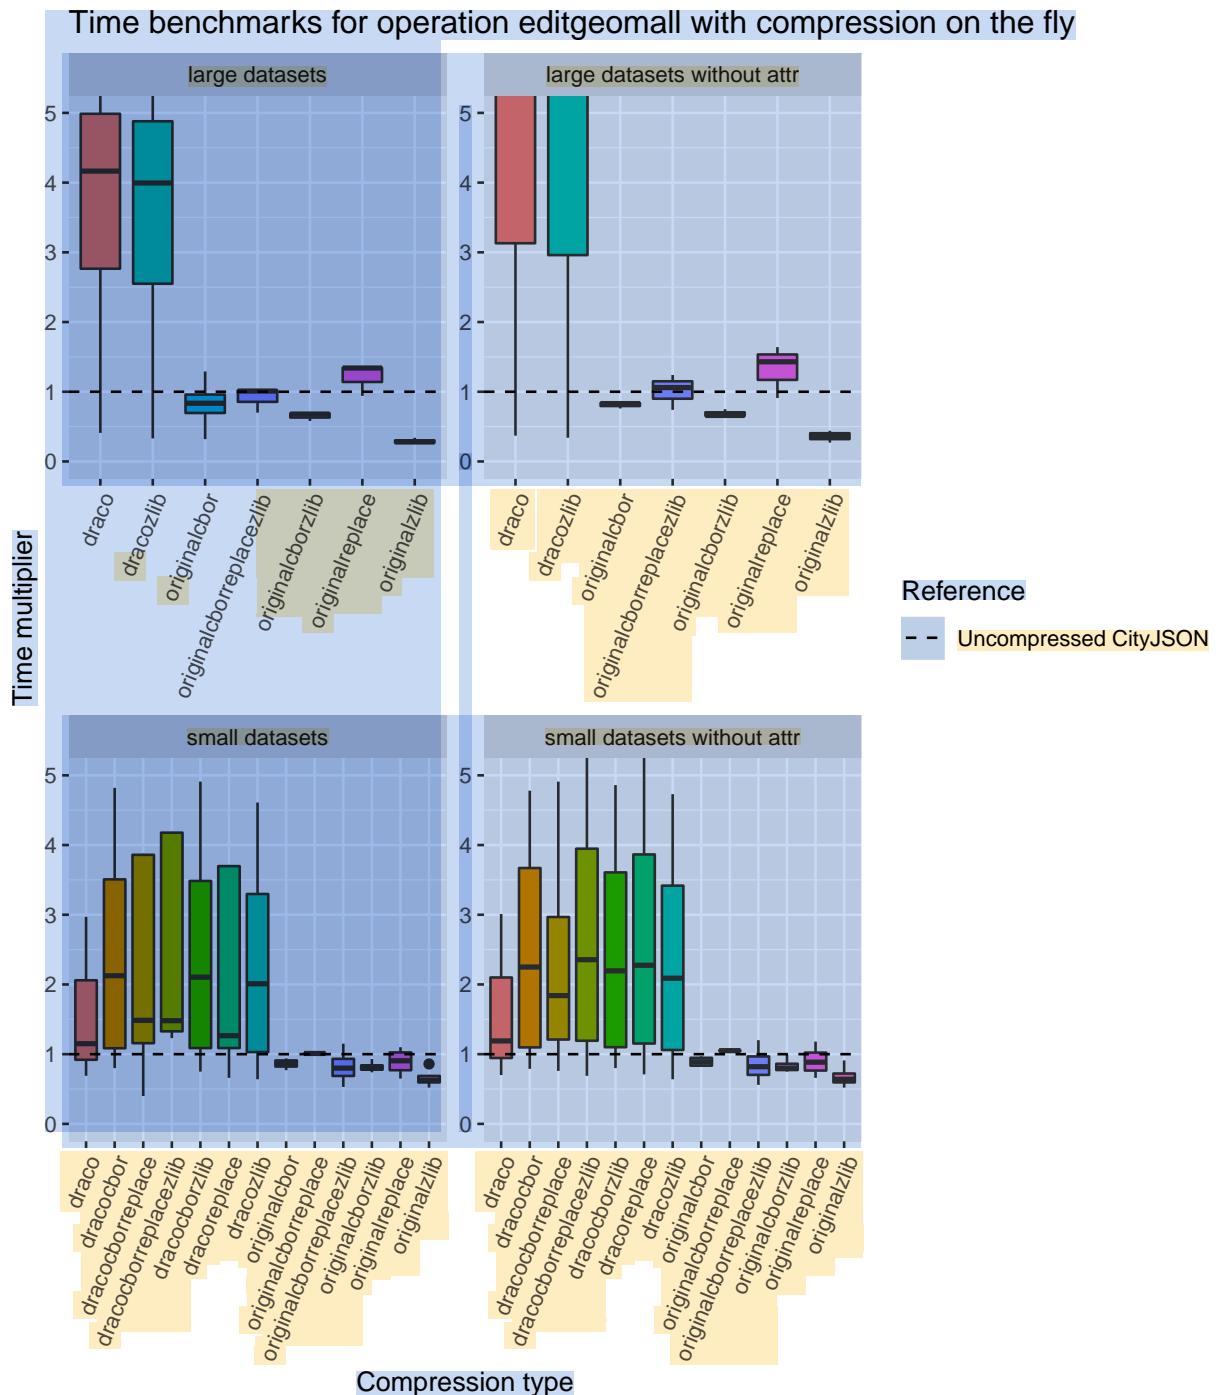


Figure 6.21: Boxplots of time performance of (on the fly) compressed datasets on editing geometry of all features compared to uncompressed CityJSON.

There was an error with multiple compression types in combination with this use case and large datasets, which is why they are missing from the charts. It is not very important, since it was mostly a problem with Draco which has already proven to perform badly with on the fly compression.

What happens in this use case is that basically full datasets are compressed after all being edited

6 Benchmarking results and analysis

with the same algorithm. The results are thus not so interesting and should come close to the theoretical performance gain. With the large datasets original-zlib and original-cbor-zlib do well, and original-cbor as well when attributes are not included. As for the small datasets, these three do well with both small dataset types. They are relatively uncomplicated compression types and their good performance can be explained by the balance between file size and processing time, despite compression time being included.

On the fly compression can improve the geometry editing time of CityJSON datasets on all city objects to 28%, 37%, 62%, or 64% (when respectively using large datasets, large datasets without attributes, small datasets, and small datasets without attributes) of the time that it takes to perform the same operation with uncompressed CityJSON.

Table 6.61: Median performance with editgeomall on large datasets, compression on the fly

Compression type	median
originalzlib	0.280
originalcborzlib	0.680
originalcbor	0.835
originalcborreplaced zlib	1.010
originalreplace	1.340
draco zlib	3.995
draco	4.165

Table 6.62: Median performance with editgeomall on large datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.37
originalcborzlib	0.66
originalcbor	0.83
originalcborreplaced zlib	1.06
originalreplace	1.43
draco zlib	5.58
draco	5.89

Table 6.63: Median performance with editgeomall on small datasets, compression on the fly

Compression type	median
originalzlib	0.625
originalcborreplaced zlib	0.800
originalcborzlib	0.800
originalcbor	0.870
originalreplace	0.905
originalcborreplaced	1.010
draco	1.150
dracoreplace	1.265
dracocborreplaced zlib	1.480
dracocborreplace	1.485
draco zlib	2.010
dracocborzlib	2.105
dracocbor	2.125

Table 6.64: Median performance with editgeomall on small datasets without attributes, compression on the fly

Compression type	median
originalzlib	0.640
originalcborzlib	0.800
originalcborreplaced zlib	0.820
originalreplace	0.885
originalcbor	0.890
originalcborreplaced	1.050
draco	1.190
dracocborreplaced	1.840
draco zlib	2.090
dracocborzlib	2.195
dracocbor	2.250
dracoreplace	2.275
dracocborreplaced zlib	2.355

6.9 Conclusion and recommendations

Boxplots with overview of results, averaged over all operations,
compression in advance

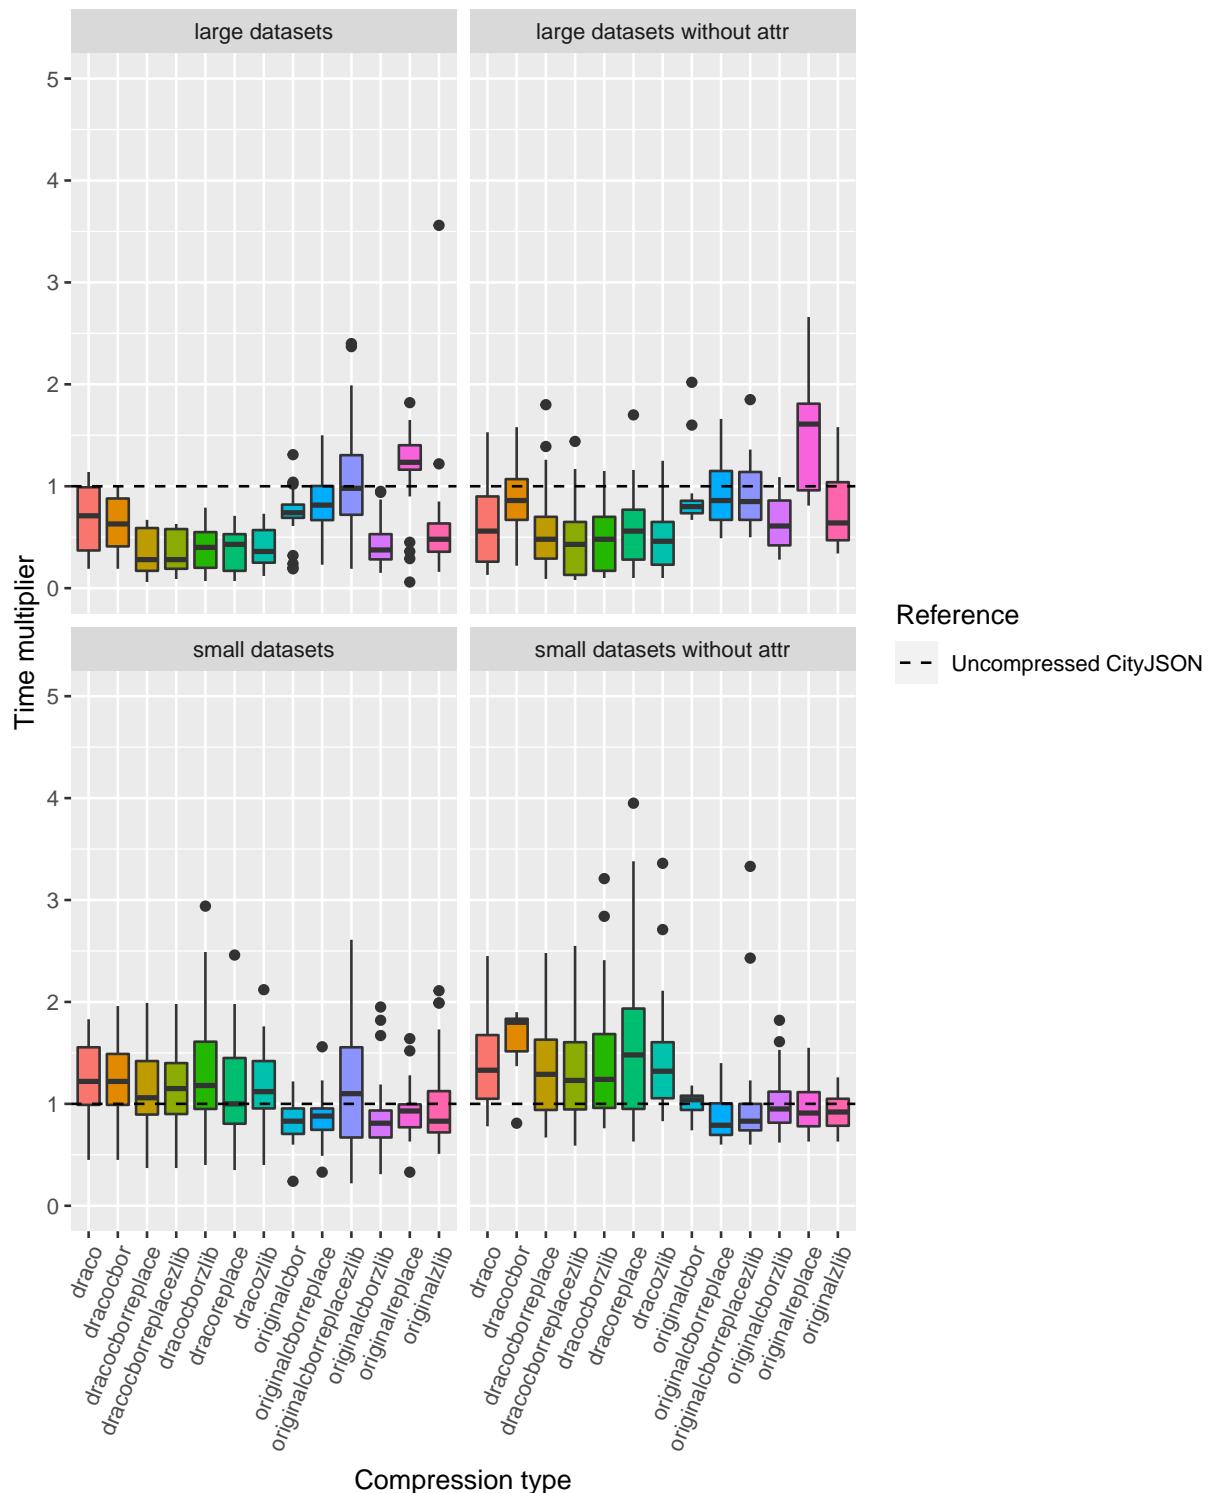


Figure 6.22: Overview of the performance of all compression types by dataset type. The results are averaged over every operation.

Boxplots with overview of results, averaged over all operations,
compression on the fly

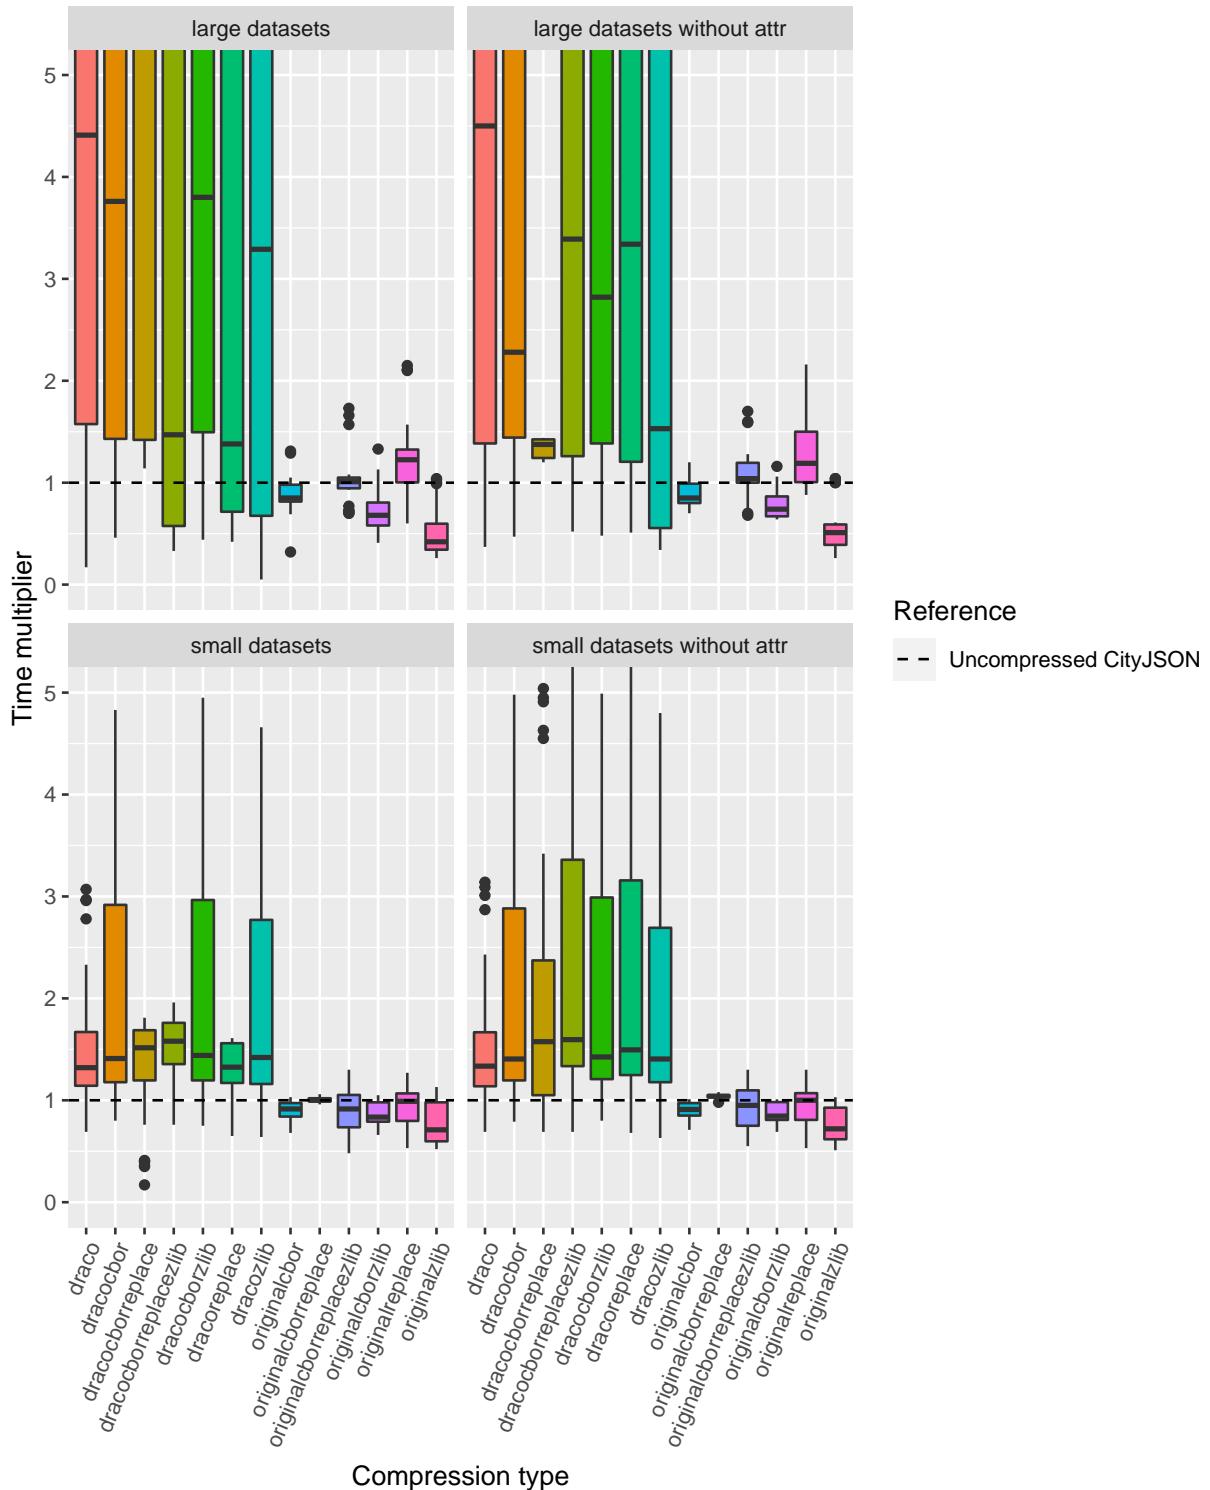


Figure 6.23: Overview of the performance of all compression types by dataset type. The results are averaged over every operation.

On Figures 6.22 and 6.23, the difference between operations is removed here such that appropriate compression types for general purpose can be found. It is possible to pick a best one based on specific cases—for a dataset type and operation combination that is known beforehand. You can see in the presented charts which compression type you should then choose for the best performance. But sometimes a software developer would want to write flexible software that works well with a larger variety of use cases, rather than creating it for visualisation or editing only. Additionally, implementing multiple compression types for CityJSON will make it more complex as a file type, because all software that uses it would need to be able to handle all the compressed versions. For these reasons these overviews are given, such that compression types that perform well in most cases can be identified.

Compression in advance

With the large datasets with attributes, the inclusion of Draco geometry is clearly beneficial for performance. This is especially so when it is combined with additional compression techniques. Draco-cbor-replace, draco-cbor-replace-zlib, draco-cbor-zlib, draco-replace, and draco-zlib all perform under the baseline with every operation. This is also true for original-cbor-zlib. Original-cbor and original-zlib are also good most of the time. This is similar for the large datasets without attributes. But generally, the results are less good and all of them do reach above the baseline for some specific operations. This indicates that the compression of attributes is actually of importance, since if they are not there, compression has less benefits as compared to attributes being included in the datasets.

Regarding the small datasets it is not as beneficial to compress datasets, both with and without attributes. In Chapter 1 I had already stated that the main goal of the thesis was to improve the performance with large files, and it is what I had expected to be improved the most. But original-cbor, original-cbor-replace, original-cbor-zlib, original-replace, and original-zlib still have median outcomes under the baseline with the small datasets that include attributes. This is similar for the ones without attributes, except that original-cbor-replace-zlib proved to also be beneficial overall, and original-cbor is a bit above the baseline. Draco compression is most of the time not useful with these dataset types since it adds overhead that is not compensated for by the lowered transmission time.

Compression on the fly

As opposed to compression in advance, the use of Draco is never a good idea when compressing data on the fly. This is due to the long compression times (see Section 6.1) of Draco, which is partially because of the intermediate conversion of CityJSON geometries to OBJ. As for compression types that retain original geometry, the boxplots also reach above the uncompressed CityJSON baseline when large datasets are processed. This means that there is none that is beneficial for any operation. But still, original-cbor, original-zlib, and original-cbor-zlib are useful if you look at the median results.

This also holds true with the small datasets, with the difference being that original-cbor-replace-zlib has a median just below the baseline as well. But most of the time, compression is still more useful with large datasets. In general the inclusion of attributes does not make a noticeable difference in the results.

Discussion

When interpreting the results it has also to be kept in mind that the baseline for the performance of small datasets is already small. A time multiplier of 0.5 looks nice, but in the end you are often

6 Benchmarking results and analysis

talking about milliseconds saved. However, this can add up when a lot of data is processed. Also, for bigger datasets, the time saved can be seconds to even minutes.

Concludingly, I have deliberately not singled out "best" compression types in the benchmarking conclusion because I think it is not possible based on these graphs, as there are no clear winners. In the previous Sections the best compression types for specific use cases can be seen, but I think a compression type that is actually implemented needs to do well all-round, because the implementation of multiple ones makes the development of software extra complicated. Additionally, these results only show performance but not the ease of implementation of the compression types, which is discussed in Chapter 7.

Generally, compression gives the greatest benefits when large datasets are used. The removal of attributes from datasets does not show any noticeable differences in outcome. While Draco compression does show good results, it does not work well all-round, especially when you consider small files and with on the fly compression. The replace method often is not beneficial either, and it does require the traversal of all city objects for both compression and decompression, which turns out to not be efficient. Considering both all-round performance and ease of use, I think that original-cbor, original-zlib, and original-cbor-zlib are the most suitable best compression types. They improve the general performance as can be seen from their median values in Figure 6.22. At the same time they are easy to implement, as both CBOR and zlib only need one line of code for encoding of decoding (see Section 5.4.1). Additionally they are widely supported and have libraries in many programming languages [CBOR, 2020; zlib, 2020].

7 Conclusion

7.1 Performed work and summary of achieved results

In the end I have tested 16 CityJSON datasets, represented in 14 different ways, on 9 operations in two different ways: with the datasets prepared in advance, and with them being compressed on the fly. This makes for 4032 benchmarks that have been condensed into 18 boxplot charts in Chapter 6, accompanied by additional benchmarks on (de)compression time and file size that help explaining the results. For every use case, the best-performing compression type has been found. I have also created two overview graphs for the two server implementations, which indicate what compression types have a good all-round performance compared to original CityJSON.

In some cases, I have achieved results of as good as 9% of the time that it would take the same operation with uncompressed CityJSON—when buffering the city objects of all larger datasets (that include attributes), with a file that is compressed beforehand with the draco-cbor-zlib type. This is however a very specific use case, and there are also cases in which compression is not worth it. For example, when querying one city object and compressing it on the fly.

The research question of the thesis is as follows:

To what extent can the implementation of different compression techniques improve CityJSON's performance on the web, considering file size, visualisation, querying, spatial analysis, and editing performance?

You could theoretically answer the question by stating the numbers of the best benchmarking results for every use case. I however think it does not make sense to state specific numbers to answer this question as I believe it does not matter for the actually useful conclusions of this thesis. The use cases are too specific and it is not a good idea to implement multiple compression types for specific cases only. It turned out to be more important what would be compression types that are worth implementing in general.

It is hard to pick the best general compression types based on performance alone, as there are none that outperform others in every case. But in addition to that, the performance itself is not the only important aspect of a compression type, since it also needs to be easily implementable by developers. CityJSON is an easy to use format and it should remain that way, as further discussed in Section 7.2. For these reasons, I think CBOR, zlib, or a combination of these two is the best to use. They are two compression techniques that are widely implemented and proved to perform well on average, regardless of the two types of server implementations and there is also not a large difference to be seen with them between different dataset types.

Draco compression only showed good results when compressing the data in advance and using it on larger datasets. For smaller datasets it adds too much overhead in (de)coding time, and for compressing it on the fly the compression time is too long. In addition to that, it proved hard to use as a developer. The C++ library has to be built beforehand, and the compression algorithms in Python have to call the command line to compress an OBJ file to a specified path. That also means that CityJSON geometries have to be converted to OBJ before this can be done, making the compression inefficient.

Using the JavaScript decoder, by default the Draco blob is decoded into one big mesh, as opposed to individual meshes for every city object. This is unwanted, since in a 3D city model you want

to be able to distinguish between features. I had to dive deep into the code myself to enable this distinction, and even then it does not work properly (see Section 5.1.3). To me it seems that Draco is very useful for visualisation only applications, but not when you want to do spatial analysis or be able to show the attributes of individual features—so not for geographical information system (GIS) applications. The latter part can change when this is actually implemented into the JavaScript decoder, and analysis can become easier if it would be possible to directly compress CityJSON geometries into Draco and back, which I think will not happen.

The replace method proved to not be very efficient. I believe this is due to all city objects having to be traversed for both compression and decompression (see Section 5.4.3). I had made up this technique myself, and the results do not indicate that it is worth it to further investigate. It is also not very easy to implement because of the need of recursive functions to be able to use it.

Again, CBOR and zlib do prove to be useful compression techniques for CityJSON, and can also be combined. But it would also be possible to use other binary JSON formats than CBOR (such as BSON, MessagePack, or UBJSON) and other general purpose techniques than zlib. I think this can be investigated a bit further before any kind of compression would actually be implemented for CityJSON.

7.2 CityJSON on the web and compression

I believe that it is best to use i3s or b3dm on the web when high visualisation performance is required. This is because of them being created with this specific purpose in mind. They already have tiling implemented. The large advantage that CityJSON has over these two is its ease of use. Therefore, I think this advantage should be preserved as much as possible and i3s and b3dm can be seen as complementary file formats rather than competitors.

CityJSON facilitates the creation of 3D models in a user-friendly way. This is (at least currently) easier than creating i3s and b3dm from scratch, since for instance cpio and 3dfier can be utilised to do create CityJSON datasets. CityJSON datasets can be inspected easily and the structure is more straightforward. Doing analysis with them is easier as well because of its simpler structure and geometries can be accessed and read easily, as opposed to having to work with the—often binary—glTF.

i3s and b3dm lack an extensive data model and need to be converted from other file types, which would then also have to be created first. This could be done from CityJSON as well when converters for it are written.

What I thus see as the ideal 3D model workflow is creating CityJSON datasets at first. It can be used to easily inspect the data, check its quality, and do analysis if wanted. Then if one is to publish the data for viewing on the web, it could be converted to one of the other formats.

That is if high performance is needed. CityJSON can be visualised on the web as well in a relatively simple manner as well, through converting the geometries to a three.js mesh. It definitely can be useful for web visualisation as well, let alone other web use such as dissemination.

CityJSON can thus perfectly exist together with i3s and b3dm as it has its own advantages. The focus should be kept on keeping it easy to use for regular users. Compression can make this more complicated, at least for developers, as there are extra steps to be taken to read or store a dataset. However, the use of massive datasets on the web can still be troublesome, which data compression could relieve. I therefore think that it is still feasible to implement compression functionalities for it, but with the strict requirement that it needs to be easy to use.

For me, the inclusion of Draco compression is out of the question. According to my experience it is not easy to work with, at least not when using it in liaison with 3D city models. This is because

Draco by default compresses all objects of an OBJ into one big object. It is possible to keep the separation between objects, but the JavaScript decoder is not able to handle this well out of the box and it will result in a less good compression. To make it somewhat work, I had to dive into the source code and make alterations which was time consuming. In addition to that, offline users would have to compile the Draco library on their computer in order to be able to read and write datasets that include Draco geometries, which is troublesome as well. Draco is hard to use with any operation and it is not suitable for compression on the fly because of its long encoding time.

It is however a good idea to include the possibility for parsing to data to and from CBOR. To use it only a lightweight library is required, and it already exists for most languages. Only one simple line of code is necessary to encode or decode it from and to JSON. While not compressing the data size as far as other compression method combinations, it still always results in a smaller file size. On top of that, the results show good performance for it, both when compressing it in advance and on the fly. Zlib is also suitable to use, and a combination of CBOR and zlib is a nice possibility as well.

7.3 Limitations

There are many variables in the benchmarking that have influenced the results.

The choice of datasets could be improved, as I picked ones that were easy to find. It also seemed that the smaller datasets were more error prone in testing as the benchmarked times were smaller. This might have been the cause in results seeming unexplainable. The plan was to analyse the differences in performances between the selected datasets. This however made it too complex as there was too much data and too many differing results to explain. Some of them made few sense, so I had decided to average them into the two groups of smaller and larger datasets.

In addition to that, I had removed the attributes of the datasets and did separate tests with those. But there are actually many more variable characteristics of datasets that can influence the results. Not only file size and the presence of attributes, but also the amount and type of attributes, their length, language, the inclusion of metadata, textures, semantics, LoDs, and so on. Testing all of these takes a lot of time, both in creating the datasets and benchmarking them, but also the results would be very complex to interpret well. While interesting, I did not analyse the dataset characteristics any further because of this reason.

Draco files are downloaded separately rather than loading them as embedded into compressed CityJSON files. This is due to a limitation in the JavaScript Draco decoder, which can only load a file from a path and can not be served a blob. While I have adapted the results to the additional time that it takes to download the separate file, I would imagine this adds some extra overhead to the process.

Two compression types that I had wanted to implement (Huffman coding and smaz) did not work in the end. It would have required more time to do this, but they might have proved to be useful.

Lastly, I have used a Flask server (in Python) for the processing of files. This is probably not as efficient (regarding performance) as 3DCityDB, which stores data in a database and processes it on the fly as requests are made. Using it could give different results, but at least the on the fly compression implementation gives an indication on it.

7.4 Future work

To improve the workings of CityJSON with Draco, the encoding CityJSON geometries into Draco geometries can be improved by implementing direct encoding, rather than having to convert CityJSON to OBJ first. It was seen with both the encoding times and the results with on the fly

7 Conclusion

compression that this really hampered the performance. Besides that it would make its use easier and if the JavaScript Draco decoder is adapted to handle individual objects well this would improve the developer experience as well.

Investigating the queryability of CBOR (or another binary JSON format) can also be of use, especially when datasets are compressed in advance. The server could then just retrieve the city objects that are requested directly, and only transmit these to the client. In addition to that, other binary formats might prove to yield better results than CBOR, so this should be investigated before implementing it. This is also the case with zlib.

Creating a server implementation with 3DCityDB and testing which compression types work best with it would be nice since the performance is likely to be better than with a Flask server. As stated in Section 7.3, the results with on the fly compression do give an indication on which compression types are the most suitable for further investigation.

8 Streaming

This is an additional Chapter on the streaming of CityJSON datasets. It is extra work that does not fit well in the other Chapters as it distracts from the main parts of the thesis. First I explain what the OGC API is, then how CityJSON could be streamed complying with it, and lastly how compression is relevant to it.

8.1 OGC API

The OGC API [Open Geospatial Consortium, 2020] is a new specification that defines how to provide online access to geospatial data. This is done in a modular way, with several building blocks that can be used together as needed when creating a RESTful Web API for geospatial information. It succeeds the older well-known OGC standards such as WMS and WFS. With CityJSON being an encoding for 3D features we focus on the OGC API Features block, which could be seen as the successor of WFS2 (and it was formerly named WFS3) [OGC API, 2020; Open Geospatial Consortium, 2020].

This block consists of multiple modules as well (on CRS, Contextual Query Language (CQL), and simple transactions) and has the possibility to be extended with more modules. In contrast with WFS, OGC API Features is created with other file encodings than XML in mind. This enables to make it work with JSON and HTML easily, which is encouraged but not mandatory as any encoding may be used [Open Geospatial Consortium, 2020, 2019a]. Because of this, it is possible to make CityJSON compliant with the specifications of the API. Ultimately, any kind of geospatial data consisting of features can be served following the guidelines of the OGC API, as long as a specification is created for it.

The API works through visiting specifically formatted links that indicate what kind of data is requested from a server. The base URL of the server (with path "/") needs to provide links to:

1. The API definition—the capabilities of the server
2. The Conformance declaration—URIs referring to the standards that the server implements
3. The Collections—the datasets that are available on the server

By visiting specifically formatted links, Collections can be queried as desired and are served in the requested encoding. Information on the Collections is found following the /collections path, with /collections/{collectionId} showing more detailed information on a specific Collection. /collections/collectionId/items will return the full dataset, and with /collections/collectionId/item/itemId one specific feature can be retrieved.

It is mandatory to provide a function to query the features of a Collection on a bounding box in for example the following way: /collections/{collectionId}/?bbox=160.6,-55.95,-170,-25.89. Similarly to WFS, it should be possible to include a limit parameter in the link to only receive a maximum amount of features. Combining this with an offset parameter, it is possible to stream the data. For instance /collections/{collectionId}/item/{itemId}/limit=10&offset=0 would return the first 10 features of the requested dataset. By increasing the offset by 10 in every new request, the full dataset is traversed. Other specific queries could be formatted in a similar manner

8 Streaming

as well. These are the most important requirements, but not all are listed and there are also more recommended implementation details [Open Geospatial Consortium, 2019b].

Currently the use of the API has been mostly described in liaison with GeoJSON [GeoJSON], which is recommended, and the WGS84 CRS which is the default. A server would return a GeoJSONCollection or GeoJSONFeature when data is requested. It is however possible to use a different file encoding and a different CRS when the CRS module is used [Open Geospatial Consortium, 2019b]. Section 8.2 therefore explains how CityJSON could be used in accordance with the OGC API.

8.2 Streaming CityJSON

GeoJSON [GeoJSON] does not support 3D data, and for OGC API to work with 3D data a new specification would have to be written. A CityJSON object can already be seen as a collection of features. It can therefore be argued that its original structure is already OGC API compliant since the specification allows for flexibility in data structures that servers return. However, the streaming of datasets is part of the OGC API specification since it is possible to add `limit` and `offset` parameters to a link.

For this purpose a draft [Ledoux, 2020] has been made to make CityJSON compliant to the OGC API as it is implemented now (with GeoJSON), and to enable streaming. Because CityJSON has the list of vertices at the end of the file rather than per city object, it is not possible to stream it in the fashion that it is done with GeoJSON. With GeoJSON, the receiver of the stream is fed with GeoJSON Features or FeatureCollections in line-delimited JSON. Once received, it can be used immediately as it already includes the complete geometry.

In order to enable CityJSON to stream in a similar way, the structure of listing 8.1 is proposed by Ledoux [2020]. The "type" is now a CityJSONFeature rather than a CityJSON object. "CityObjects" contains a parent and its children, with the "id" of the CityJSONFeature being the one of the parent. All their vertices can now be accessed immediately, without needing to have the whole file. A CityJSONFeature can thus be loaded in the same manner as a CityJSON, thus avoiding the need to adapt the software that will use it. In Section 8.3 I investigate how compression could be implemented for streaming.

Listing 8.1: Snippet of OGC API compliant CityJSON structure. Source: Ledoux [2020]

```
{  
  "type": "CityJSONFeature",  
  "id": "myid", //-- to clearly identify which of the CityObjects is the "main" one  
  "CityObjects": {},  
  "vertices": [],  
  "appearance": {}  
}
```

8.3 Streaming CityJSON and compression

Subsection 8.2 has introduced a proposal for a structure in which CityJSON could be streamed over the web. Streaming is related to the efficient web use of 3D geoinformation, but it is not within the scope of this thesis to test streaming performance. Despite that I have looked into how compression can potentially contribute to efficient streaming to see if it is worthwhile to investigate it further in the future.

Listing 8.1 shows how CityJSONFeatures look like, which can be streamed individually as they contain their own array of vertices. This gives the possibility to compress the geometry of a CityJSONFeature with Draco. There are two ways to do this: compressing every City Object of a Feature into individual Draco geometries (which would then be a parent with its children), and compressing all City Objects of a Feature together in one Draco geometry. The former option removes the need of having to separate objects within the Draco file, thus object IDs not needing to be stored inside. With the second option this does have to be done (just like when compressing complete CityJSON datasets)—but only if the Feature contains multiple City Objects—but having one Draco file might be more efficient as the header of the file is relatively smaller.

A caveat is that datasets need to have "transform" to be able to be compressed properly with Draco, while the streaming implementation currently does not allow this Ledoux [2020]. Additionally, since Draco blobs can not be stored in JSON it is necessary to parse the file into a binary format, for which CBOR is already used throughout the thesis. CBOR on its own however can also potentially improve the performance.

The file sizes of four streaming variants of the datasets used in this thesis are shown in Table 8.1. The CityJSONFeatures of every dataset are aggregated in a CityJSONFeatureCollection, akin to GeoJSON's FeatureCollection. The "collection" column this represents uncompressed CityJSON datasets that can be used for streaming. "collection-cbor" is compressed with CBOR, "collection-cbor-draco" includes individual Draco blobs for every City Object of a CityJSONFeature, and "collection-cbor-draco2" has one Draco blob per CityJSONFeature.

Dataset	collection	collection-cbor	collection-cbor-draco	collection-cbor-draco2
Delft	1.2	0.7	0.6	0.6
Den Haag	3.3	2.4	2.4	2.3
Rotterdam	3.9	2.5	2.6	2.6
Montréal	5.4	3.5	2.9	2.9
Singapore	54.5	34.4	32.1	32.1
TU Delft campus	70.6	41.8	25.5	25.5
New York	109.8	69.7	70.2	70.2
Zürich	284.2	182.5	163.2	165.6

Table 8.1: File size in MB for CityJSONCollections of datasets and compressed versions

It is shown that using CBOR already compresses the files, and possibly increases the reading and loading speed as compared to JSON (see Section 3.6.1 for an explanation on this). The addition of Draco compression is only worth it for datasets in which City Objects have parent-child relationships. The way in which Draco is incorporated is almost completely insignificant—the second Draco implementation gives a slightly better result for Den Haag, but a worse result with Zürich.

Concludingly, compression can be beneficial for streaming as well. But while the implementation with CBOR could easily be tested, this is not the case when using Draco. Out of the box it is not possible to directly load Draco blobs in the JavaScript decoder, so this would need to be implemented first.

A Datasets and characteristics

These are snippets of the eight CityJSON datasets used in this thesis. One City Object is shown and only a small part of the geometry. In orange are the attributes and in green the geometries.

A.1 Delft

Listing A.1: Snippet of Delft dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "geographicalExtent": [
      84616.468, 447422.999, -0.47,
      85140.839, 447750.636, 13.8
    ],
    "referenceSystem": "urn:ogc:def:crs:EPSG::7415"
  },
  "CityObjects": {
    "b0a8da4cc-2d2a-11e6-9a38-393caa90be70": {
      "attributes": {
        "bgt_status": "bestaand",
        "bronhouder": "G0503",
        "class": "dek",
        "creationdate": "2014-07-09",
        "eindregistratie": "",
        "hoortbijtypeoverbrugging": "waardeOnbekend",
        "inonderzoek": "0",
        "lokaalid": "G0503.032e68f09d7049cce0532ee22091b28c",
        "lv_publicatiедatum": "2016-06-07T16:22:15.000",
        "namespace": "NL.IMGeo",
        "overbruggingisbeweegbaar": "0",
        "plus_status": "geenWaarde",
        "relatievehoogteligging": "1",
        "terminationdate": "",
        "tijdstipregistratie": "2016-05-17T13:43:18.000"
      },
      "geometry": [
        {
          "boundaries": [
            [ [ 11337, 7, 9 ] ],
            [ [ 11337, 43, 7 ] ],
            [ [ 14, 15, 13 ] ], ...
          ],
          "lod": 1,
          "type": "MultiSurface"
        }
      ],
      "type": "Bridge"
    }
  },
  "transform": {
    "scale": [
      0.001, 0.001, 0.001
    ],
    "translate": [
      84616.468, 447422.999, -0.47
    ],
    "vertices": [
      [ 85012.343, 447455.577, -0.27 ],
      [ 85010.804, 447448.808, -0.27 ],
      [ 85013.832, 447447.447, -0.27 ], ...
    ]
  }
}
```

A.2 Den Haag

Listing A.2: Snippet of Den Haag dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "CityObjects": {"GUID_DBDA5F53-7DD5-4C2F-BE7F-51F29A0CBA16_1": {
    {"type": "BuildingPart",
      "attributes": {
        "roofType": "1000",
        "RelativeEavesHeight": 3.133,
        "RelativeRidgeHeight": 3.133,
        "AbsoluteEavesHeight": 7.717,
        "AbsoluteRidgeHeight": 7.717
      },
      "geometry": [
        {"type": "Solid",
          "boundaries": [
            [[[10915, 10991, 10992, 10916]], [[10991, 10993, 10994, 10992]], [[10993, 10995, 10996, 10994]], ...],
            {"semantics": {
              "values": [[0, 1, 2, 3, 4, 5]],
              "surfaces": [
                {"type": "WallSurface"},
                {"type": "WallSurface"},
                {"type": "WallSurface"},
                {"type": "WallSurface"},
                {"type": "RoofSurface", "Direction": 0, "Slope": 90},
                {"type": "GroundSurface"}]
            },
            "material": {
              "": {"values": [
                [2313, 2313, 2313, 2313, 2314, 2315]]}
            },
            "lod": 2
          ],
          "parents": ["GUID_DBDA5F53-7DD5-4C2F-BE7F-51F29A0CBA16"]},
          ...
        },
        "vertices": {
          [[506451, 10377, 5928], [506316, 403, 5941], [511459, 5304, 5921], ...]
        },
        "metadata": {
          "geographicalExtent": [
            78248.66,
            457604.591,
            2.463,
            79036.024,
            458276.439,
            37.481
          ],
          "referenceSystem": "urn:ogc:def:crs:EPSG::7415"
        },
        "transform": {
          "scale": [
            0.001,
            0.001,
            0.001
          ],
          "translate": [
            78248.66,
            457604.591,
            2.463
          ]
        }
      }
    }
  }
}
```

A Datasets and characteristics

A.3 Montréal

Listing A.3: Snippet of Montréal dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "geographicalExtent": [
      300060.181, 5040722.785, 8.444,
      301073.701, 5042282.343, 129.092
    ]
  },
  "CityObjects": {
    "B-201391182746-80FE6BE00A9C": {
      "type": "Building",
      "geometry": [
        {
          "type": "MultiSurface",
          "boundaries": [
            [
              [
                [
                  22432, 22433, 22434
                ]
              ],
              [
                [
                  22434, 22435, 22436
                ]
              ],
              ...
            ],
            "semantics": {
              "values": [
                0, 0, 0,
                ...
              ],
              "surfaces": [
                {
                  "type": "GroundSurface"
                },
                {
                  "type": "WallSurface"
                },
                ...
              ],
              "texture": {
                "Rhino_texturing": {
                  "values": [
                    [
                      [
                        238, 63606, 63606, 63606
                      ]
                    ],
                    [
                      [
                        238, 63606, 63606, 63606
                      ]
                    ],
                    ...
                  ]
                },
                "lod": 2
              }
            ]
          }
        ],
        "vertices": [
          [
            300762.679, 5041624.941, 14.274
          ],
          [
            300749.896, 5041617.688, 14.274
          ],
          [
            300755.943, 5041628.802, 14.274
          ],
          ...
        ]
      ]
    }
  }
}
```

A.4 New York

Listing A.4: Snippet of New York dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "geographicalExtent": [
      985320.346, 219464.155, -1.5,
      1009996.302, 259992.617, 677.862
    ]
  },
  "CityObjects": {
    "gml_7KKGXJDCBAXCTELB700MHK04I376NBA6NGHC": {
      "type": "Building",
      "attributes": {
        "SOURCE_ID": "13210006270",
        "BIN": "1061273",
        "DOITT_ID": "168309"
      },
      "geometry": [
        {
          "type": "MultiSurface",
          "boundaries": [
            [
              [
                [
                  327517,
                  327518,
                  327519,
                  ...
                ],
                ...
              ],
              ...
            ],
            "semantics": {
              "values": [
                0,
                1,
                2,
                ...
              ],
              "surfaces": [
                {
                  "type": "GroundSurface",
                  "SOURCE_ID": "13210006270",
                  "BIN": "1061273",
                  "DOITT_ID": "168309"
                },
                ...
              ],
              ...
            }
          ],
          "lod": 2
        }
      ]
    },
    "vertices": [
      [
        19709059, 40053117, 21766
      ],
      [
        19699777, 40061095, 21766
      ],
      [
        19682224, 40040679, 21766
      ],
      ...
    ]
  }
}
```


A.5 Rotterdam

Listing A.5: Snippet of Rotterdam dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "referenceSystem": "urn:ogc:def:crs:EPSG::28992",
    "geographicalExtent": [
      92824.468, 436417.288, -7.436, 93947.37, 437464.569, 126.071
    ],
    "presentLoDs": {
      "1.0": 546, "2.0": 550
    }
  },
  "CityObjects": {
    "ID_0599100000643092": {
      "address": {
        "CountryName": "Nederland", "LocalityName": "Rotterdam",
        ...
        "location": {
          "type": "MultiPoint",
          "boundaries": [
            20
          ],
          "lod": 0
        }
      },
      "type": "Building",
      "attributes": {
        "yearOfConstruction": 1983, "creationDate": "2019-06-11", "gebouwnummer": "0599100000643092",
        ...
      },
      "geographicalExtent": [
        93460.98, 437063.18, 3.01130000000232, 93471.11, 437072.8, 9.69260888055718
      ],
      "geometry": [
        {
          "type": "Solid",
          "boundaries": [
            [
              [ 0, 1, 2, 3 ],
              ...
            ]
          ],
          "lod": 1
        },
        {
          "type": "Solid",
          "boundaries": [
            [
              [ 2, 11, 9, 6, 4, 3 ],
              ...
            ]
          ],
          "semantics": {
            "values": [
              [ 0, 1, 2,
                ...
              ]
            ],
            "surfaces": [
              {
                "type": "GroundSurface", "id": "UUID_d7d650ac-6751-4285-aaac-ed6b4929d23c",
                "creationDate": "2019-06-11"
              },
              ...
            ],
            "texture": {
              "RhinoCity texturing": {
                "values": [
                  [ [ 0, 0, 1, ... ] ],
                  ...
                ]
              }
            },
            "lod": 2
          }
        }
      ],
      "vertices": [
        [
          646642, 651352, 14924
        ],
        ...
      ]
    }
  }
}
```

A.6 Singapore

Listing A.6: Snippet of Singapore dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "datasetTitle": "3D city model of public housing (HDB) buildings in Singapore",
    "datasetReferenceDate": "2019-08-25",
    "geographicLocation": "Singapore, Republic of Singapore",
    "referenceSystem": "urn:ogc:def:crs:EPSG::3414",
    "geographicalExtent": [
      11474.615816611371, 28054.808157231186, 0,
      45326.44585339671, 48758.78176700817, 141.3
    ],
    "datasetPointOfContact": {
      "contactName": "NUS Urban Analytics Lab, National University of Singapore",
      "emailAddress": "filip.nus.edu.sg",
      "contactType": "organization",
      "website": "https://ual.sg/"
    },
    "metadataStandard": "ISO 19115 - Geographic Information - Metadata",
    "metadataStandardVersion": "ISO 19115:2014(E)"
  },
  "CityObjects": {
    "way172927954": {
      "type": "Building",
      "attributes": {
        "osm_timestamp": "2016-06-25T01:20:18",
        "osm_version": 3,
        "osm_changeset": 40274256,
        "osm_user": "JaLooNz",
        "osm_uid": 741163,
        "osm_building": "residential",
        "osm_addr:city": "Singapore",
        "osm_addr:street": "Beach Road",
        ....
      },
      "geometry": [
        {
          "type": "Solid",
          "lod": 1,
          "boundaries": [
            [
              [
                [
                  [
                    [0, 1, 2, 3]
                  ]
                ],
                [
                  [
                    [
                      [1, 4, 5, 2]
                    ]
                  ],
                  ...
                ]
              ]
            ]
          }
        }
      ],
      "vertices": [
        [
          199673130715612, 36984818430307, 0
        ],
        [
          199701511315904, 36869046893490, 0
        ],
        [
          199701511315904, 36869046893490, 461000000000
        ],
        ...
      ]
    }
  }
}
```

A.7 TU Delft campus

Listing A.7: Snippet of TU Delft campus dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "geographicalExtent": [
      84527.3237110656, 444659.804912295, 0.0,
      86603.7691803278, 447427.906238934, 0.0
    ],
    "referenceSystem": "urn:ogc:def:crs:EPSG::7415"
  },
  "CityObjects": {
    "b00a922cc-e6d0-11e7-951f-610a7ca84980": {
      "attributes": {
        "bgt-functie": "fietspad",
        "bgt-fysiekvoorkomen": "open verharding",
        "bgt-status": "bestaand",
        "bronhouder": "G0503",
        "creationdate": "2017-11-22",
        "eindregistratie": "",
        "inonderzoek": "0",
        "lokaalid": "G0503.028ec3c1e25143e3b667dba12130f554",
        "lv-publicatiedatum": "2017-12-04T11:39:48.000",
        "namespace": "NL.IMGeo",
        "plus-functiewegdeel": "waardeOnbekend",
        "plus-fysiekvoorkomenwegdeel": "betonstraatstenen",
        "plus-status": "geenWaarde",
        "relatievehoogteligging": "0",
        "terminationdate": "",
        "tijdstipregistratie": "2017-12-04T11:13:31.000",
        "wegdeeloptalud": "0",
        "class": "Wegdeel",
        "function": "open verharding"
      },
      "geometry": [
        {
          "boundaries": [
            [
              [
                [
                  [0,
                   1,
                   2]
                ]
              ],
              [
                [
                  [
                    [1,
                     3,
                     2]
                  ]
                ]
              ],
              [
                [
                  [
                    [
                      [1,
                       2,
                       3]
                    ]
                  ]
                ]
              ],
              [
                [
                  [
                    [
                      [
                        "lod": 1,
                        "type": "MultiSurface"
                      ]
                    ]
                  ]
                ]
              ],
              [
                [
                  [
                    [
                      [
                        [
                          [
                            [
                              "type": "Road"
                            ]
                          ]
                        ]
                      ]
                    ]
                  ]
                ]
              ]
            ],
            "vertices": [
              [
                1618304, 572011, 98810
              ],
              [
                1625102, 574863, 98980
              ],
              [
                1617445, 574088, 98810
              ],
              ...
            ]
          ]
        }
      ]
    }
  }
}
```


A.8 Zürich

Listing A.8: Snippet of Zurich dataset

```
{
  "type": "CityJSON",
  "version": "1.0",
  "metadata": {
    "geographicalExtent": [
      2677116.375, 1241839.025, 0.0,
      2689381.984, 1254150.95, 1044.25
    ],
    "presentLoDs": {
      "2.0": 145865
    }
  },
  "CityObjects": {
    "UUID_93fc5bae-4446-4336-9ff8-6679ebfdfde3": {
      "type": "BuildingPart",
      "attributes": {
        "creationDate": "2017-01-23",
        "Geomtype": 1
      },
      "geographicalExtent": [
        2678701.758, 1252228.936, 509.796,
        2678714.997, 1252238.97, 519.178
      ],
      "parents": [
        "UUID_c9884c4e-1cac-47f5-b88b-6fb074c0ae50"
      ],
      "geometry": [
        {
          "type": "MultiSurface",
          "boundaries": [
            [
              [
                [
                  0, 1, 2, 3
                ]
              ]
            ],
            [
              [
                [
                  3, 4, 5, 0
                ]
              ]
            ],
            ...
          ],
          "semantics": {
            "values": [
              0,
              1,
              2,
              ...
            ],
            "surfaces": [
              {
                "type": "GroundSurface"
              },
              {
                "type": "WallSurface"
              },
              ...
            ],
            "lod": 2
          }
        }
      ],
      "vertices": [
        [
          1587680, 10389911, 509796
        ],
        [
          1585382, 10394788, 509796
        ],
        [
          1596325, 10399945, 509796
        ]
      ]
    }
  }
}
```

B List of used tools

B.1 Shapely

Shapely [Shapely, 2020] is a Python library used for the analysis and manipulation of planar features. I use it to perform buffer operations in the client.

B.2 cpio

Standing for CityJSON/io, [3D geoinformation research group at TU Delft, 2019] is a Python command line interface (CLI) for processing and manipulating CityJSON datasets.

B.3 Flask

Flask [Flask, 2020] is a lightweight Python web framework with which web applications can be built. It allows harp.gl to work together with cpio, as Flask acts as a server and enables the preparation of web pages.

B.4 three.js

A JavaScript library for the visualisation of 3D graphics in the browser, utilising the lower level Web Graphics Library (WebGL) [Three.js, 2020a]. It is used in this thesis to visualise CityJSON datasets.

B.5 harp.gl

An open source 3D web map rendering engine for JavaScript [HERE, 2019]. It utilises WebGL and three.js. Used for the visualisation of CityJSON datasets.

B.6 Turf.js

It is a library written in JavaScript for geospatial analysis [Turf.js, 2020]. I have used it to buffer geometries in the client.

B.7 Proj4js

Proj4js [Proj4js, 2020] is a JavaScript library that enables the transformation of coordinates to other CRSS.

B.8 Selenium

Selenium [Muthukadan, 2020] can drive web browsers in an automated way. It has Python bindings and I have used it for the benchmarking part of the thesis.

B.9 Draco

An open-source library by Google [Google, 2019c] that is used to compress and decompress 3D geometric meshes for the purpose of improved storage and transmission efficiency. It has functionalities in C++, JavaScript, and WebAssembly. Used for the compression of geometries of CityJSON datasets.

C Reproducibility self-assessment

I have made all code and raw benchmarking results available on GitHub, and have added the sources to the datasets that I have used. This means that anyone could create the same compressed datasets as I have. Dataset characteristics and theoretical results could thus be entirely reproduced. However, the benchmarking results have some uncertainties in them, despite always being iterated for 10 times. It is dependent on computer specifications and mostly randomness (at least to me—I can not completely explain the inner workings of a server, browser, and computer). So it is not possible to exactly recreate these results, but if every test is iterated in the same way with similar computer specifications, it is very likely that the results are close enough for the analysis to remain the same.

But I could have maintained my code better on GitHub and in general. Parts are hard to read and not always properly commented, or have outdated ones. So if one is to not reuse my code and do the research from scratch (only looking at my code rather than copying), it is likely that there will be some differences in the recreation.

D Full benchmark results

task	method	delft	denhaag	hdlb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
bufferall	draco	NA	1.61	0.27	0.81	NA	1.57	NA	NA
bufferall	dracochbor	NA	1.61	0.23	0.81	NA	1.77	NA	NA
bufferall	dracochborreplace	NA	1.54	0.12	0.82	NA	1.56	NA	NA
bufferall	dracochborreplacezlib	NA	1.49	0.14	0.88	NA	1.76	NA	NA
bufferall	dracochborzlib	NA	1.65	0.09	0.87	NA	2.31	NA	NA
bufferall	dracoreplace	NA	1.59	0.14	0.8	NA	1.63	NA	NA
bufferall	dracozlib	NA	1.49	0.12	0.78	NA	1.54	NA	NA
bufferall	original	NA	1	1	1	NA	1	NA	NA
bufferall	originalcbor	NA	0.85	1.02	0.76	NA	1.07	NA	NA
bufferall	originalcborreplace	NA	1.16	0.78	0.7	NA	0.71	NA	NA
bufferall	originalcborreplacezlib	NA	1.1	0.9	0.57	NA	0.61	NA	NA
bufferall	originalcborzlib	NA	0.96	0.87	0.7	NA	0.88	NA	NA
bufferall	originalreplace	NA	1.16	1.14	0.85	NA	0.65	NA	NA
bufferall	originalzlib	NA	0.72	0.77	0.72	NA	0.95	NA	NA
bufferone	draco	1.7	1.65	1.07	1.3	1.11	1.47	0.65	1.04
bufferone	dracochbor	1.69	1.71	0.95	1.54	0.99	1.39	0.63	0.95
bufferone	dracochborreplace	1.85	1.54	0.67	1	0.63	1.45	0.57	0.63
bufferone	dracochborreplacezlib	1.75	1.48	0.63	1.15	0.62	1.38	0.51	0.63
bufferone	dracochborzlib	2.94	1.61	0.59	1.32	0.61	1.4	0.5	0.55
bufferone	dracoreplace	1.98	1.82	0.71	1.05	0.43	1.32	0.43	0.44
bufferone	dracozlib	1.76	1.48	0.68	1.12	0.73	1.41	0.23	0.68
bufferone	original	1	1	1	1	1	1	1	1
bufferone	originalcbor	0.95	0.96	0.83	0.87	0.75	0.89	0.7	0.83
bufferone	originalcborreplace	0.96	0.95	0.71	0.68	0.58	0.73	0.78	0.74
bufferone	originalcborreplacezlib	0.98	0.98	0.75	0.68	0.55	0.75	0.75	0.71
bufferone	originalcborzlib	0.87	0.85	0.39	0.84	0.32	0.68	0.43	0.4
bufferone	originalreplace	1.11	1.04	1.22	0.86	0.9	0.8	1.22	1.23
bufferone	originalzlib	0.86	0.8	0.51	0.71	0.39	0.82	0.5	0.46
editattrall	draco	0.95	0.98	0.66	0.8	0.9	1.02	0.32	0.6
editattrall	dracochbor	1.07	1.03	0.57	0.96	0.52	0.92	0.29	0.43
editattrall	dracochborreplace	0.94	0.74	0.24	0.66	0.12	0.9	0.16	0.09
editattrall	dracochborreplacezlib	1.12	0.77	0.25	0.74	0.09	1.14	0.17	0.17
editattrall	dracochborzlib	1.14	0.68	0.28	1.04	0.24	1.14	0.16	0.17

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudn3d	Zurich
editattrall	dracoreplace	0.93	0.81	0.56	0.69	0.53	1	0.29	0.12
editattrall	dracozlib	1.08	0.97	0.36	0.93	0.34	1.02	0.25	0.52
editattrall	original	1	1	1	1	1	1	1	1
editattrall	originalchbor	1.01	0.88	0.7	0.71	0.73	0.69	0.65	0.79
editattrall	originalborreplace	1.23	0.73	0.82	0.89	1.13	0.98	1.39	NA
editattrall	originalchborreplacezlib	1.56	1.67	1.27	0.58	0.89	1.82	1.75	2.4
editattrall	originalchborzlib	1.19	0.77	0.26	0.59	0.25	0.68	0.34	0.48
editattrall	originalchborreplace	1.28	0.8	1.39	0.63	1.34	0.95	1.65	1.59
editattrall	originalchborzlib	1.11	0.74	0.34	1.27	0.32	0.68	0.57	0.37
editattrone	draco	1.15	1.22	0.71	0.9	0.82	1.16	0.35	0.6
editattrone	dracochbor	1.44	1.27	0.64	1	0.54	1.05	0.27	0.41
editattrone	dracocborreplace	1.28	0.88	0.24	1.06	0.1	0.97	0.17	0.08
editattrone	dracochborreplacezlib	1.32	1.01	0.27	0.68	0.09	1.32	0.17	0.11
editattrone	dracochborzlib	1.47	1.68	0.4	1.05	0.55	1.17	0.2	0.27
editattrone	dracoreplace	0.81	0.62	0.55	0.47	0.07	0.54	0.12	0.08
editattrone	dracozlib	1.6	1.16	0.38	1.04	0.34	1.18	0.26	0.44
editattrone	original	1	1	1	1	1	1	1	1
editattrone	originalchbor	0.95	0.98	0.8	0.75	0.82	0.83	0.76	1.31
editattrone	originalchborreplace	1.56	0.87	0.84	1.13	1	1.17	1.5	1.03
editattrone	originalchborreplacezlib	2.32	1.66	1.39	0.66	0.73	2.25	2.37	0.88
editattrone	originalchborzlib	1.95	1.82	0.56	0.94	0.6	0.93	0.95	NA
editattrone	originalchborreplace	1.64	0.94	1.44	0.71	1.22	1.01	1.64	0.06
editattrone	originalchborzlib	1.73	1.02	0.73	1.99	0.65	1.99	1.22	NA
editgeomall	draco	1.77	1.04	0.34	0.45	0.32	1.27	0.19	0.36
editgeomall	dracochbor	1.56	1.03	0.3	0.45	0.29	1.22	0.19	0.33
editgeomall	dracocborreplace	1.5	0.97	0.21	0.37	0.19	1.11	0.17	0.24
editgeomall	dracochborreplacezlib	1.51	0.96	0.21	0.37	0.19	1.14	0.16	0.24
editgeomall	dracochborzlib	1.61	0.95	0.18	0.4	0.18	1.18	0.15	0.07
editgeomall	dracoreplace	1.51	0.94	0.23	0.35	0.15	1.53	0.19	0.2
editgeomall	dracozlib	1.55	0.95	0.21	0.4	0.21	1.01	0.16	0.25
editgeomall	original	1	1	1	1	1	1	1	1
editgeomall	originalchbor	1.13	0.68	0.24	0.24	0.2	0.6	0.19	0.32
editgeomall	originalchborreplace	0.7	0.49	0.24	0.33	0.23	0.81	0.38	0.32

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudml3d	Zurich
editgeomall	originalcborreplacezlib	1.35	1.02	0.25	0.22	0.19	1.08	0.71	NA
editgeomall	originalcborzlib	0.87	0.57	0.17	0.31	0.15	0.68	0.24	NA
editgeomall	originalreplace	0.75	0.71	0.36	0.33	0.29	0.79	0.45	NA
editgeomall	originalzlib	0.87	0.6	0.22	0.51	0.16	0.71	0.39	0.28
editgeomone	draco	1.61	1.78	1.02	1.17	1.14	1.4	0.66	0.99
editgeomone	dracochbor	1.67	1.58	0.9	1.11	1.01	1.32	0.64	0.89
editgeomone	dracocborreplace	1.6	1.39	0.6	0.89	0.62	1.18	0.55	0.6
editgeomone	dracocborreplacezlib	1.64	1.42	0.61	0.89	0.61	1.21	0.53	0.58
editgeomone	dracocborzlib	1.63	1.45	0.49	1.02	0.55	1.23	0.48	0.57
editgeomone	dracoreplace	1.62	1.38	0.66	0.85	0.44	1.4	0.58	0.45
editgeomone	dracozlib	1.63	1.43	0.57	1.02	0.67	1.26	0.53	0.64
editgeomone	original	1	1	1	1	1	1	1	1
editgeomone	originalcbor	0.95	1.22	0.8	0.75	0.81	0.78	0.74	1.04
editgeomone	originalcborreplace	0.9	0.88	0.81	0.93	0.86	1.08	1.45	1.25
editgeomone	originalcborreplacezlib	2.13	1.55	1.49	0.66	0.78	2.61	1.99	NA
editgeomone	originalcborzlib	1.19	1.03	0.57	0.92	0.59	0.89	0.94	NA
editgeomone	originalreplace	0.97	0.98	1.21	0.92	1.22	1.13	1.82	1.32
editgeomone	originalzlib	1.12	1	0.7	2.11	0.65	0.9	3.56	0.85
queryall	draco	0.94	1.12	0.75	0.94	0.8	0.93	0.37	0.67
queryall	dracochbor	0.95	1.1	0.63	0.84	0.67	0.86	0.35	0.58
queryall	dracocborreplace	0.9	0.92	0.34	0.63	0.28	0.74	0.26	0.28
queryall	dracocborreplacezlib	0.91	0.94	0.34	0.64	0.28	0.74	0.25	0.27
queryall	dracocborzlib	0.93	0.95	0.23	0.77	0.2	0.77	0.19	0.22
queryall	dracoreplace	0.89	0.9	0.39	0.57	0.1	0.68	0.17	0.13
queryall	dracozlib	0.91	0.96	0.31	0.77	0.33	0.81	0.24	0.32
queryall	original	1	1	1	1	1	1	1	1
queryall	originalcbor	0.77	1.11	0.72	0.74	0.74	0.68	0.67	0.83
queryall	originalcborreplace	0.94	0.79	0.66	0.76	0.67	0.93	1	0.84
queryall	originalcborreplacezlib	1.22	1.31	1.23	0.6	0.59	1.43	0.98	0.98
queryall	originalcborzlib	1.13	0.77	0.28	0.65	0.27	0.6	0.35	0.45
queryall	originalreplace	1.1	0.94	1.24	0.7	1.03	0.92	1.53	1.33
queryall	originalzlib	0.77	0.8	0.36	1.13	0.32	0.63	0.57	0.38
queryone	draco	1.33	1.42	1.03	1.06	1.13	1.24	0.66	0.98

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudn3d	Zurich
queryone	dracochbor	1.28	1.43	0.9	0.98	0.99	1.19	0.64	0.88
queryone	dracocborreplace	1.21	1.24	0.6	0.8	0.61	1.06	0.55	0.59
queryone	dracocborreplacezlib	1.22	1.25	0.6	0.78	0.6	1.05	0.54	0.58
queryone	dracocborzlib	1.26	1.28	0.48	0.9	0.53	1.09	0.48	0.59
queryone	dracoreplace	1.23	1.23	0.66	0.76	0.43	0.99	0.56	0.49
queryone	dracozlib	1.26	1.28	0.58	0.91	0.66	1.12	0.54	0.63
queryone	original	1	1	1	1	1	1	1	1
queryone	originalbor	0.76	1.11	0.74	0.68	0.74	0.7	0.69	0.82
queryone	originalborreplace	0.84	0.8	0.68	0.7	0.66	0.91	1.01	0.81
queryone	originalborreplacezlib	1.23	1.55	1.34	0.56	0.6	0.91	0.99	0.99
queryone	originalborzlib	0.81	0.76	0.29	0.61	0.26	0.63	0.35	0.5
queryone	originalreplace	0.83	0.93	1.17	0.65	1.08	0.95	1.57	1.32
queryone	originalzlib	0.78	0.79	0.35	1.16	0.32	0.69	0.59	0.38
visualise	draco	1.83	1.74	0.86	1	1.01	1.54	0.31	0.86
visualise	dracochbor	1.96	1.29	0.78	0.92	0.75	1.35	0.44	0.59
visualise	dracocborreplace	1.99	1.19	0.48	1.19	0.55	1.67	0.27	0.06
visualise	dracocborreplacezlib	1.98	1.31	0.47	1.2	0.41	1.96	0.28	0.44
visualise	dracocborzlib	2.02	2.05	0.61	0.8	0.79	2.49	0.31	0.42
visualise	dracoreplace	1.5	1.2	0.71	1.16	0.53	2.46	0.32	0.37
visualise	dracozlib	2.12	1.3	0.47	0.92	0.56	1.14	0.24	0.3
visualise	original	1	1	1	1	1	1	1	1
visualise	originalbor	0.86	1.1	0.68	0.68	0.76	0.69	0.61	0.99
visualise	originalborreplace	0.92	0.79	0.63	0.84	0.84	1	1.16	0.94
visualise	originalborreplacezlib	0.98	1.52	1.01	0.6	0.64	2.37	1.23	1.65
visualise	originalborzlib	1.67	0.8	0.33	0.64	0.36	0.66	0.54	0.46
visualise	originalreplace	1.52	0.95	1.25	0.68	1.19	0.98	1.24	1.58
visualise	originalzlib	1.41	0.83	0.4	1.33	0.62	0.7	0.56	0.63

Table D.1: Full benchmark results of compression in advance, datasets with attributes

task	method	delft	denhaag	hdlb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
bufferall	draco	NA	1.71	0.26	0.78	NA	1.54	NA	NA
bufferall	dracochbor	NA	1.81	0.22	0.81	NA	1.86	NA	NA
bufferall	dracochborreplace	NA	1.55	0.32	0.73	NA	1.47	NA	NA
bufferall	dracochborreplacezlib	NA	1.57	0.13	0.93	NA	1.94	NA	NA
bufferall	dracochborzlib	NA	1.74	0.1	0.86	NA	2.84	NA	NA
bufferall	dracoreplace	NA	1.74	0.13	0.77	NA	1.53	NA	NA
bufferall	dracozlib	NA	1.58	0.12	0.91	NA	1.64	NA	NA
bufferall	original	NA	1	1	NA	1	NA	NA	NA
bufferall	originalcbor	NA	0.83	0.88	0.74	NA	1.14	NA	NA
bufferall	originalcborreplace	NA	1.14	0.78	0.67	NA	0.69	NA	NA
bufferall	originalcborreplacezlib	NA	1.23	0.85	2.43	NA	3.33	NA	NA
bufferall	originalcborzlib	NA	0.97	0.8	0.62	NA	0.92	NA	NA
bufferall	originalreplace	NA	1.25	1.14	1.55	NA	0.63	NA	NA
bufferall	originalzlib	NA	0.82	0.7	0.63	NA	1.05	NA	NA
bufferone	draco	1.91	1.7	1.04	1.42	1.01	1.42	1.53	0.99
bufferone	dracochbor	1.8	1.9	0.82	1.66	0.9	1.37	1.58	0.99
bufferone	dracochborreplace	2.01	1.55	0.59	1	0.58	1.47	1.39	0.59
bufferone	dracochborreplacezlib	1.92	1.49	0.63	1.17	0.54	1.33	1.44	0.6
bufferone	dracochborzlib	3.21	1.61	0.57	1.34	0.51	1.53	1.15	0.5
bufferone	dracoreplace	2.04	2.16	0.73	1.17	0.41	1.27	0.93	0.45
bufferone	dracozlib	1.85	1.6	0.65	1.14	0.6	1.32	1.25	0.6
bufferone	original	1	1	1	1	1	1	1	1
bufferone	originalcbor	1.08	0.94	0.76	0.87	0.67	0.83	1.6	0.84
bufferone	originalcborreplace	1	1.06	0.71	0.65	0.49	0.77	1.66	0.71
bufferone	originalcborreplacezlib	0.96	1.02	0.65	0.66	0.5	0.82	1.85	0.65
bufferone	originalcborzlib	0.9	0.82	0.46	0.93	0.28	0.76	1.01	0.41
bufferone	originalreplace	1.42	1.09	1.12	0.87	0.81	0.76	2.66	1.14
bufferone	originalzlib	0.93	0.9	1.58	0.68	0.36	0.81	1.26	0.48
editattrall	draco	1	0.89	0.24	0.88	0.55	0.93	0.13	0.41
editattrall	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
editattrall	dracochborreplace	0.96	0.8	0.43	0.67	0.29	0.77	0.15	0.08
editattrall	dracochborreplacezlib	0.97	0.83	0.24	0.59	0.08	0.77	0.11	0.09
editattrall	dracochborzlib	0.95	0.81	0.23	0.76	0.13	0.82	0.11	0.17

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
editattrall	dracoreplace	1.1	0.82	0.43	0.63	0.29	0.79	0.19	0.12
editattrall	dracozlib	1.13	0.95	0.25	1.01	0.23	1.01	0.12	0.22
editattrall	original	1	1	1	1	1	1	1	1
editattrall	originalchbor	1.12	0.94	0.69	0.81	0.73	1.04	0.7	0.72
editattrall	originalcorreplace	1.06	0.84	1.35	0.63	0.75	0.72	1.14	NA
editattrall	originalcborreplacezlib	0.83	0.77	1.29	0.62	0.75	0.75	1.13	NA
editattrall	originalcborzlib	1.06	0.92	0.47	0.76	0.29	0.74	0.41	0.33
editattrall	originalchborreplace	0.91	0.83	1.83	0.73	1.03	0.77	1.81	NA
editattrall	originalchborzlib	1.03	0.8	0.49	0.72	0.34	0.75	0.47	0.51
editattrone	draco	1.21	1.08	0.26	0.94	0.56	0.96	0.13	0.41
editattrone	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
editattrone	dracocborreplace	1.23	0.96	0.32	0.71	0.1	0.8	0.12	0.09
editattrone	dracocborreplacezlib	1.23	1	0.26	0.64	0.08	0.8	0.11	0.08
editattrone	dracocborzlib	1.24	1.03	0.26	0.86	0.17	0.91	0.12	0.23
editattrone	dracoreplace	0.78	1	0.28	0.67	0.1	0.87	0.14	0.12
editattrone	dracozlib	1.58	1.16	0.27	1.06	0.23	1.05	0.12	0.22
editattrone	original	1	1	1	1	1	1	1	1
editattrone	originalchbor	1.04	1.04	0.85	1.06	0.81	1.12	2.02	0.8
editattrone	originalcorreplace	1.22	1	1.37	0.7	0.67	0.76	1.11	0.91
editattrone	originalcborreplacezlib	1.07	0.96	1.31	0.69	0.66	0.77	1.14	0.89
editattrone	originalcborzlib	1.47	1.42	0.93	1.08	0.62	0.95	1.09	0.73
editattrone	originalchborreplace	1.25	1.02	1.89	0.78	0.96	0.8	1.67	1.65
editattrone	originalchborzlib	1.2	1.06	1.05	1.13	0.64	0.94	1.08	0.75
editgeomall	draco	2.45	2.08	1.16	1.63	1.11	1.94	0.67	0.41
editgeomall	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
editgeomall	dracocborreplace	2.48	1.96	1.8	2.09	0.91	1.8	0.7	0.6
editgeomall	dracocborreplacezlib	2.55	1.99	1.17	1.44	0.65	1.83	0.68	0.6
editgeomall	dracocborzlib	2.41	1.92	1.15	1.46	0.7	1.77	0.8	0.15
editgeomall	dracoreplace	3.95	3	1.7	1.71	0.91	2.57	0.99	0.63
editgeomall	dracozlib	3.36	2.01	1.18	1.51	0.71	1.84	0.65	0.13
editgeomall	original	1	1	1	1	1	1	1	1
editgeomall	originalchbor	NA	NA	NA	NA	NA	NA	NA	NA
editgeomall	originalcorreplace	1.04	0.98	1.32	0.67	0.67	0.79	1.15	0.89

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudml3d	Zurich
editgeomall	originalcborreplacezlib	1.05	0.92	1.36	0.66	0.68	0.81	1.15	0.89
editgeomall	originalcborzlib	1.82	1.43	0.86	1.1	0.61	1.02	1.05	0.72
editgeomall	originalreplace	1.15	0.96	1.86	0.78	0.95	0.91	1.72	1.21
editgeomall	originalzlib	1.22	1.05	1.04	1.15	0.65	0.97	1.09	0.82
editgeomone	draco	1.86	1.65	0.86	1.31	0.91	1.48	0.48	0.72
editgeomone	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
editgeomone	dracochborreplace	2.05	1.58	1.14	1.4	0.59	1.27	0.48	0.4
editgeomone	dracochborreplacezlib	1.99	1.64	0.88	1.04	0.44	1.32	0.49	0.4
editgeomone	dracocborzlib	2.05	1.67	0.83	1.23	0.5	1.38	0.66	0.13
editgeomone	dracoreplace	3.38	1.97	1.08	1.48	0.56	1.67	0.66	0.43
editgeomone	dracozplib	2.11	1.61	0.85	1.19	0.51	1.38	0.46	0.45
editgeomone	original	1	1	1	1	1	1	1	1
editgeomone	originalcbor	1.12	1.03	0.82	1.01	0.93	1.11	0.75	1.12
editgeomone	originalcborreplace	1.11	1.04	1.32	0.67	0.67	0.74	1.11	0.89
editgeomone	originalcborreplacezlib	1.1	0.94	1.32	0.69	0.67	0.74	1.1	0.89
editgeomone	originalcborzlib	1.61	1.53	0.89	1.14	0.61	0.96	1.03	0.71
editgeomone	originalreplace	1.14	0.99	1.86	0.78	0.95	0.91	2.65	1.22
editgeomone	originalzlib	1.26	1.05	1	1.14	0.64	0.92	1.04	1.08
queryall	draco	1.2	1.06	0.26	0.94	0.56	1.04	0.13	0.41
queryall	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
queryall	dracochborreplace	1.2	0.92	0.32	0.7	0.09	0.83	0.11	0.08
queryall	dracochborreplacezlib	1.23	0.96	0.25	0.64	0.08	0.85	0.11	0.07
queryall	dracocborzlib	1.23	0.97	0.25	0.82	0.13	0.94	0.1	0.16
queryall	dracoreplace	1.29	1.07	0.3	0.66	0.1	0.9	0.13	0.13
queryall	dracozplib	1.36	0.98	0.25	0.83	0.15	0.92	0.1	0.13
queryall	original	1	1	1	1	1	1	1	1
queryall	originalcbor	1.18	1.05	0.86	0.91	0.84	1.09	0.68	0.73
queryall	originalcborreplace	1.01	0.94	1.14	0.6	0.54	0.75	0.84	0.67
queryall	originalcborreplacezlib	0.99	0.85	1.12	0.6	0.54	0.76	0.84	0.68
queryall	originalcborzlib	1.36	0.89	0.49	0.77	0.29	0.81	0.42	0.33
queryall	originalreplace	1.26	0.96	1.77	0.75	0.91	0.87	1.48	1.16
queryall	originalzlib	1.03	0.82	0.53	0.76	0.34	0.76	0.47	0.43
queryone	draco	1.76	1.53	0.82	1.22	0.9	1.53	0.48	0.71

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudn3d	Zurich
queryone	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
queryone	dracocborreplace	1.81	1.43	1.26	1.29	0.57	1.34	0.48	0.4
queryone	dracochborreplacezlib	1.8	1.49	0.83	0.97	0.43	1.43	0.47	0.39
queryone	dracocborzlib	1.78	1.42	0.8	1.12	0.48	1.45	0.53	0.52
queryone	dracoreplace	2.71	1.9	1.16	1.38	0.6	1.83	0.66	0.54
queryone	dracozlib	2.71	1.47	0.78	1.13	0.49	1.46	0.46	0.44
queryone	original	1	1	1	1	1	1	1	1
queryone	originalbor	1	0.96	0.8	0.95	0.8	0.94	0.7	0.71
queryone	originalborreplace	1	0.91	1.15	0.6	0.53	0.78	0.86	0.66
queryone	originalborreplacezlib	1.02	0.86	1.11	0.62	0.54	0.78	0.85	0.66
queryone	originalborzlib	1.07	0.9	0.48	0.78	0.3	0.81	0.42	0.33
queryone	originalreplace	1.26	0.96	1.77	0.75	0.91	1.03	1.64	1.14
queryone	originalzlib	1.07	0.84	0.54	0.74	0.34	0.79	0.48	0.43
visualise	draco	1.77	1.28	0.68	1.14	0.79	1.33	0.25	0.61
visualise	dracochbor	NA	NA	NA	NA	NA	NA	NA	NA
visualise	dracocborreplace	1.73	1.18	0.77	1.2	0.53	1.68	0.24	0.35
visualise	dracocborreplacezlib	1.73	1.26	0.68	0.89	0.39	1.2	0.25	0.34
visualise	dracocborzlib	1.7	1.22	0.7	1.03	0.43	1.24	0.23	0.13
visualise	dracoreplace	2.46	1.49	0.77	1.01	0.56	1.52	0.36	0.1
visualise	dracozlib	1.96	1.3	0.68	0.99	0.44	1.25	0.23	0.39
visualise	original	1	1	1	1	1	1	1	1
visualise	originalbor	1.06	1.07	0.79	1.01	0.77	1.08	0.88	0.74
visualise	originalborreplace	1.4	0.92	1.15	0.63	0.58	0.73	0.86	0.7
visualise	originalborreplacezlib	1.01	0.87	1.11	0.64	0.58	0.74	0.85	0.69
visualise	originalborzlib	1.17	0.98	0.69	0.83	0.39	0.81	0.67	0.44
visualise	originalreplace	1.26	1.06	1.61	0.78	0.91	0.81	1.45	1.2
visualise	originalzlib	1.05	0.87	0.64	0.78	0.43	0.76	0.7	0.44

Table D.2: Full benchmark results of compression in advance, datasets without attributes

task	method	delft	denhaag	hdbl	montreal	newyork	RotterdamLoD1LoD2	tucnl3d	Zurich
bufferall	draco	2.33	2.33	2.32	1.19	2.2	NA	2.76	NA
bufferall	dracochbor	2.34	2.43	2.89	1.28	2.29	2.18	2.68	NA
bufferall	dracochborreplace	1.14	0.17	NA	1.81	NA	6.87	NA	NA
bufferall	dracochborreplacezlib	1.09	0.79	8.69	1.88	0.57	7.9	0.75	NA
bufferall	dracochorzlib	2.37	2.48	2.15	1.28	2.11	2.19	2.51	NA
bufferall	dracoreplace	1.02	0.78	7.78	0.83	0.65	7.21	1	NA
bufferall	dracozlib	2.28	2.38	0.43	1.1	0.41	2	0.53	NA
bufferall	original	1	1	1	1	1	1	1	NA
bufferall	originalcbor	0.84	0.76	0.76	0.97	0.81	0.68	0.69	NA
bufferall	originalcborreplace	1.02	NA	NA	NA	NA	NA	NA	NA
bufferall	originalcborreplacezlib	1.3	0.93	0.93	0.48	0.72	0.66	0.98	NA
bufferall	originalcborzlib	0.85	0.72	0.63	0.78	0.71	0.66	0.68	NA
bufferall	originalreplace	1.15	1.06	1.14	0.53	0.87	0.7	1.15	NA
bufferall	originalzlib	0.76	0.57	0.41	0.58	0.39	0.53	0.53	NA
bufferone	draco	1.45	1.4	1.39	1.38	1.37	NA	NA	NA
bufferone	dracochbor	1.5	1.4	1.39	1.39	1.4	1.43	1.53	NA
bufferone	dracochborreplace	1.62	NA	1.42	1.5	1.42	NA	1.41	NA
bufferone	dracochborreplacezlib	1.64	1.96	1.47	1.47	1.43	1.72	1.41	NA
bufferone	dracochorzlib	1.46	1.44	1.49	1.44	1.5	1.46	1.5	NA
bufferone	dracoreplace	1.53	1.61	1.38	1.46	1.38	1.56	1.42	NA
bufferone	dracozlib	1.48	1.43	1.41	1.4	1.41	1.48	1.44	0.05
bufferone	original	1	1	1	1	1	1	1	1
bufferone	originalcbor	1	0.98	0.98	0.98	1.05	1	0.98	0.84
bufferone	originalcborreplace	0.99	NA	NA	NA	NA	NA	NA	NA
bufferone	originalcborreplacezlib	1.05	0.99	1.02	0.99	1.01	1.01	1.03	NA
bufferone	originalcborzlib	1.05	0.98	1	0.99	1.13	1.02	1	0.41
bufferone	originalreplace	1.06	0.98	1.02	1.02	1.01	1.03	0.99	1.24
bufferone	originalzlib	1.13	1	0.98	0.98	0.99	0.99	1	0.47
editattrall	draco	1.36	1.2	39.54	1.03	22.35	NA	2.13	14.07
editattrall	dracochbor	1.38	1.29	41.66	1.09	52.04	1.31	3.76	23.83
editattrall	dracochborreplace	1.48	NA	42.83	0.76	25.05	NA	23.59	29.39
editattrall	dracochborreplacezlib	1.46	1.53	47.35	0.76	1.69	1.37	1.69	32.21
editattrall	dracochorzlib	1.41	1.27	30.37	1.12	38.75	1.36	3.8	27.1

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
editattrall	dracoreplace	1.41	1.27	41.99	0.79	1.98	1.35	2.86	27.5
editattrall	dracozlib	1.39	1.23	16.13	0.96	20.95	1.29	2.86	17.69
editattrall	original	1	1	1	1	1	1	1	1
editattrall	originalcbor	1.02	1	0.81	0.91	0.92	0.96	0.84	0.8
editattrall	originalcborreplace	1.06	NA	NA	NA	NA	NA	NA	NA
editattrall	originalcborreplacezlib	1.06	1.08	1.57	0.77	1.66	1.17	1.73	NA
editattrall	originalcborzlib	1.01	0.95	0.66	0.95	0.7	1.04	0.85	0.47
editattrall	originalcborreplace	1.12	1.09	2.11	0.81	2.1	1.27	2.15	1.57
editattrall	originalzlib	1.13	0.87	0.56	0.72	0.43	0.94	0.52	0.38
editattrone	draco	1.15	2.96	4.38	0.72	5.48	NA	NA	9.68
editattrone	dracochbor	1.17	3.19	4.6	0.8	5.86	4.83	0.46	NA
editattrone	dracocborreplace	1.39	0.41	NA	1.53	NA	10.16	NA	NA
editattrone	dracocborreplacezlib	1.37	1.29	13.97	1.63	0.42	11.57	0.42	NA
editattrone	dracocborzlib	1.18	3.19	5.61	0.82	6.34	4.95	0.44	NA
editattrone	dracoreplace	1.28	1.22	12.4	0.65	0.5	10.72	0.76	NA
editattrone	dracozlib	13.25	2.93	4.51	0.64	5.47	4.52	0.36	10.22
editattrone	original	1	1	1	1	1	1	1	1
editattrone	originalcbor	0.94	0.91	0.86	0.82	0.88	0.83	0.77	1.31
editattrone	originalcborreplace	1	NA	NA	NA	NA	NA	NA	NA
editattrone	originalcborreplacezlib	1.14	0.87	0.96	0.55	0.7	0.71	1.05	NA
editattrone	originalcborzlib	0.96	0.8	0.57	0.77	0.56	0.8	0.7	NA
editattrone	originalreplace	1.14	1	1.28	0.68	0.87	0.8	1.37	0.6
editattrone	originalzlib	1.11	0.62	0.33	0.52	0.26	0.59	0.36	NA
editattrone	draco	1.15	2.97	4.78	0.69	5.61	NA	0.41	3.55
editgeomall	dracochbor	1.18	3.07	NA	0.8	NA	4.82	NA	NA
editgeomall	dracocborreplace	1.41	0.4	NA	1.56	NA	10.76	NA	NA
editgeomall	dracocborreplacezlib	1.36	1.23	NA	1.6	NA	11.91	NA	NA
editgeomall	dracocborzlib	1.2	3.01	NA	0.75	NA	4.91	NA	NA
editgeomall	dracoreplace	1.3	1.23	NA	0.66	NA	10.89	NA	NA
editgeomall	dracozlib	1.16	2.86	4.7	0.64	5.42	4.61	0.33	3.29
editgeomall	original	1	1	1	1	1	1	1	1
editgeomall	originalcbor	0.94	0.9	0.85	0.77	0.82	0.84	1.29	0.32
editgeomall	originalcborreplace	1.01	NA	NA	NA	NA	NA	NA	NA

(continued)

task	method	delft	denhaag	hdlb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
editgeomall	originalcborreplacezlib	1.15	0.86	1.01	0.53	0.7	0.74	1.03	NA
editgeomall	originalcborzlib	0.93	0.79	0.58	0.74	0.68	0.81	0.7	NA
editgeomall	originalahreplace	1.1	1	1.35	0.65	0.94	0.81	1.34	NA
editgeomall	originalalzlib	0.86	0.62	0.29	0.52	0.26	0.63	0.34	0.27
editgeomone	draco	1.12	3.07	4.44	0.71	5.54	NA	NA	10.6
editgeomone	dracochbor	1.16	3.08	11.7	0.8	13.23	4.65	0.98	NA
editgeomone	dracochborreplace	1.39	0.4	NA	1.59	NA	10.27	NA	NA
editgeomone	dracochborreplacezlib	1.32	1.28	13.3	1.65	0.33	11.5	0.36	NA
editgeomone	dracocborzlib	1.16	3.1	9.02	0.8	9.95	4.78	0.96	NA
editgeomone	dracoreplace	1.25	1.23	11.89	0.67	0.42	10.6	0.65	NA
editgeomone	dracozlib	1.1	2.92	4.33	0.66	5.39	4.66	0.33	12.44
editgeomone	original	1	1	1	1	1	1	1	1
editgeomone	originalcbor	0.93	0.89	0.81	0.79	0.85	0.82	0.76	1.04
editgeomone	originalcborreplace	0.99	NA	NA	NA	NA	NA	NA	NA
editgeomone	originalcborreplacezlib	1.24	0.88	0.93	0.59	1.06	0.72	1	NA
editgeomone	originalcborzlib	0.94	0.8	0.58	0.81	0.68	0.78	0.68	NA
editgeomone	originalreplace	1.1	0.95	1.25	0.67	0.91	0.85	1.31	1.29
editgeomone	originalalzlib	1.08	0.63	0.34	0.52	0.27	0.6	0.35	0.85
dracochbor	draco	1.19	2.78	3.98	0.73	4.53	NA	0.83	31.26
dracochbor	dracochbor	1.18	2.88	9.38	0.84	10.29	4.68	3.04	44.22
dracochbor	dracochborreplace	1.41	0.36	10.07	1.54	5.09	9.74	16.45	52.75
dracochbor	dracochborreplacezlib	1.34	1.26	11.08	1.6	0.56	10.97	1.6	59.56
dracochbor	dracocborzlib	1.15	2.95	7.35	0.85	7.58	4.45	3.02	47.17
dracochbor	dracoreplace	1.3	1.28	9.85	0.73	0.62	10.11	2.4	47.98
dracochbor	dracozlib	1.16	2.68	3.79	0.68	4.38	4.3	0.82	30.48
originalcbor	original	1	1	1	1	1	1	1	1
originalcbor	originalcbor	0.92	0.91	0.88	0.84	0.89	0.87	0.92	0.83
originalcbor	originalcborreplace	1.04	NA	NA	NA	NA	NA	NA	NA
originalcbor	originalcborreplacezlib	1.17	0.9	1.02	0.65	0.77	0.76	1.05	NA
originalcbor	originalcborzlib	0.99	0.82	0.67	0.82	0.68	0.8	0.85	0.45
originalcbor	originalreplace	1.12	0.99	1.21	0.72	0.94	0.79	1.2	1.31
originalcbor	originalalzlib	0.91	0.7	0.48	0.59	0.46	0.64	0.72	0.38
originalcbor	draco	1.4	1.43	1.21	1.28	1.14	NA	NA	0.17

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudnl3d	Zurich
queryone	dracochbor	1.42	1.42	1.31	1.29	1.17	1.05	1.46	NA
queryone	dracocborreplace	1.56	NA	1.27	1.36	1.14	NA	1.18	NA
queryone	dracocborreplacezlib	1.56	1.88	1.49	1.39	1.27	1.63	1.16	NA
queryone	dracocborzlib	1.44	1.43	1.31	1.32	1.26	1.37	1.3	NA
queryone	dracoreplace	1.46	1.61	1.28	1.35	1.22	1.45	1.23	NA
queryone	dracozlib	1.38	1.44	1.3	1.28	1.24	1.38	1.23	0.19
queryone	original	1	1	1	1	1	1	1	1
queryone	originalcbor	0.95	0.99	1	0.98	1	1	1.01	0.84
queryone	originalcborreplaced	0.96	NA	NA	NA	NA	NA	NA	NA
queryone	originalcborreplacedzlib	1.08	1.05	1.07	0.97	1.08	0.99	0.99	NA
queryone	originalcborzlib	1	1	0.99	0.94	1.33	0.98	1.03	0.49
queryone	originalreplace	0.99	1.06	1	0.96	1.03	1.01	1.01	1.33
queryone	originalreplacezlib	1.01	1.03	1.03	0.96	1.04	0.98	1	0.38
visualise	draco	NA	NA	5.83	NA	5.33	NA	0.46	9.03
visualise	dracochbor	2.37	3.03	11.03	0.96	11.9	4.62	0.86	19.17
visualise	dracocborreplace	1.69	0.35	11.35	1.68	5.99	9.58	4.48	23.7
visualise	dracocborreplacezlib	1.67	1.39	12.34	1.66	0.58	11.09	0.46	26.5
visualise	dracocborzlib	1.48	3.05	8.32	0.92	9.09	4.68	0.85	21.66
visualise	dracoreplace	1.56	1.39	11.36	0.84	0.66	9.98	0.67	22.53
visualise	dracozlib	1.41	2.74	4.16	0.79	5.11	4.47	0.4	8.65
visualise	original	1	1	1	1	1	1	1	1
visualise	originalcbor	0.94	0.89	0.83	0.95	0.89	0.83	0.83	0.99
visualise	originalcborreplaced	0.99	NA	NA	NA	NA	NA	NA	NA
visualise	originalcborreplacedzlib	1.01	0.87	0.97	0.6	0.73	0.74	1.01	NA
visualise	originalcborzlib	0.96	0.79	0.62	0.79	0.61	0.77	0.79	0.57
visualise	originalreplace	1.11	0.96	1.25	0.71	0.92	0.77	1.26	1.57
visualise	originalzlib	0.87	0.7	0.36	0.57	0.34	0.65	0.55	0.61

Table D.3: Full benchmark results of compression on the fly, datasets with attributes

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudn13d	Zurich
bufferall	draco	2.38	2.43	2.94	1.22	2.16	NA	2.84	NA
bufferall	dracochbor	2.34	2.59	2.52	1.98	2.2	2.82	NA	NA
bufferall	dracochborreplace	1.08	1.04	NA	1	NA	3.42	NA	NA
bufferall	dracochborreplacezlib	1.09	1.66	4.72	0.93	5.04	4.11	0.94	NA
bufferall	dracochorzlib	2.4	2.43	2.9	1.36	2.14	2.35	2.82	NA
bufferall	dracoreplace	1.04	1.63	4.71	0.88	5.04	4.06	0.9	NA
bufferall	dracozlib	2.3	2.38	0.5	1.16	0.42	2.1	0.56	NA
bufferall	original	1	1	1	1	1	1	1	NA
bufferall	originalalchbor	0.85	0.76	0.83	0.81	0.77	0.71	0.7	NA
bufferall	originalalcorreplace	1.04	NA	NA	NA	NA	NA	NA	NA
bufferall	originalalcorreplacezlib	1.12	0.95	1	0.55	0.73	0.68	1.04	NA
bufferall	originalalczlib	0.84	0.76	0.7	0.85	0.64	0.69	0.72	NA
bufferall	originalalreplace	1.13	1.07	1.17	0.53	0.92	0.77	1.19	NA
bufferall	originalalzlib	0.77	0.56	0.56	0.6	0.41	0.6	0.54	NA
bufferone	draco	1.43	1.38	1.4	1.42	1.37	NA	NA	-2304.48
bufferone	dracochbor	1.47	1.41	1.46	1.4	1.48	1.51	1.54	NA
bufferone	dracochborreplace	1.6	NA	1.44	1.65	1.43	NA	1.41	NA
bufferone	dracochborreplacezlib	1.61	1.92	1.53	1.47	1.38	1.72	1.44	NA
bufferone	dracochorzlib	1.46	1.4	1.6	1.4	1.52	1.45	1.46	NA
bufferone	dracoreplace	1.5	1.65	1.42	1.42	1.39	1.55	1.41	NA
bufferone	dracozlib	1.42	1.38	1.43	1.43	1.53	1.46	1.41	-2527.15
bufferone	original	1	1	1	1	1	1	1	NA
bufferone	originalalchbor	0.97	0.99	1.06	0.98	1.06	0.99	1.07	NA
bufferone	originalalcorreplace	1.02	NA	NA	NA	NA	NA	NA	NA
bufferone	originalalcorreplacezlib	1.09	0.99	1.01	0.96	1.01	1.02	1.02	NA
bufferone	originalalczlib	0.99	0.95	1	0.97	1.06	1	1	NA
bufferone	originalalreplace	1.04	1	1	1	1.04	1.04	1.02	NA
bufferone	originalalzlib	1	0.99	1.01	0.95	1.04	1.01	1	NA
editattrall	draco	1.37	1.25	22.87	1	31.39	NA	2.36	0.33
editattrall	dracochbor	1.37	1.26	16.64	1.06	29.24	1.3	2.04	11.94
editattrall	dracochborreplace	1.5	NA	NA	0.9	NA	NA	NA	NA
editattrall	dracochborreplacezlib	1.43	1.54	16.95	0.76	28.33	1.39	2.06	11.2
editattrall	dracochorzlib	1.41	1.26	17.11	1.09	28.43	1.31	NA	11.53

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1Lod2	tudn13d	Zurich
editatfull	dracoreplace	1.42	1.3	16.76	0.77	28.54	1.37	1.97	11.45
editatfull	dracozlib	1.38	1.22	17.47	0.99	28.42	1.28	2.18	NA
editatfull	original	1	1	1	1	1	1	1	1
editatfull	originalalchor	1.01	0.97	0.75	0.89	0.89	0.95	1.01	0.73
editatfull	originalalchorreplace	1.08	NA	NA	NA	NA	NA	NA	NA
editatfull	originalalchorreplacezlib	1.29	1.08	1.59	0.78	1.6	1.21	1.7	NA
editatfull	originalalchorzlib	0.99	0.95	0.67	0.91	0.84	1.01	0.89	0.34
editatfull	originalalreplace	1.13	1.12	2.16	0.84	1.99	1.3	2.09	NA
editatfull	originalalzlib	1.02	0.85	0.51	0.74	0.41	0.92	0.51	0.52
editatrone	draco	1.14	3.14	8.1	0.69	5.18	NA	NA	9.22
editatrone	dracochbor	1.15	3.27	8.43	0.8	6.6	4.81	0.47	NA
editatrone	dracochborreplace	1.33	2.39	NA	0.69	NA	4.95	NA	NA
editatrone	dracochborreplacezlib	1.34	3.49	9.46	0.7	7.37	5.84	0.53	NA
editatrone	dracozlib	1.2	3.18	10.28	0.82	6.16	4.96	0.48	NA
editatrone	dracoreplace	1.25	3.24	9.72	0.68	7.26	5.69	0.53	NA
editatrone	dracozlib	1.18	2.96	8.76	0.63	5.11	4.75	0.35	11.97
editatrone	original	1	1	1	1	1	1	1	1
editatrone	originalalchor	0.92	0.9	0.75	0.83	0.83	0.86	0.87	0.79
editatrone	originalalchorreplace	1.03	NA	NA	NA	NA	NA	NA	NA
editatrone	originalalchorreplacezlib	1.17	0.94	1.28	0.57	0.68	0.73	1.04	NA
editatrone	originalalchorzlib	1	0.81	0.67	0.77	0.66	0.79	0.72	0.73
editatrone	originalalreplace	1.07	1	1.72	0.63	0.92	0.82	1.38	1.59
editatrone	originalalzlib	0.9	0.65	0.45	0.51	0.26	0.61	0.36	0.73
editgeomall	draco	1.19	3.01	8.61	0.7	5.89	NA	0.37	10.25
editgeomall	dracochbor	1.2	3.3	NA	0.79	NA	4.78	NA	NA
editgeomall	dracochborreplace	1.36	2.32	NA	0.76	NA	4.91	NA	NA
editgeomall	dracochborreplacezlib	1.36	3.35	NA	0.69	NA	5.74	NA	NA
editgeomall	dracozlib	1.2	3.19	NA	0.8	NA	4.86	NA	NA
editgeomall	dracoreplace	1.3	3.25	NA	0.71	NA	5.71	NA	NA
editgeomall	dracozlib	1.2	2.98	8.36	0.64	5.58	4.73	0.34	NA
editgeomall	original	1	1	1	1	1	1	1	1
editgeomall	originalalchor	0.97	0.94	0.83	0.82	0.85	0.84	0.76	NA
editgeomall	originalalchorreplace	1.05	NA	NA	NA	NA	NA	NA	NA

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1LoD2	tudn13d	Zurich
editgeomall	originalcborreplacezlib	1.2	0.89	1.24	0.56	0.74	0.75	1.06	NA
editgeomall	originalcborlib	0.99	0.82	0.65	0.75	0.66	0.78	0.75	0.72
editgeomall	originalreplace	1.18	0.97	1.64	0.66	0.91	0.8	1.43	1.21
editgeomall	originalzlib	0.91	0.66	0.44	0.52	0.27	0.62	0.37	0.81
editgeomone	draco	1.13	3.09	7.91	0.73	5.29	NA	NA	10.45
editgeomone	dracochbor	1.14	3.25	9.39	0.83	7.56	4.98	0.58	NA
editgeomone	dracochborreplace	1.34	2.4	NA	0.81	NA	5.04	NA	NA
editgeomone	dracochborreplacezlib	1.32	3.46	9.35	0.73	7.37	5.84	0.52	NA
editgeomone	dracocborzlib	1.19	3.22	9.26	0.81	7.63	4.99	0.55	NA
editgeomone	dracoreplace	1.24	3.34	9.11	0.7	7.26	5.78	0.51	NA
editgeomone	dracozlib	1.15	2.85	8.4	0.67	5.42	4.8	0.34	10.14
editgeomone	original	1	1	1	1	1	1	1	1
editgeomone	originalcbor	0.93	0.88	0.73	0.85	0.85	0.85	0.86	1.19
editgeomone	originalcborreplace	1.03	NA	NA	NA	NA	NA	NA	NA
editgeomone	originalcborreplacezlib	1.18	0.86	1.24	0.57	0.7	0.75	1	NA
editgeomone	originalcborlib	0.96	0.82	0.65	0.79	0.66	0.81	0.72	0.75
editgeomone	originalreplace	1.08	1	1.59	0.66	0.88	0.79	1.35	1.28
editgeomone	originalzlib	0.87	0.66	0.43	0.53	0.27	0.62	0.36	1.11
draco	draco	1.16	2.87	5.93	0.73	4.5	NA	0.56	28.93
draco	dracochbor	1.18	2.83	6.6	0.85	6.11	4.71	0.71	29.95
draco	dracochborreplace	1.35	2.12	NA	0.79	NA	4.63	NA	NA
draco	dracochborreplacezlib	1.34	3.12	6.75	0.75	5.88	5.57	0.7	31.65
draco	dracocborzlib	1.21	2.96	6.69	0.83	5.82	4.56	0.67	30.47
draco	dracoreplace	1.28	2.95	6.68	0.8	5.93	5.38	0.68	31.21
draco	dracozlib	1.17	2.64	6	0.69	4.41	4.49	0.55	NA
queryall	original	1	1	1	1	1	1	1	1
queryall	originalcbor	0.99	0.87	0.81	0.83	0.91	0.87	0.92	0.74
queryall	originalcborreplace	1.05	NA	NA	NA	NA	NA	NA	NA
queryall	originalcborreplacezlib	1.18	0.86	1.15	0.64	0.77	0.75	1.04	NA
queryall	originalcborlib	0.98	0.84	0.79	0.8	0.81	0.79	0.8	0.33
queryall	originalreplace	1.13	0.95	1.45	0.75	0.93	0.81	1.28	1.17
queryall	originalzlib	0.9	0.65	0.61	0.58	0.45	0.65	0.53	0.44
queryone	draco	1.41	1.41	1.28	1.3	1.11	NA	NA	0.16

(continued)

task	method	delft	denhaag	hdb	montreal	newyork	RotterdamLoD1Lod2	tudnl3d	Zurich
queryone	dracochbor	1.4	1.47	1.32	1.33	1.39	1.38	1.53	NA
queryone	dracochborreplace	1.55	NA	1.34	1.7	1.2	NA	1.21	NA
queryone	dracochborreplacezlib	1.58	1.94	1.47	1.39	1.23	1.67	1.27	NA
queryone	dracochborzlib	1.43	1.49	1.44	1.33	1.33	1.42	1.22	NA
queryone	dracoreplace	1.49	1.64	1.31	1.35	1.16	1.5	1.22	NA
queryone	dracozlib	1.39	1.47	1.32	1.34	1.14	1.39	1.22	0.17
queryone	original	1	1	1	1	1	1	1	1
queryone	originalalchbor	0.99	1.01	0.99	1.01	1.05	0.98	1.2	0.71
queryone	originalalcorreplace	0.98	NA	NA	NA	NA	NA	NA	NA
queryone	originalalcorreplacezlib	1.12	1.02	1.04	1.02	1.08	0.99	1.01	NA
queryone	originalalczlib	1	1.01	0.98	0.98	1.16	1.01	1.03	0.32
queryone	originalalreplace	1	1.03	1.01	0.98	1.05	1.03	1.04	1.16
queryone	originalalzlib	1	1.03	1.04	0.99	1.01	1.01	1	0.43
visualise	draco	NA	NA	7.58	NA	5.17	NA	0.45	9.11
visualise	dracochbor	1.5	3.04	8.59	0.98	6.99	4.87	0.57	10.24
visualise	dracochborreplace	1.67	2.09	NA	0.88	NA	4.55	NA	NA
visualise	dracochborreplacezlib	1.74	3.39	8.58	0.9	6.86	5.62	0.58	10
visualise	dracochborzlib	1.6	3.08	8.39	0.97	6.92	4.82	0.55	9.95
visualise	dracoreplace	1.65	3.13	8.54	0.86	6.93	5.59	0.57	9.94
visualise	dracozlib	1.56	2.98	7.53	0.82	5.13	4.65	0.43	NA
visualise	original	1	1	1	1	1	1	1	1
visualise	originalalchbor	0.93	0.9	0.79	0.92	0.85	0.88	0.99	0.75
visualise	originalalcorreplace	1.05	NA	NA	NA	NA	NA	NA	NA
visualise	originalalcorreplacezlib	1.3	0.95	1.24	0.64	0.72	0.77	1.05	NA
visualise	originalalczlib	0.98	0.81	0.73	0.83	0.74	0.83	0.8	0.44
visualise	originalalreplace	1.14	0.99	1.55	0.71	0.94	0.83	1.27	1.21
visualise	originalalzlib	0.91	0.7	0.52	0.59	0.34	0.64	0.57	0.44

Table D.4: Full benchmark results of compression on the fly, datasets without attributes

Bibliography

- 3D geoinformation research group at TU Delft. CityJSON/io: Python CLI to process and manipulate CityJSON files, 2019. URL <https://github.com/cityjson/cjio>.
- Analytical Graphics, Inc. 3D Tiles Format Specification, 2018. URL <https://github.com/AnalyticalGraphicsInc/3d-tiles/tree/master/specification#spatial-data-structures>.
- Analytical Graphics, Inc. Batched 3D Model, 2019. URL <https://github.com/AnalyticalGraphicsInc/3d-tiles/blob/master/specification/TileFormats/Batched3DModel/README.md>.
- F. Biljecki, J. Stoter, H. Ledoux, S. Zlatanova, and A. Çöltekin. Applications of 3D city models: State of the art review. *ISPRS International Journal of Geo-Information*, 4(4):2842–2889, 2015.
- F. Biljecki, H. Ledoux, and J. Stoter. Generating 3D city models without elevation data. *Computers, Environment and Urban Systems*, 64:1–18, July 2017. doi: 10.1016/j.compenvurbsys.2017.01.001.
- F. Boersma. Online viewer for CityJSON files., 2019. URL <https://github.com/fhb1990/CityJSON-viewer>.
- P. Bourke. RLE - Run Length Encoding, 1995. URL <https://web.archive.org/web/20031020232320/http://astronomy.swin.edu.au/~pbourke/dataformats/rle/>.
- BSON. BSON homepage, 2019. URL <http://bsonspec.org/>.
- Cable. Worldwide broadband speed league 2019, 2019. URL <https://www.cable.co.uk/broadband/speed/worldwide-speed-league/>.
- CBOR. Concise Binary Object Representation (CBOR), 2013. URL <https://tools.ietf.org/html/rfc7049>.
- CBOR. CBOR. RFC 7049 Concise Binary Object Representation, 2019. URL <https://cbor.io/>.
- CBOR. Implementations, 2020. URL <https://cbor.io/impls.html>.
- Cesium. Draco compressed meshes with glTF and 3D tiles, 2018. URL <https://cesium.com/blog/2018/04/09/draco-compression/>.
- Cesium. 3D Tiles Overview, 2020. URL <https://github.com/CesiumGS/3d-tiles/blob/master/3d-tiles-overview.pdf>.
- CityJSON. Cityjson specifications 1.0.1, 2019. URL <https://www.cityjson.org/specs/1.0.1/>.
- CityJSON. A web viewer for CityJSON files, 2020. URL <https://github.com/cityjson/ninja>.
- CityJSON. Datasets, 2020. URL <https://www.cityjson.org/datasets/>.
- Dcoetzee. File:Huffman tree.svg, 2007. URL https://commons.wikimedia.org/wiki/File:Huffman_tree.svg.
- M. Deering. Geometry compression. In *Proceedings of the 22nd annual conference on Computer graphics and interactive techniques*, pages 13–20, 1995.

Bibliography

- S. Dehpour. Deep Difference and search of any Python object/data., 2020. URL <https://github.com/seperman/deepdiff>.
- C. Dempsey. Web-based GIS, 2018. URL <https://www.gislounge.com/web-based-gis/>.
- Esri. Esri Indexed 3d Scene Layer (I3S) and Scene Layer Package (*.slpk) Format Specification, 2019a. URL <https://github.com/Esri/i3s-spec/blob/master/format/Index3d%20Scene%20Layer%20Format%20Specification.md>.
- Esri. Scene Layers: Service and Package Standard, 2019b. URL <https://github.com/Esri/i3s-spec>.
- A. Feldspar. An explanation of the deflate algorithm, 1997. URL <http://www.zlib.net/feldspar.html>.
- R. T. Fielding. Representational State Transfer (REST), 2000. URL https://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm.
- Flask. Flask - web development, one drop at a time, 2020. URL <https://flask.palletsprojects.com/en/1.1.x/>.
- F. Galligan. Draco Bitstream Specification, 2019. URL <https://google.github.io/draco/spec/>.
- GeoJSON. GeoJSON. URL <https://geojson.org/>.
- Google. mesh_prediction_scheme_parallelogram_encoder.h, 2018. URL https://github.com/google/draco/blob/master/src/draco/compression/attributes/prediction_schemes/mesh_prediction_scheme_parallelogram_encoder.h.
- Google. mesh_sequential_encoder.h, 2019a. URL https://github.com/google/draco/blob/master/src/draco/compression/mesh/mesh_sequential_encoder.h.
- Google. Draco 3D data compression, 2019b. URL <https://google.github.io/draco/>.
- Google. Draco 3D data compression, 2019c. URL <https://github.com/google/draco>.
- Google. Draco C++ decoder API, 2020a. URL <https://github.com/google/draco/blob/master/src/draco/mesh/mesh.h>.
- Google. Draco JavaScript decoder API, 2020b. URL https://github.com/google/draco/blob/master/src/draco/javascript/emscripten/draco_web_decoder.idl.
- Google. Mesh edgebreaker traversal predictive encoder, 2020c. URL https://github.com/google/draco/blob/master/src/draco/compression/mesh/mesh_edgebreaker_traversal_predictive_encoder.h.
- HERE. harp.gl: 3d web map rendering engine, 2019. URL <https://www.harp.gl/>.
- J. Hildebrand. Encode and decode CBOR documents, with both easy mode, streaming mode, and SAX-style evented mode., 2020. URL <https://github.com/hildjj/node-cbor>.
- D. A. Huffman. A method for the construction of minimum-redundancy codes. *Proceedings of the IRE*, 40(9):1098–1101, 1952.
- Introducing JSON. Introducing JSON, 2019. URL <https://www.json.org/json-en.html>.
- KeyCDN. What is latency and how to reduce it, 2018. URL <https://www.keycdn.com/support/what-is-latency>.
- Khronos Group. gltf0verview-2.0.0b.png, 2019. URL <https://raw.githubusercontent.com/KhronosGroup/gltf/master/specification/2.0/figures/gltf0verview-2.0.0b.png>.

- H. Ledoux. TU Delft campus dataset created with 3dfier. private communication, 2019.
- H. Ledoux. CityJSON + RESTful access + streaming, 2020. URL https://github.com/hugoledoux/cityjson_ogcapi/blob/master/best-practice.md.
- H. Ledoux, K. Ohori, K. Kumar, B. Dukai, A. Labetski, and S. Vitalis. CityJSON: A compact and easy-to-use encoding of the CityGML data model. *Open Geospatial Data, Software and Standards*, 4(1):4, 2019.
- N. Lohmann. JSON for Modern C++, 2020. URL <https://github.com/nlohmann/json>.
- Mango. Web gis, 2017. URL <https://mangomap.com/web-gis>.
- MessagePack. MessagePack, 2019. URL <https://msgpack.org/>.
- MongoDB. The database for modern applications, 2020. URL <https://www.mongodb.com/>.
- Mozilla. Base64, 2020. URL <https://developer.mozilla.org/en-US/docs/Glossary/Base64>.
- Municipality of Rotterdam. Rotterdam multi-LoDs CityGML buildings, 2020. URL <https://3d.bktudelft.nl/projects/geobim-benchmark/rotterdamlod12.html>.
- B. Muthukadan. Selenium with Python, 2020. URL <https://selenium-python.readthedocs.io/installation.html#introduction>.
- OGC API. Welcome To The OGC APIs, 2020. URL <http://www.ogcapi.org/>.
- Open Geospatial Consortium. CityGML. URL <https://www.ogc.org/standards/citygml>.
- Open Geospatial Consortium. OGC API - Features, an overview, 2019a. URL <https://github.com/opengeospatial/ogcapi-features/blob/master/overview.md>.
- Open Geospatial Consortium. OGC API - Features - Part 1: Core, 2019b. URL http://docs.opengeospatial.org/is/17-069r3/17-069r3.html#_items_.
- Open Geospatial Consortium. OGC API - Features, 2019c. URL <https://www.opengeospatial.org/standards/ogcapi-features>.
- Open Geospatial Consortium. OGC API - Features, 2020. URL <https://github.com/opengeospatial/ogcapi-features>.
- Proj4js. Proj4js: JavaScript library to transform coordinates from one coordinate system to another, including datum transformations, 2020. URL <http://proj4js.org/>.
- M. Reddy. Object Files (.obj). URL <http://www.martinreddy.net/gfx/3d/0BJ.spec>.
- J. Rossignac. 3D compression made simple: Edgebreaker with ZipandWrap on a corner-table. In *Proceedings International Conference on Shape Modeling and Applications*, pages 278–283. IEEE, 2001.
- J. Rossignac, A. Safanova, and A. Szymczak. Edgebreaker on a corner table: A simple technique for representing and compressing triangulated surfaces. In *Hierarchical and geometrical methods in scientific visualization*, pages 41–50. Springer, 2003.
- S. Sanfilippo. smaz, 2019a. URL <https://github.com/antirez/smaz>.
- S. Sanfilippo. smaz, 2019b. URL <https://github.com/antirez/smaz/blob/master/smaz.c>.
- K. Sayood. *Introduction to data compression*. Morgan Kaufmann, 2017.
- Selenium. Selenium Projects, 2020. URL <https://www.selenium.dev/projects/>.

Bibliography

- Shapely. The Shapely User Manual, 2020. URL <https://shapely.readthedocs.io/en/latest/manual.html>.
- Socket.io. Socket.io 2.0 is here, 2020. URL <https://socket.io/>.
- T. Suel. Delta compression techniques., 2019.
- G. Taubin, W. P. Horn, F. Lazarus, and J. Rossignac. Geometry coding and vrml. *Proceedings of the IEEE*, 86(6):1228–1243, 1998.
- Technische Universität München. 3dcitydb database, 2020. URL <https://www.3dcitydb.org/3dcitydb/>.
- The Linux Information Project. Binary File Definition, 2006. URL http://www.linfo.org/binary_file.html.
- Three.js. three.js, 2020a. URL <https://threejs.org/>.
- Three.js. Mesh, 2020b. URL <https://threejs.org/docs/#api/en/objects/Mesh>.
- Three.js. BufferGeometry, 2020c. URL <https://threejs.org/docs/#api/en/core/BufferGeometry>.
- C. Touma and C. Gotsman. Triangle mesh compression. In *Proceedings of the Graphics Interface 1998 Conference, June 18-20, 1998, Vancouver, BC, Canada*, pages 26–34, June 1998. URL <http://graphicsinterface.org/wp-content/uploads/gi1998-4.pdf>.
- Turf.js. Welcome to Turf.js, 2020. URL <https://turfjs.org/>.
- UBJSON. Universal binary json specification: The universally compatible format specification for binary json., 2020. URL <http://ubjson.org/>.
- WebGL. Getting Started, 2011. URL https://www.khronos.org/webgl/wiki/Getting_Started.
- S. Yura. CBOR bindings for Python with (at the moment) simple streaming interface , 2016. URL <https://github.com/funny-falcon/flunn>.
- I. Zderadicka. Comparison of json like serializations – json vs ubjson vs messagepack vs cbor, 2017. URL <http://zderadicka.eu/comparison-of-json-like-serializations-json-vs-ubjson-vs-messagepack-vs-cbor/>.
- zlib. zlib - A Massively Spiffy Yet Delicately Unobtrusive Compression Library, 2020. URL <https://zlib.net/>.

Colophon

This document was typeset using L^AT_EX, using the KOMA-Script class `scrbook`. The main font is Palatino.

